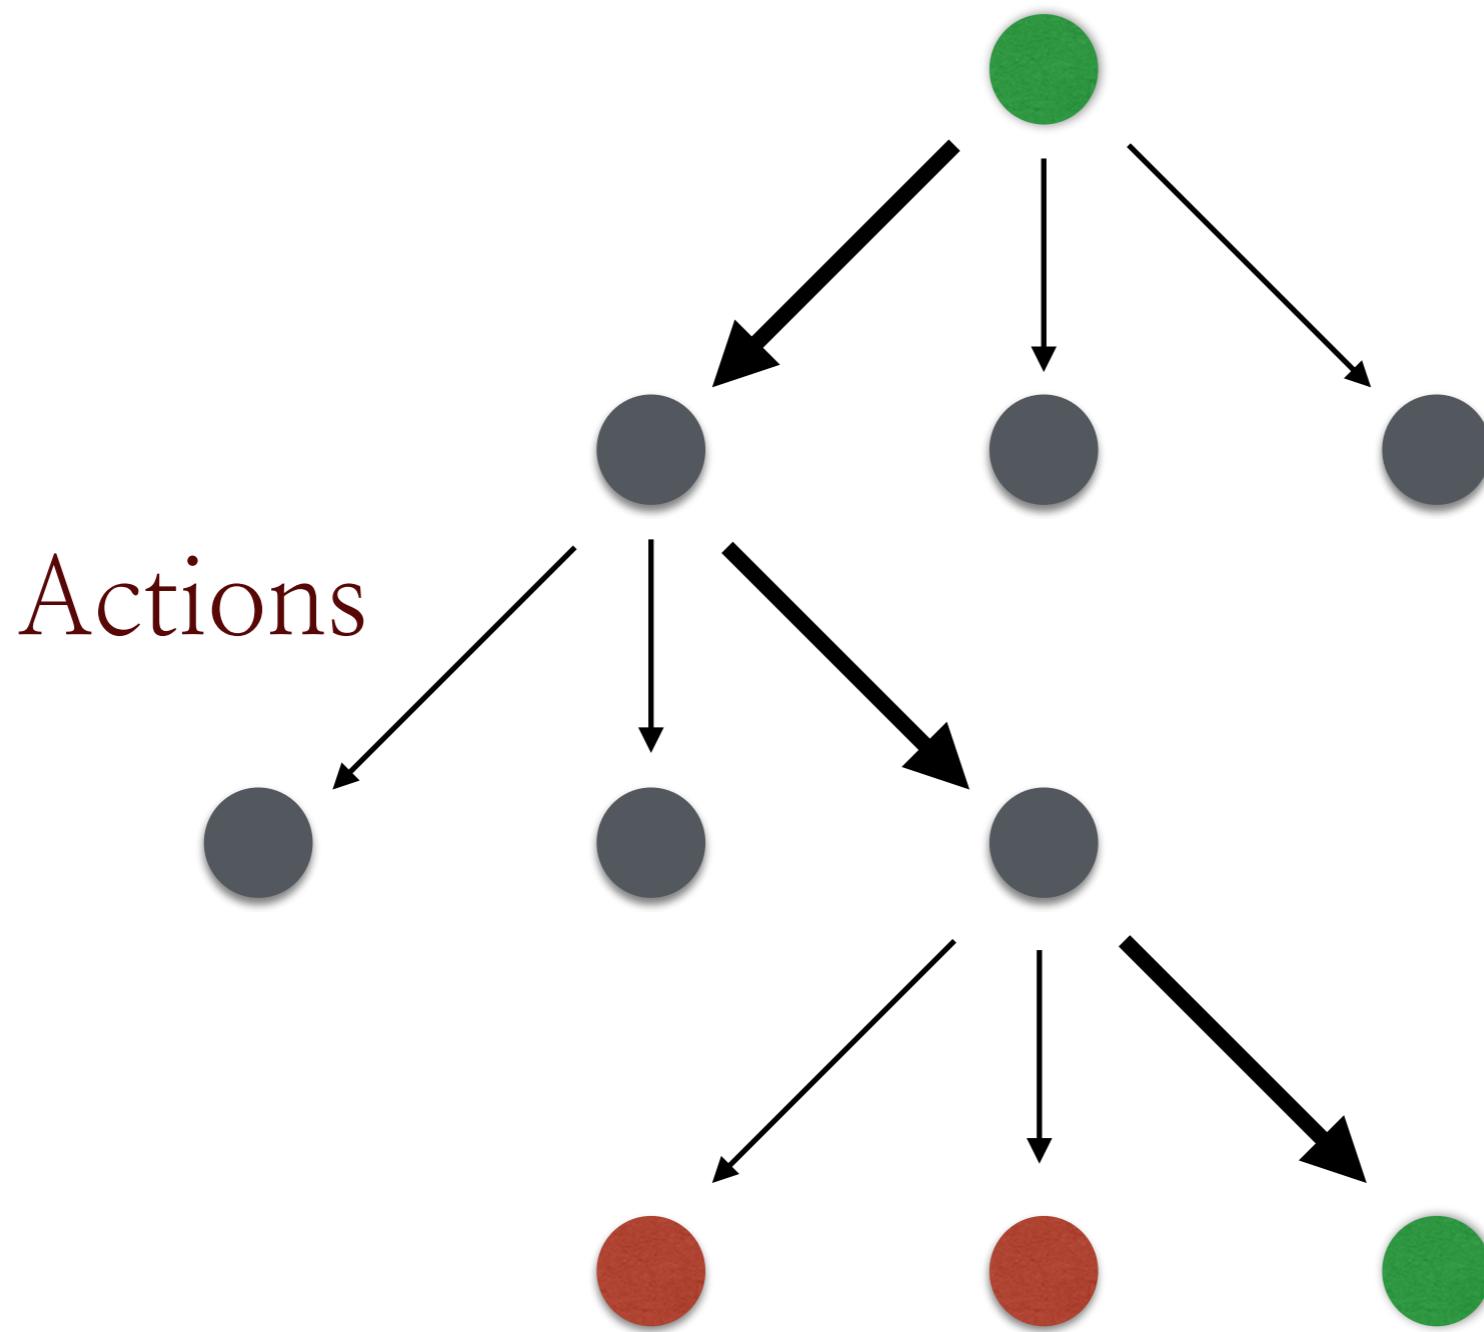


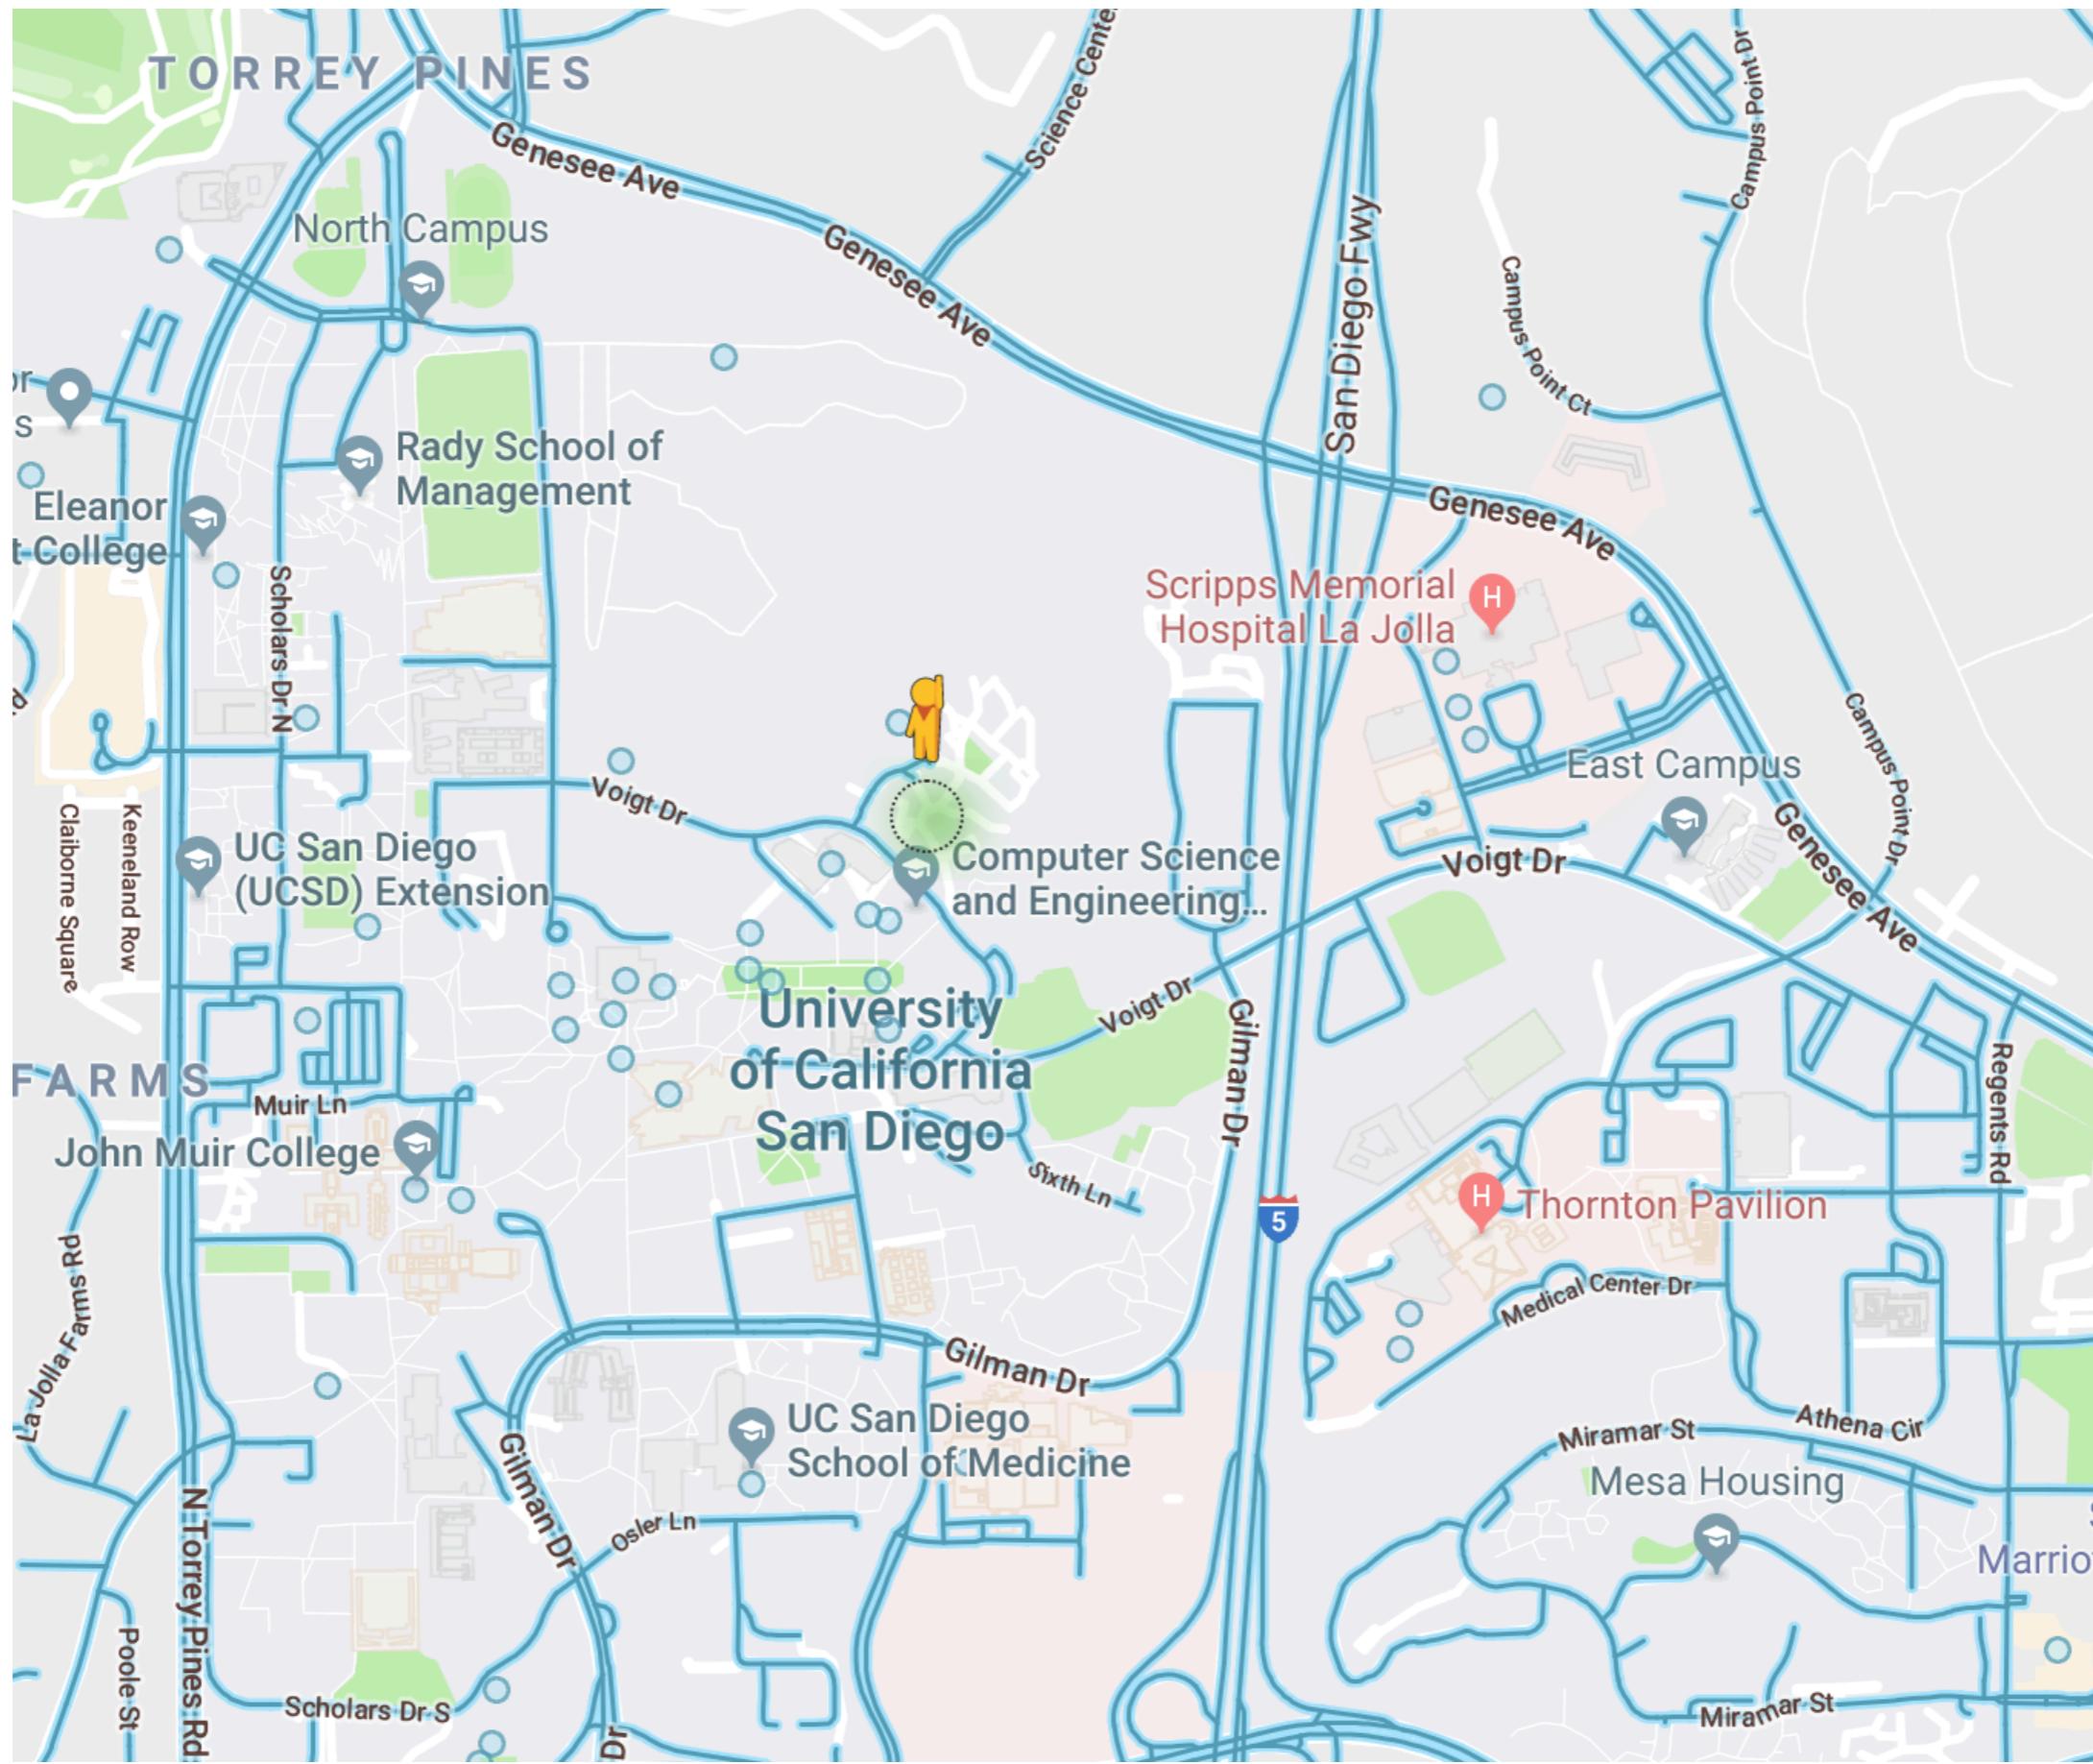
# Classical Search

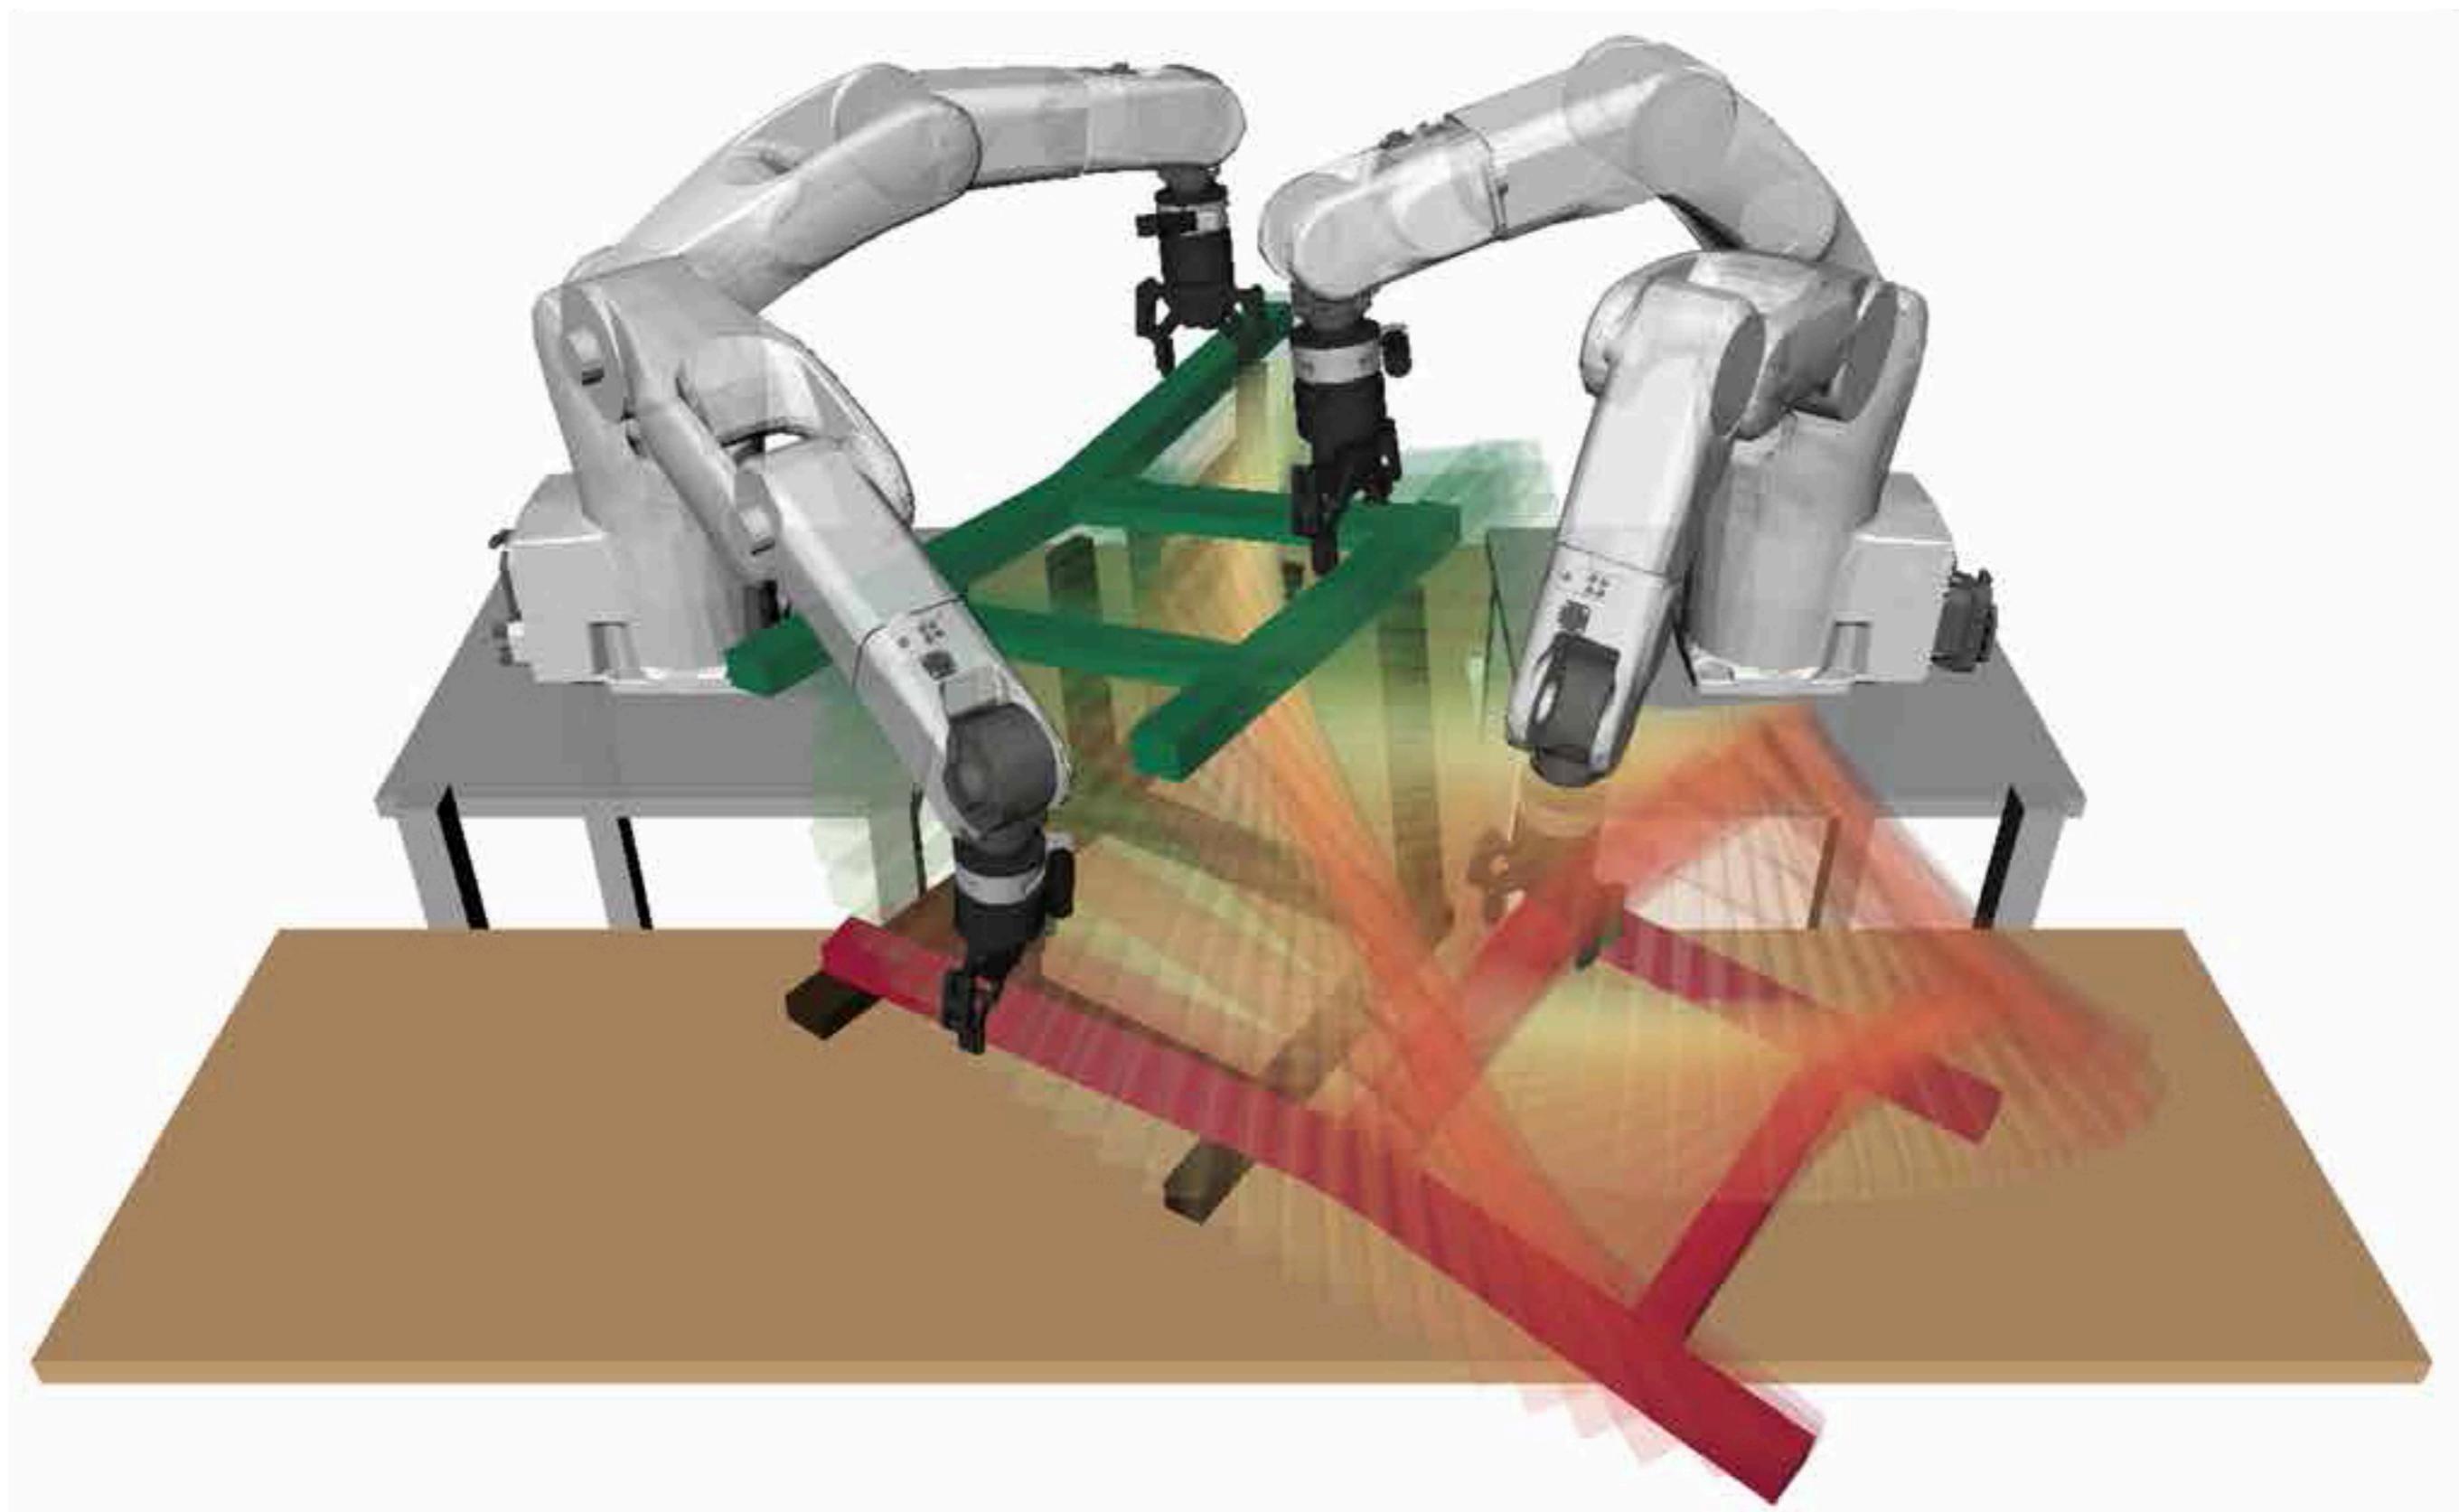
UCSD CSE 257  
Sicun Gao

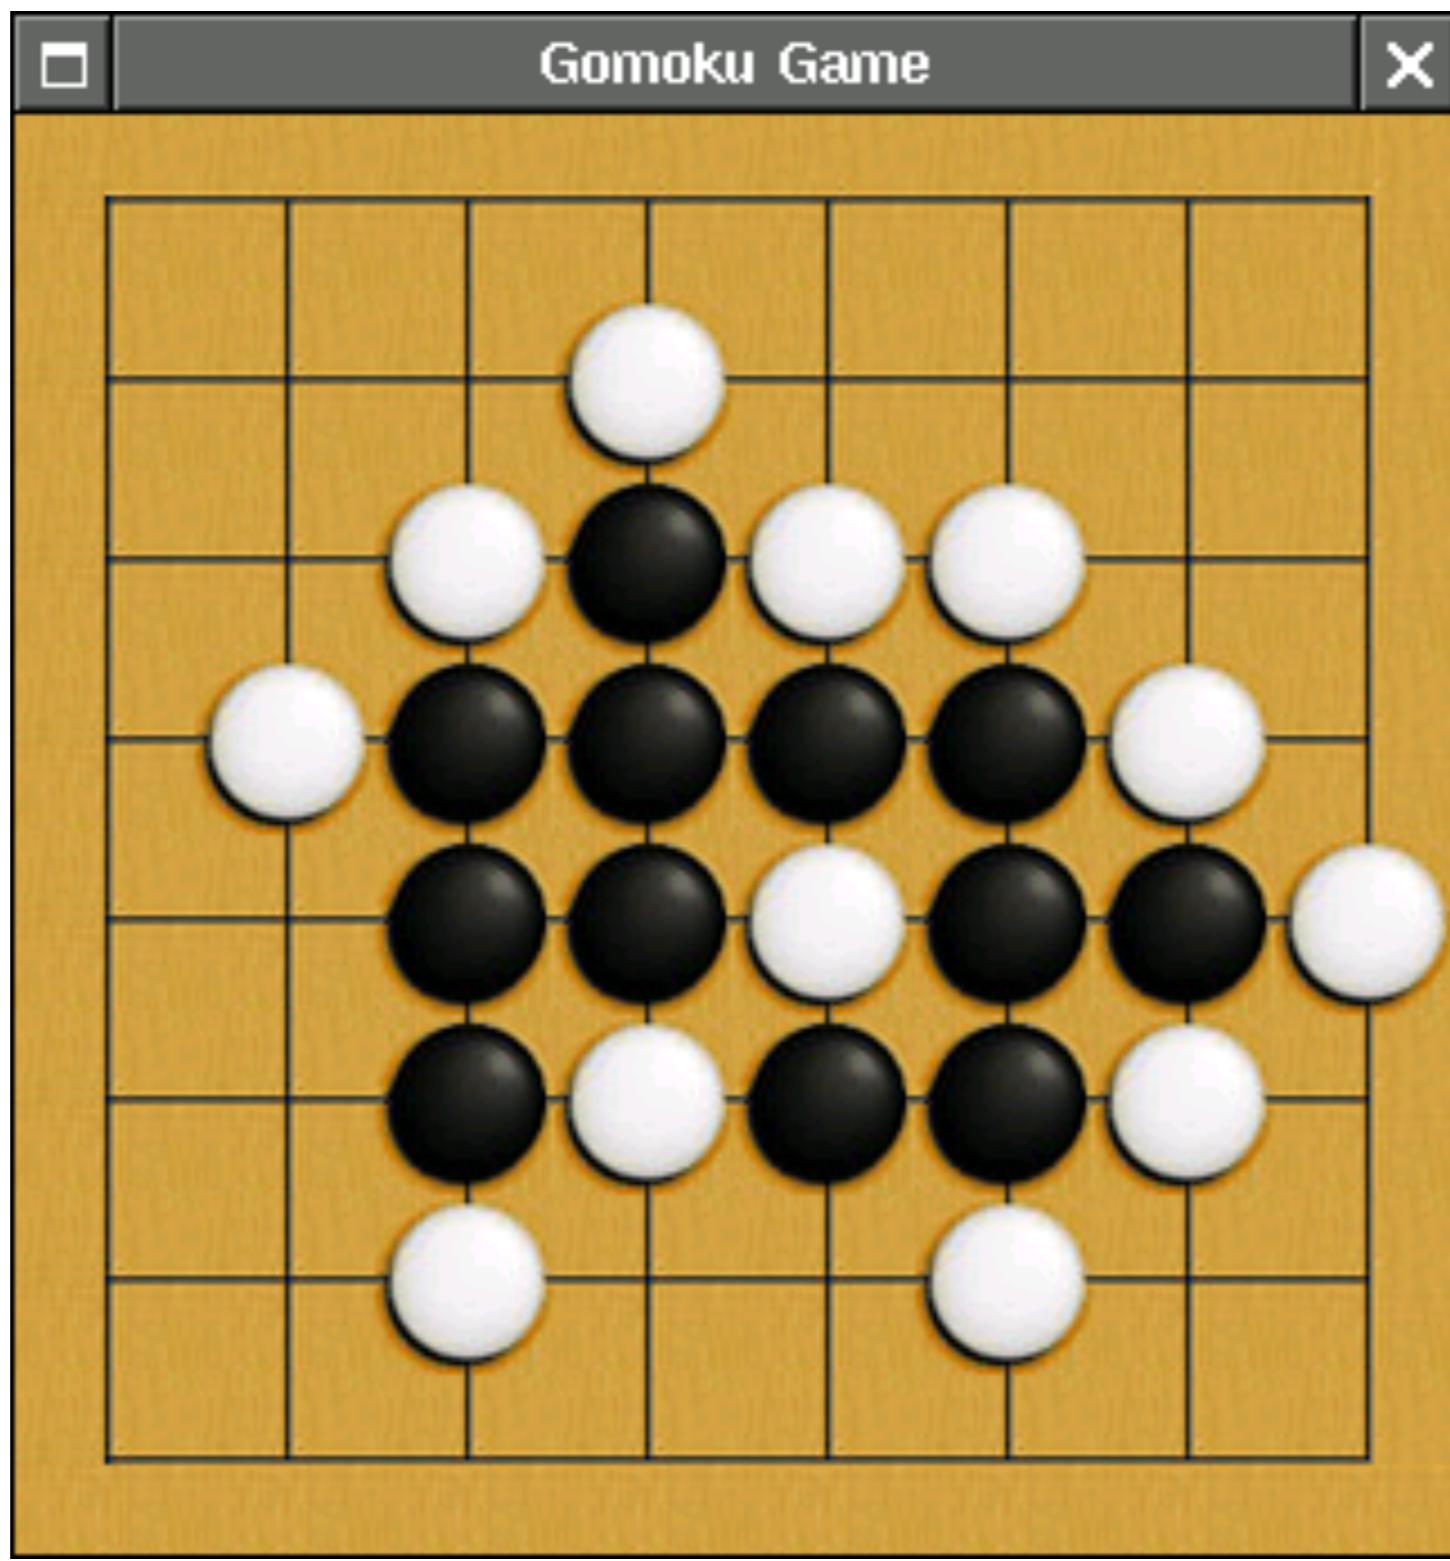
Initial State



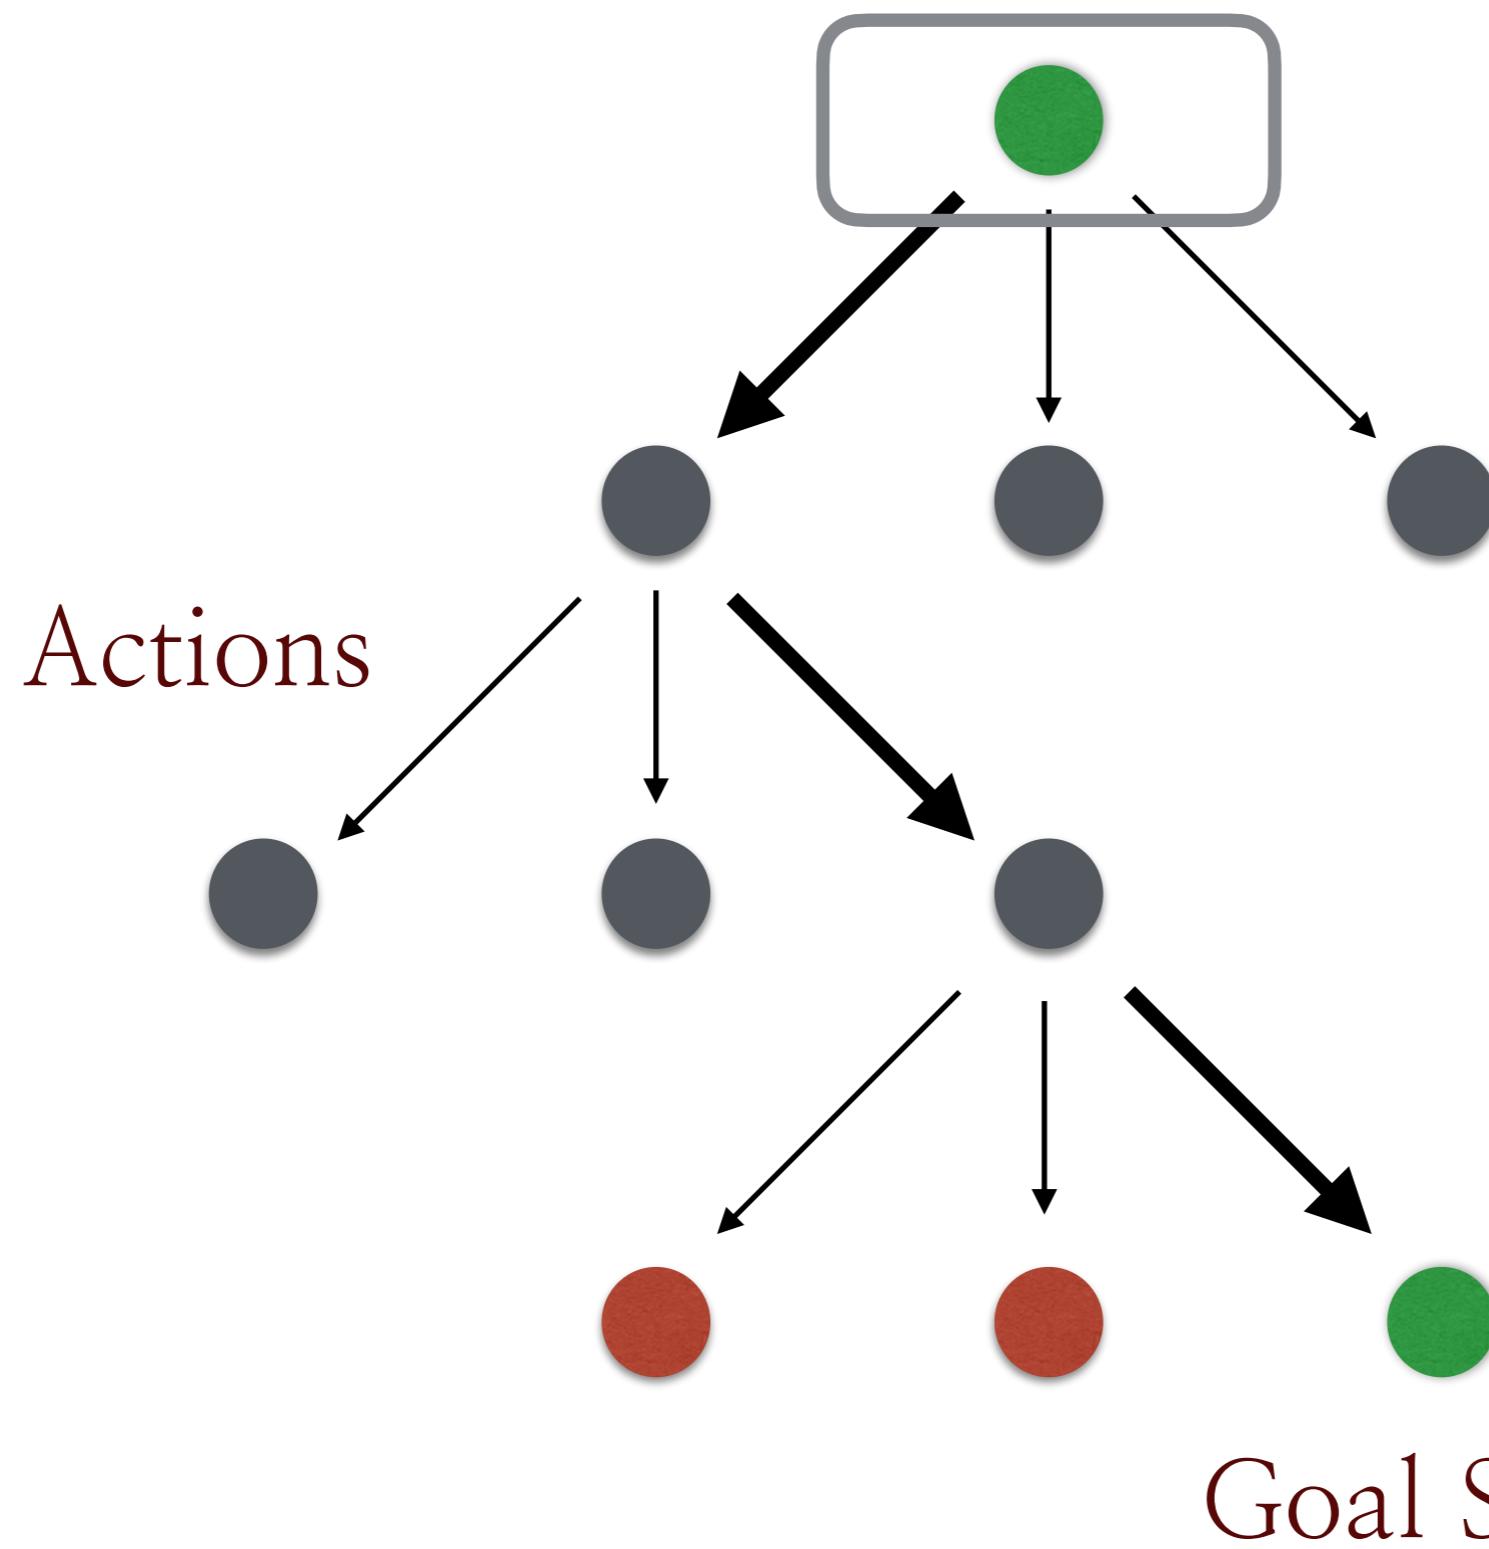
Goal State





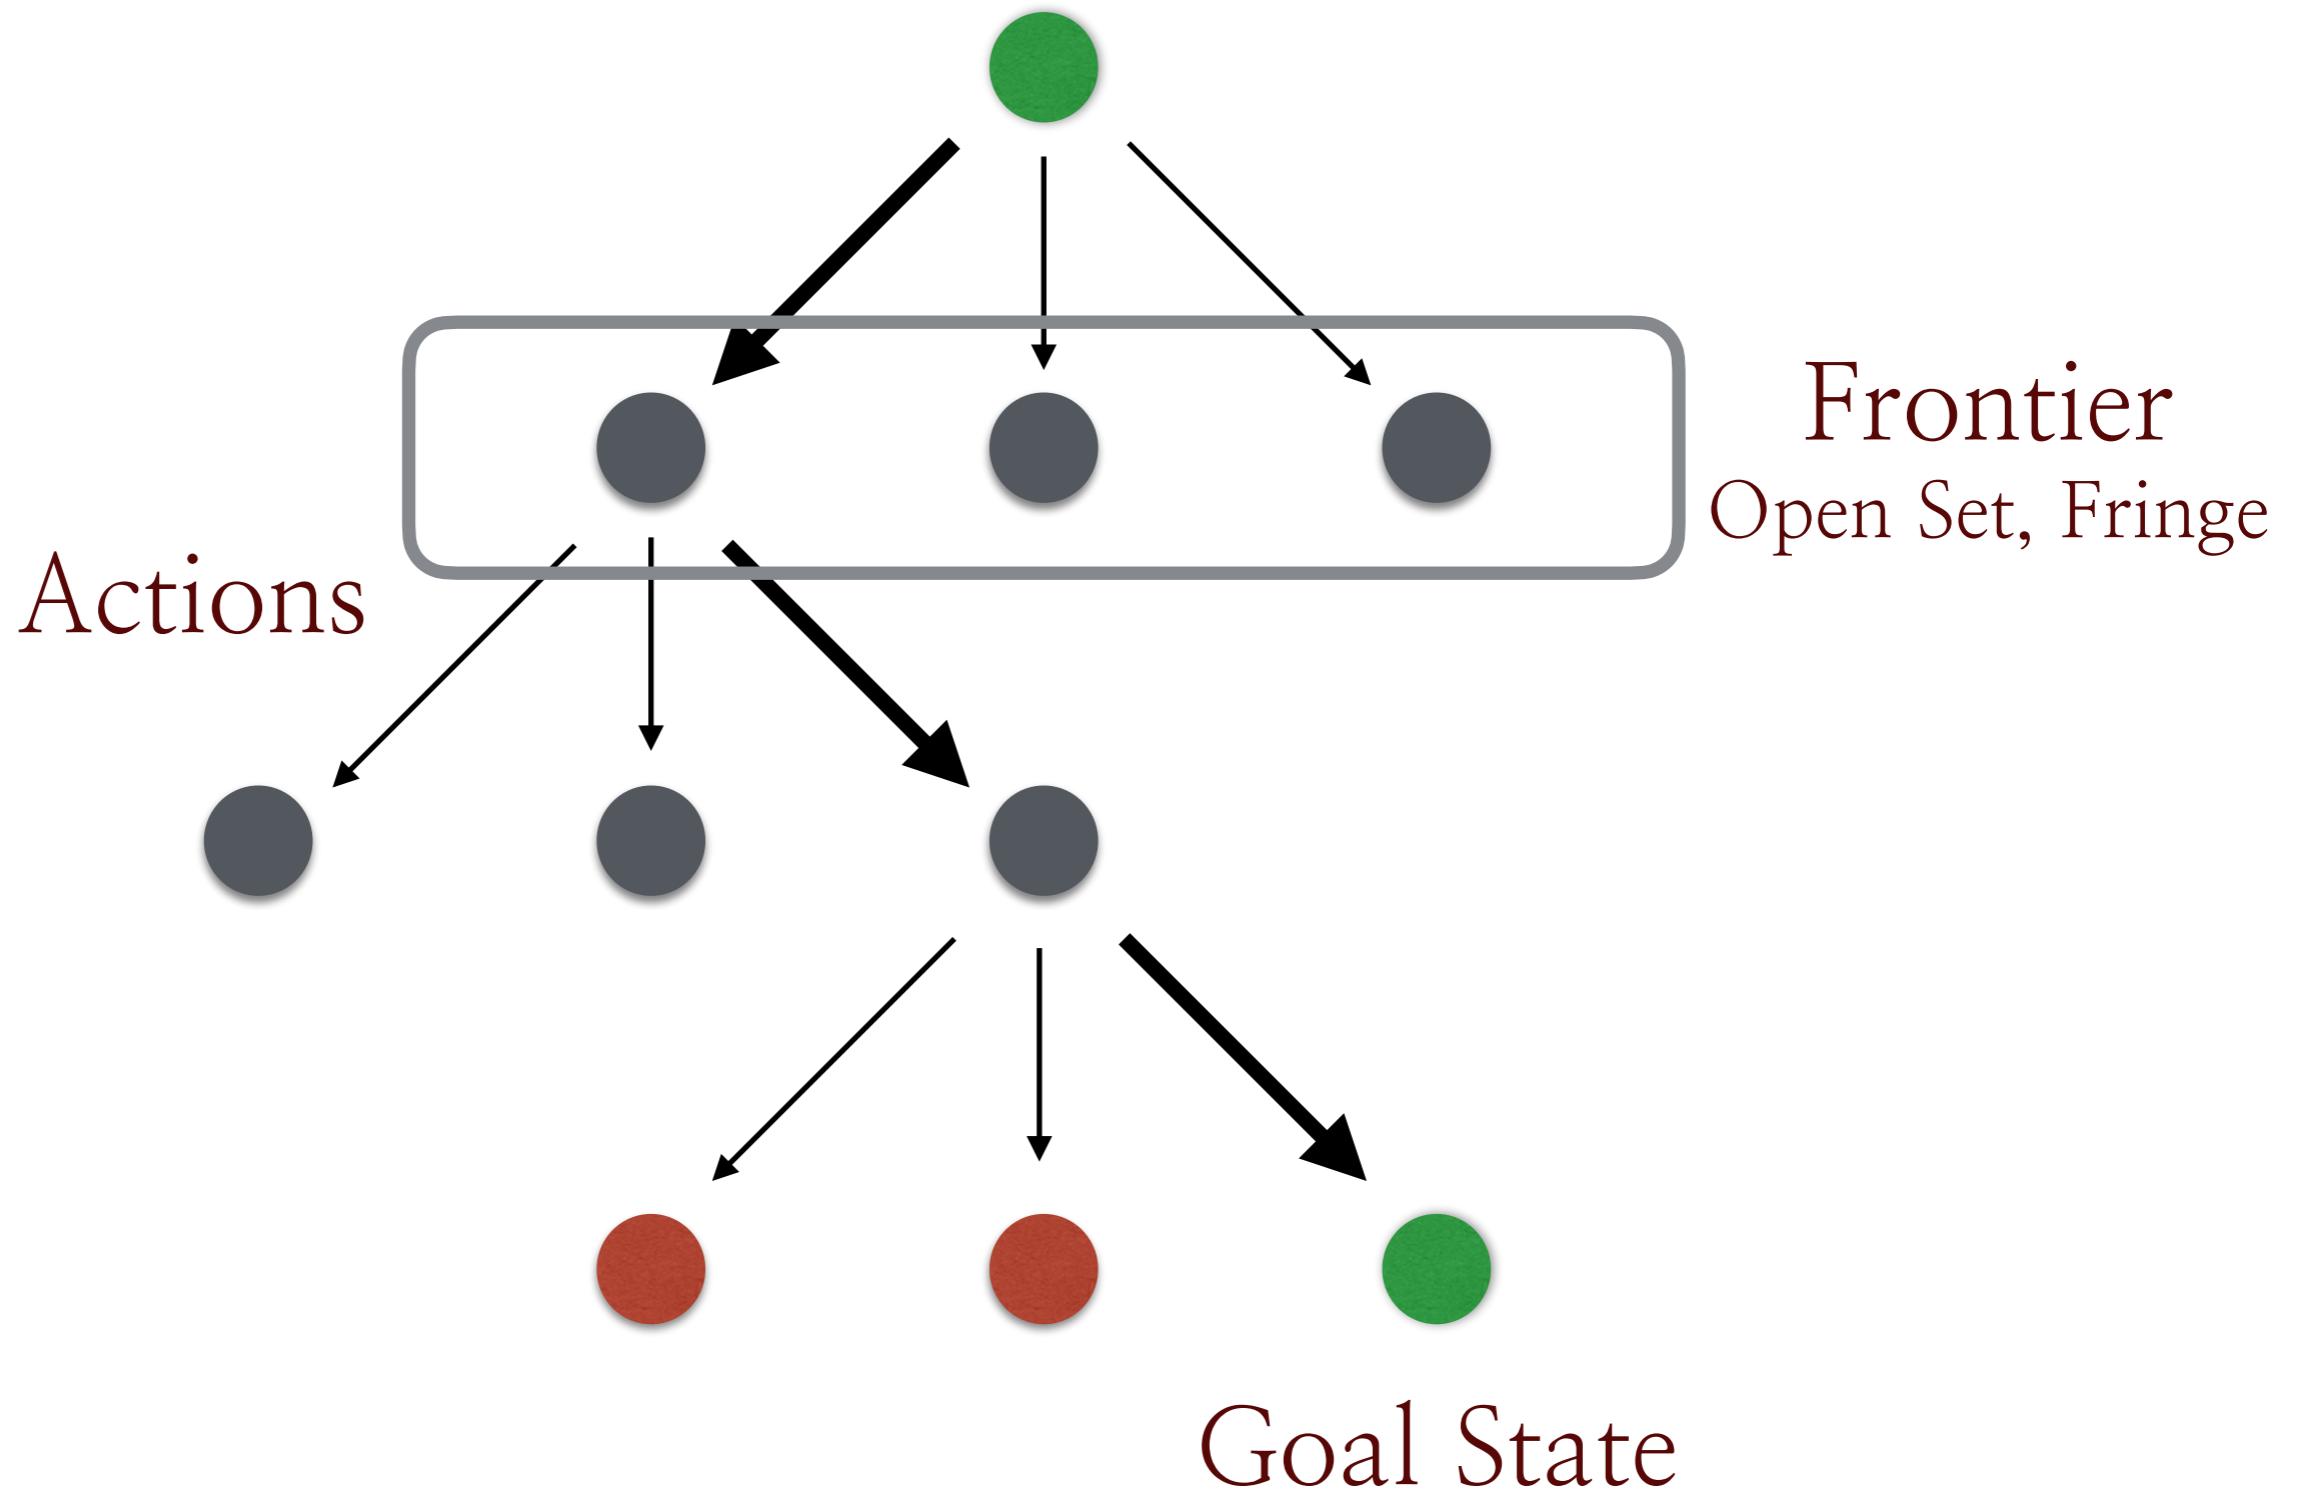


Initial State

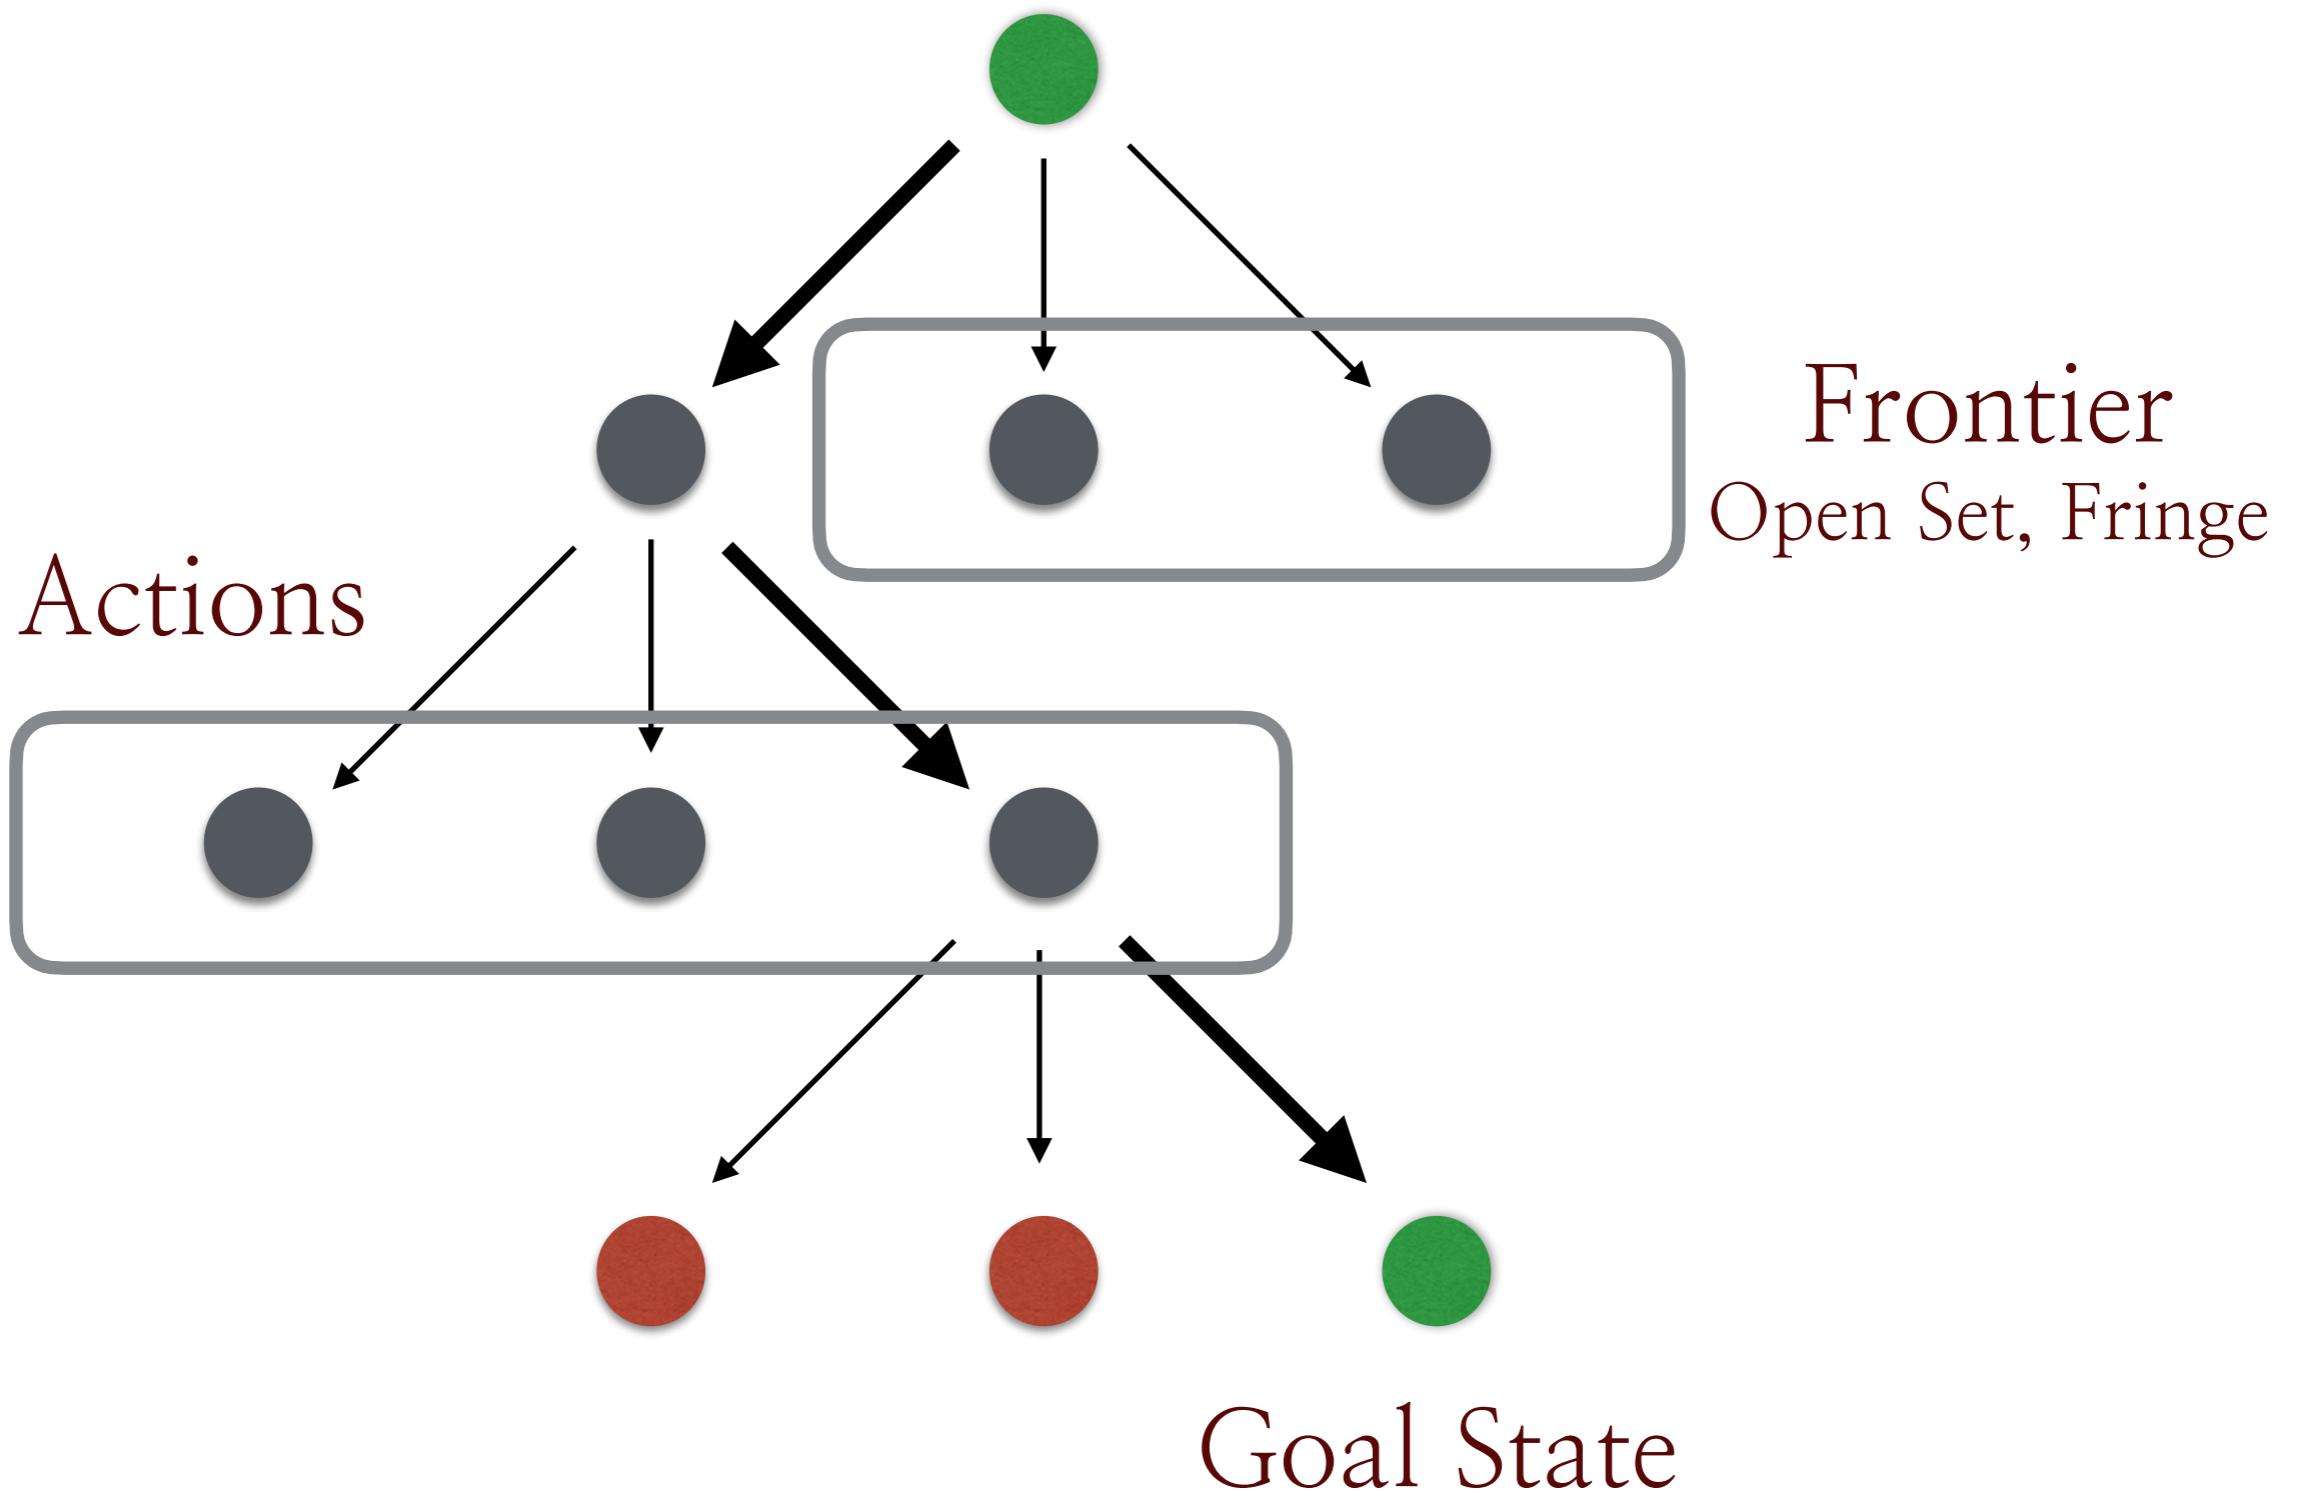


Frontier  
Open Set, Fringe

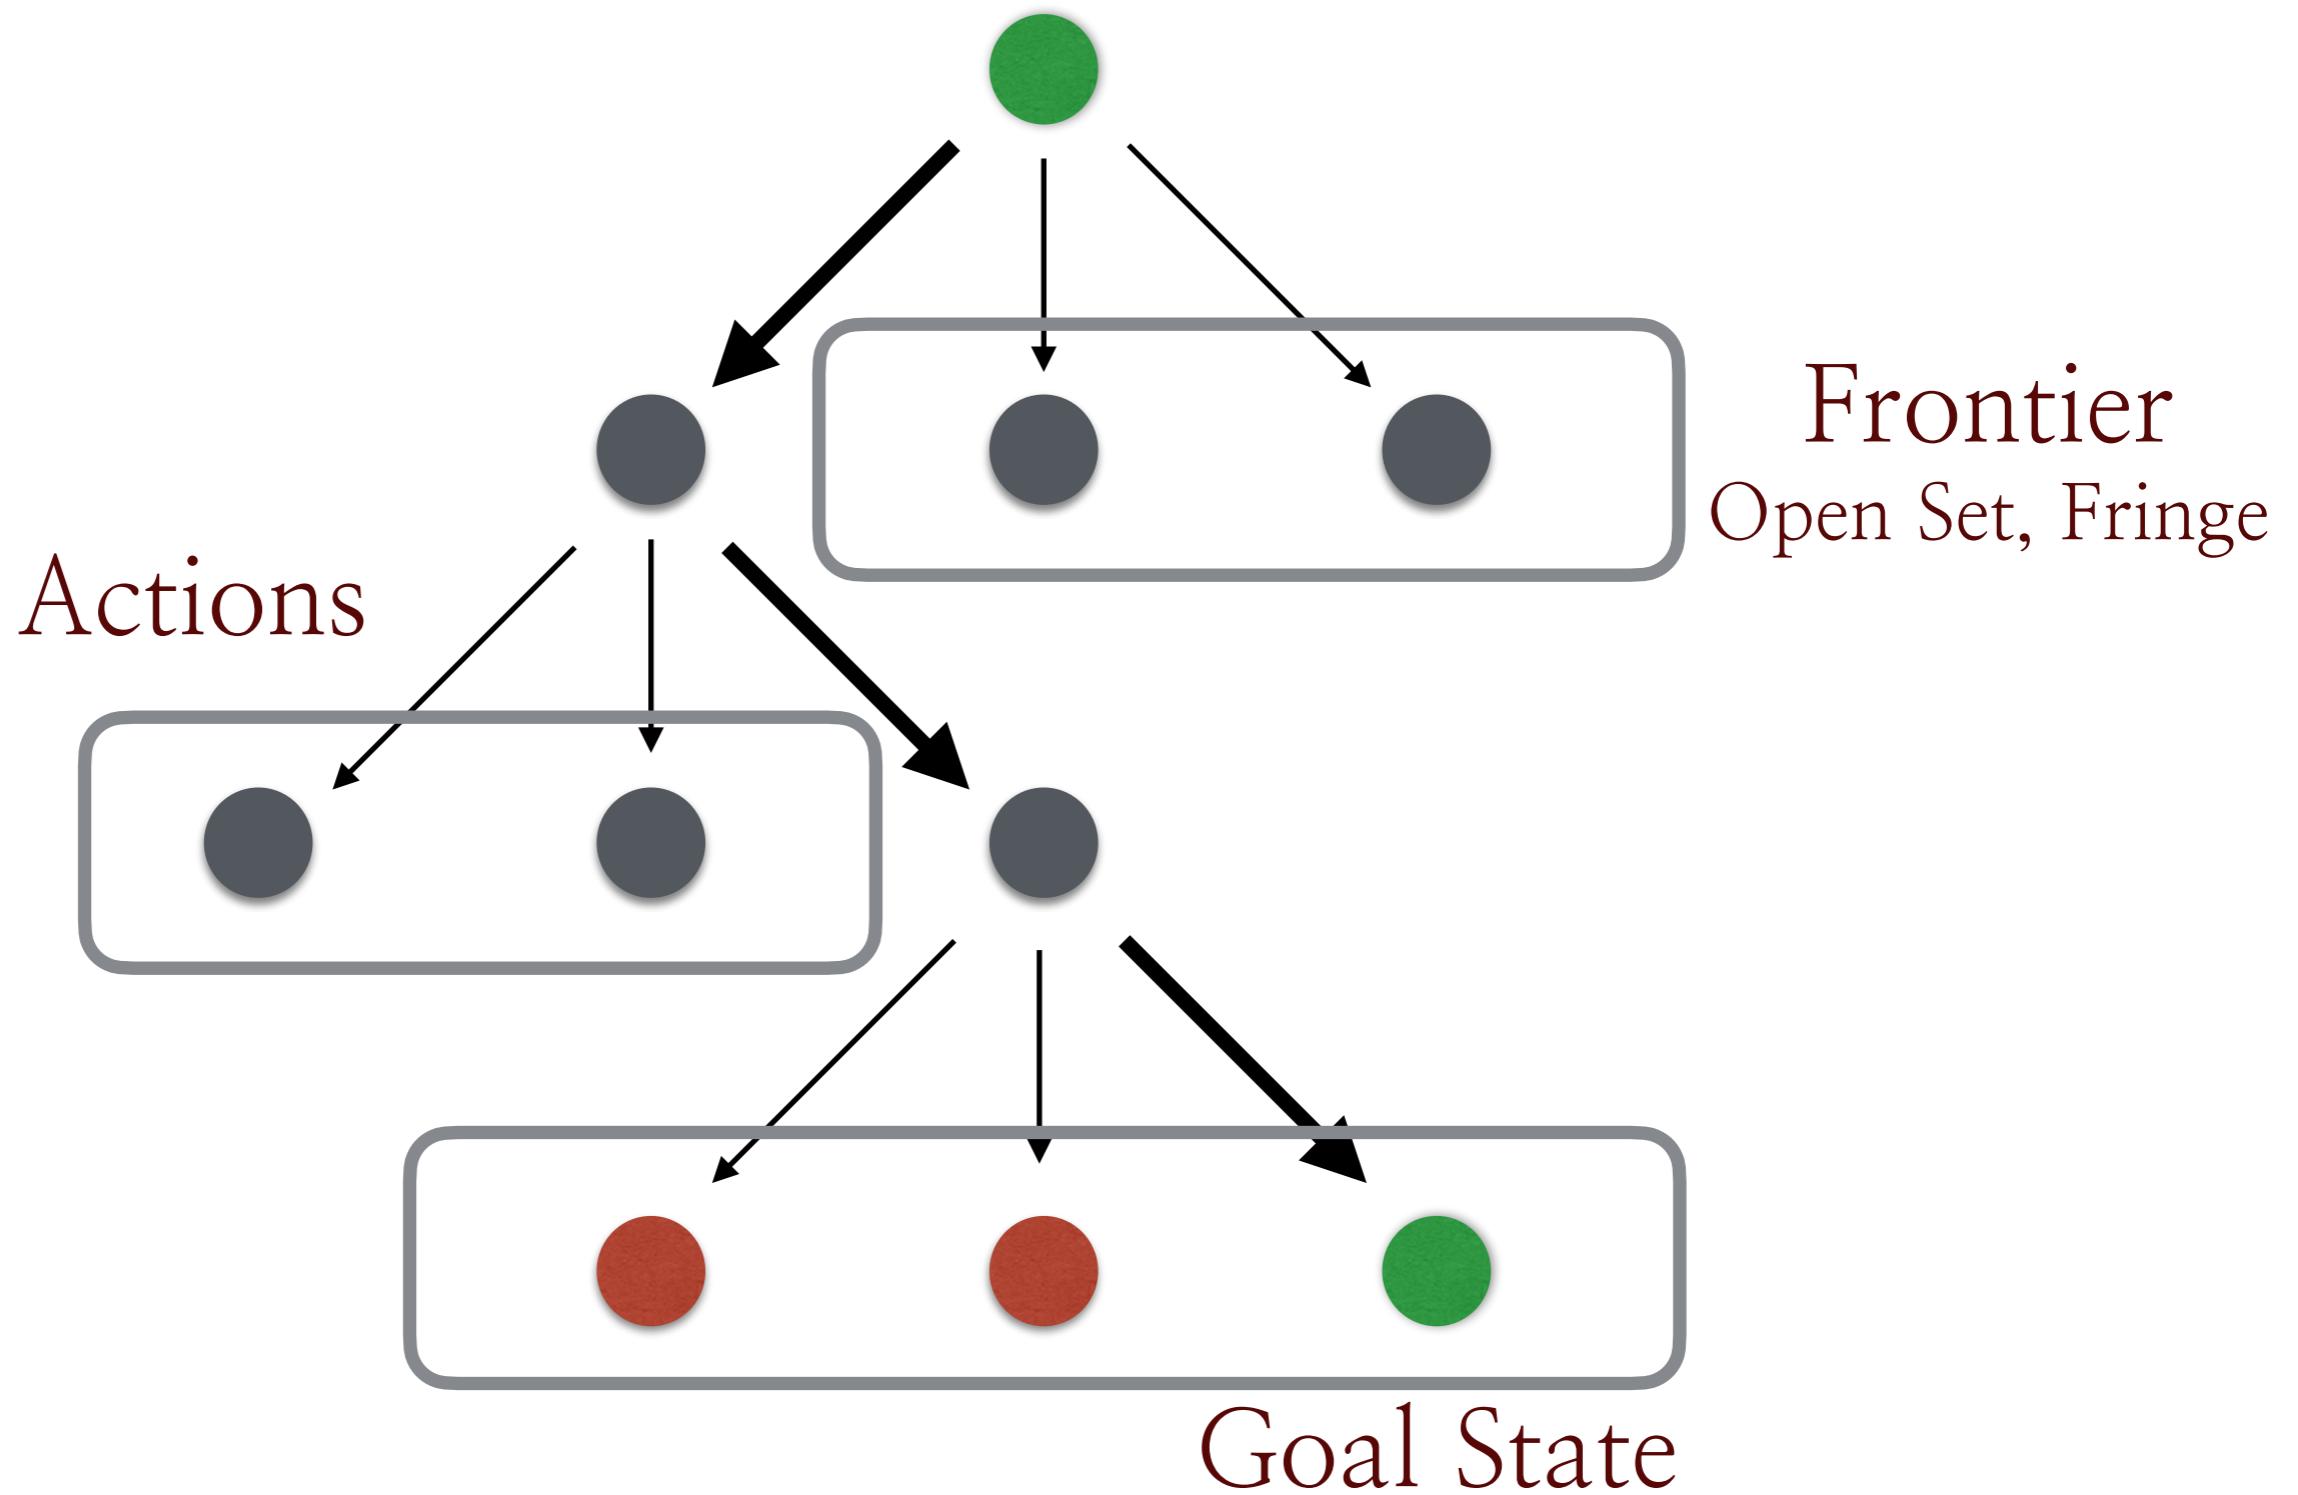
Initial State



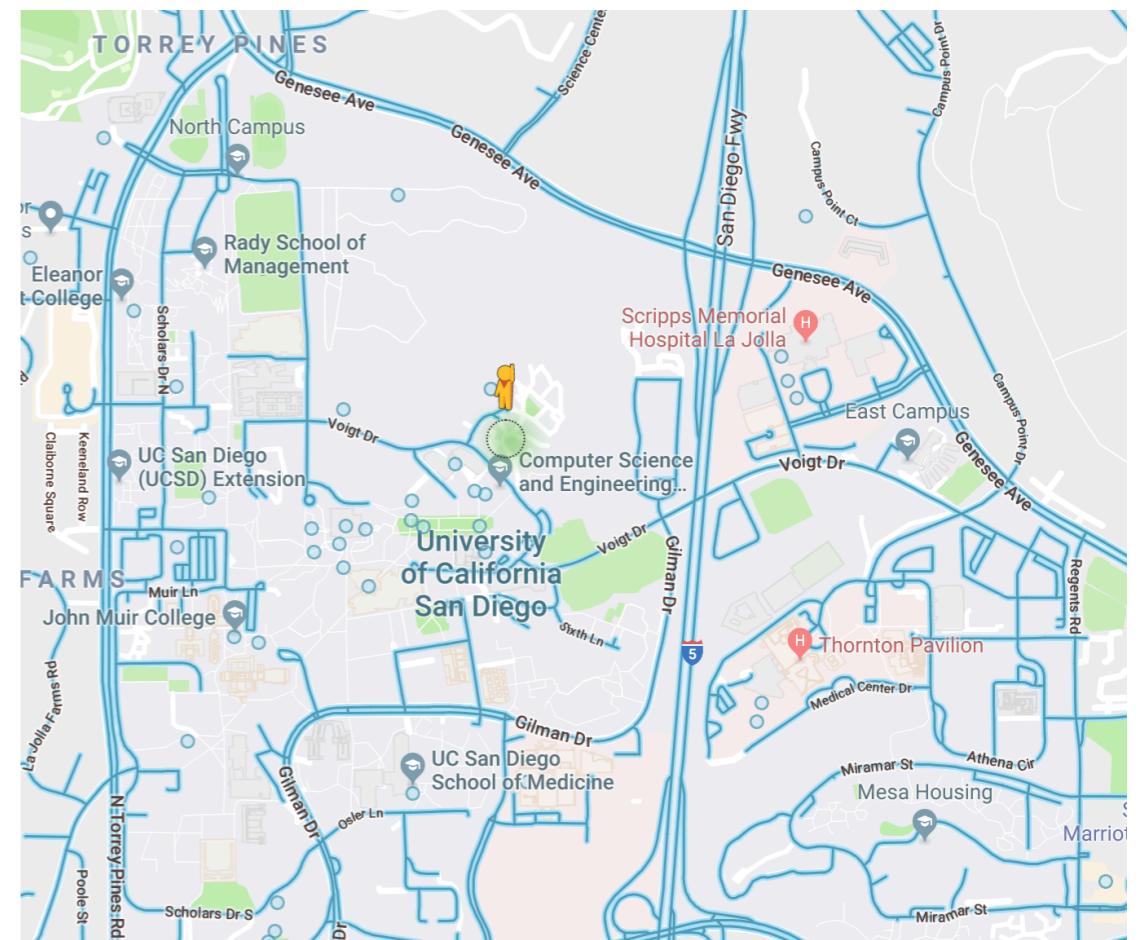
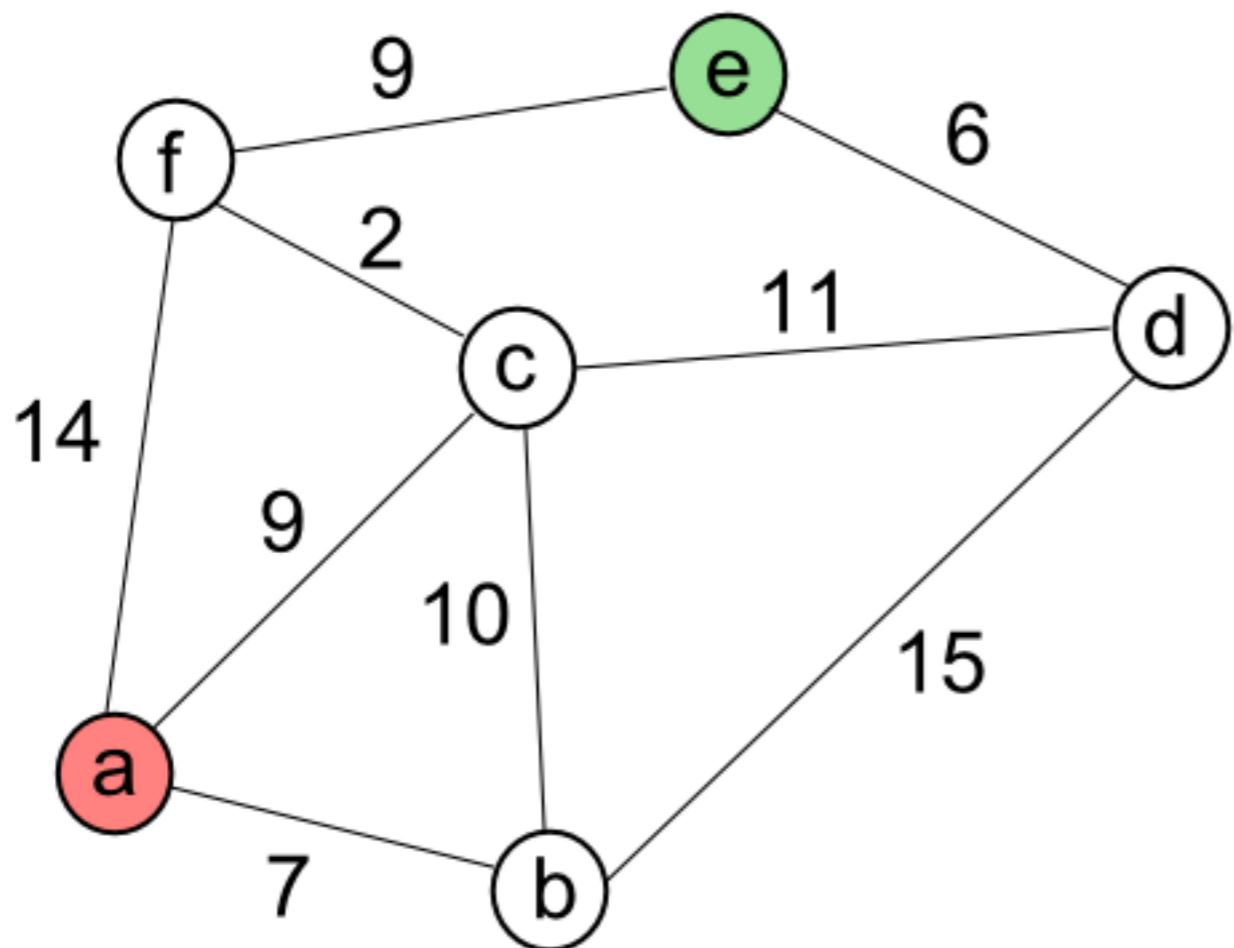
Initial State



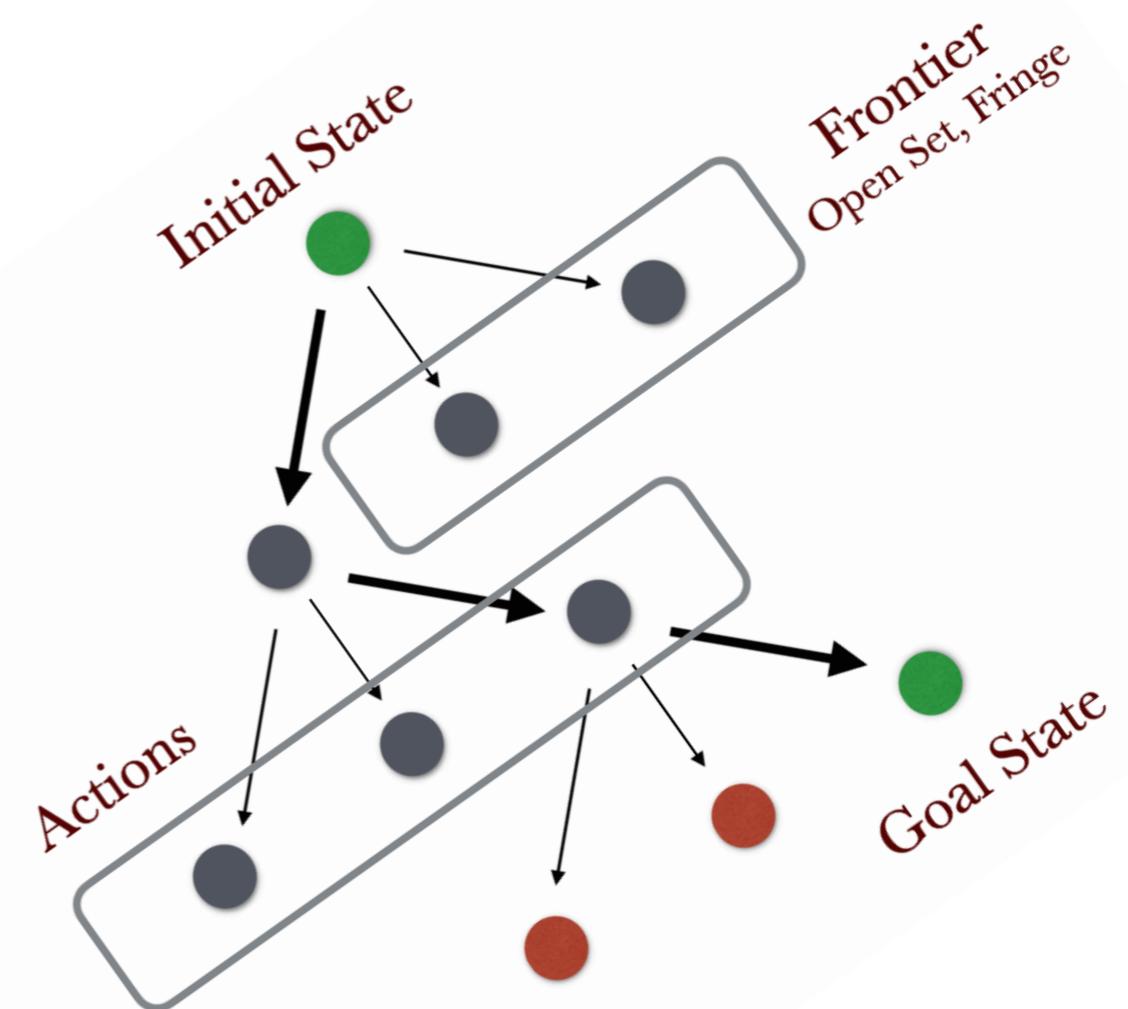
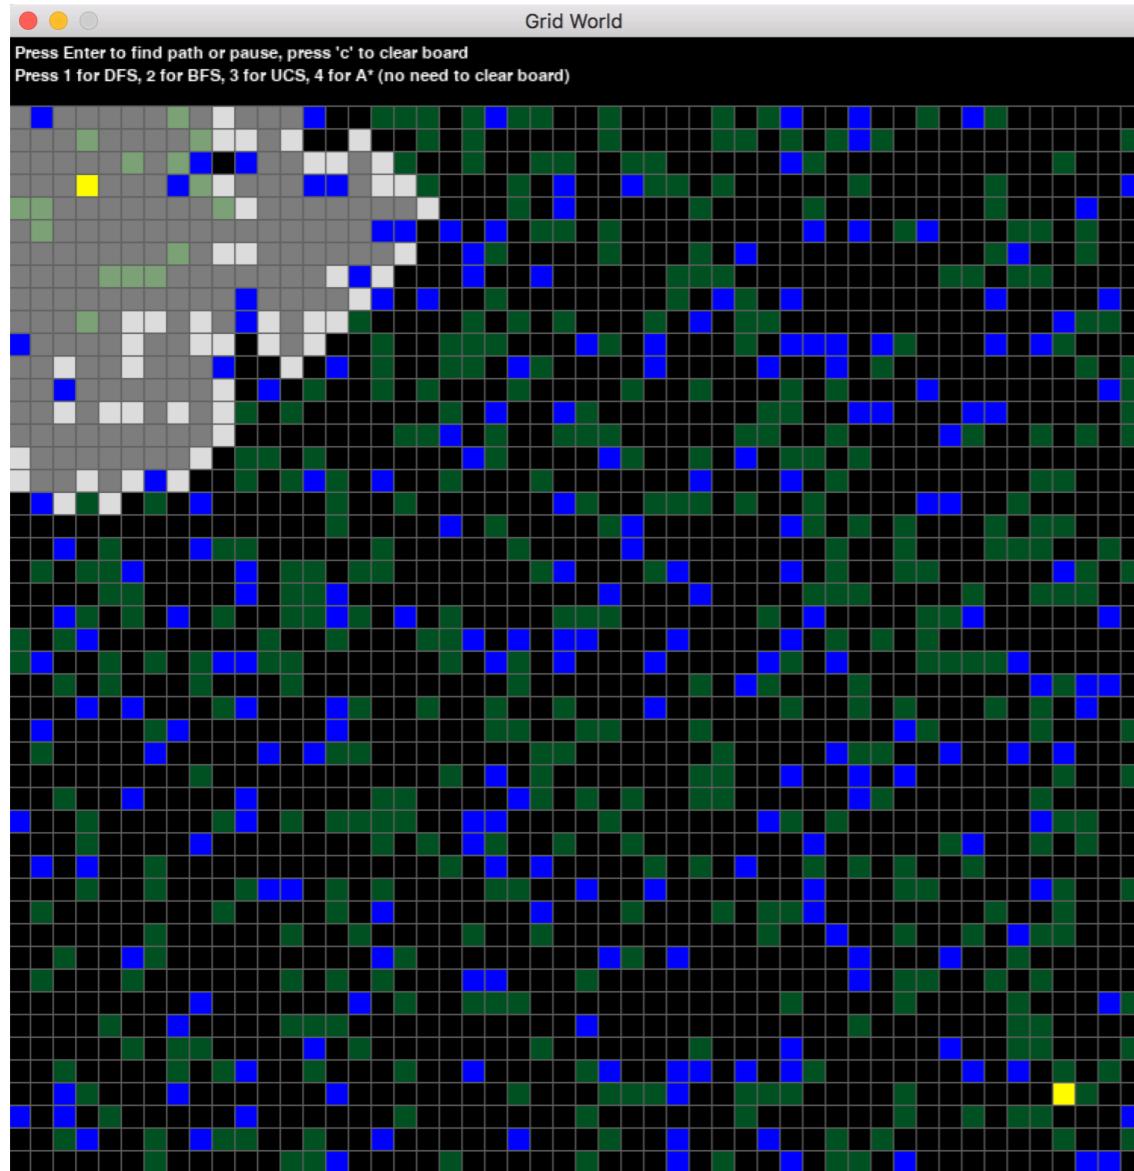
DFS? BFS?



# Search in Graphs



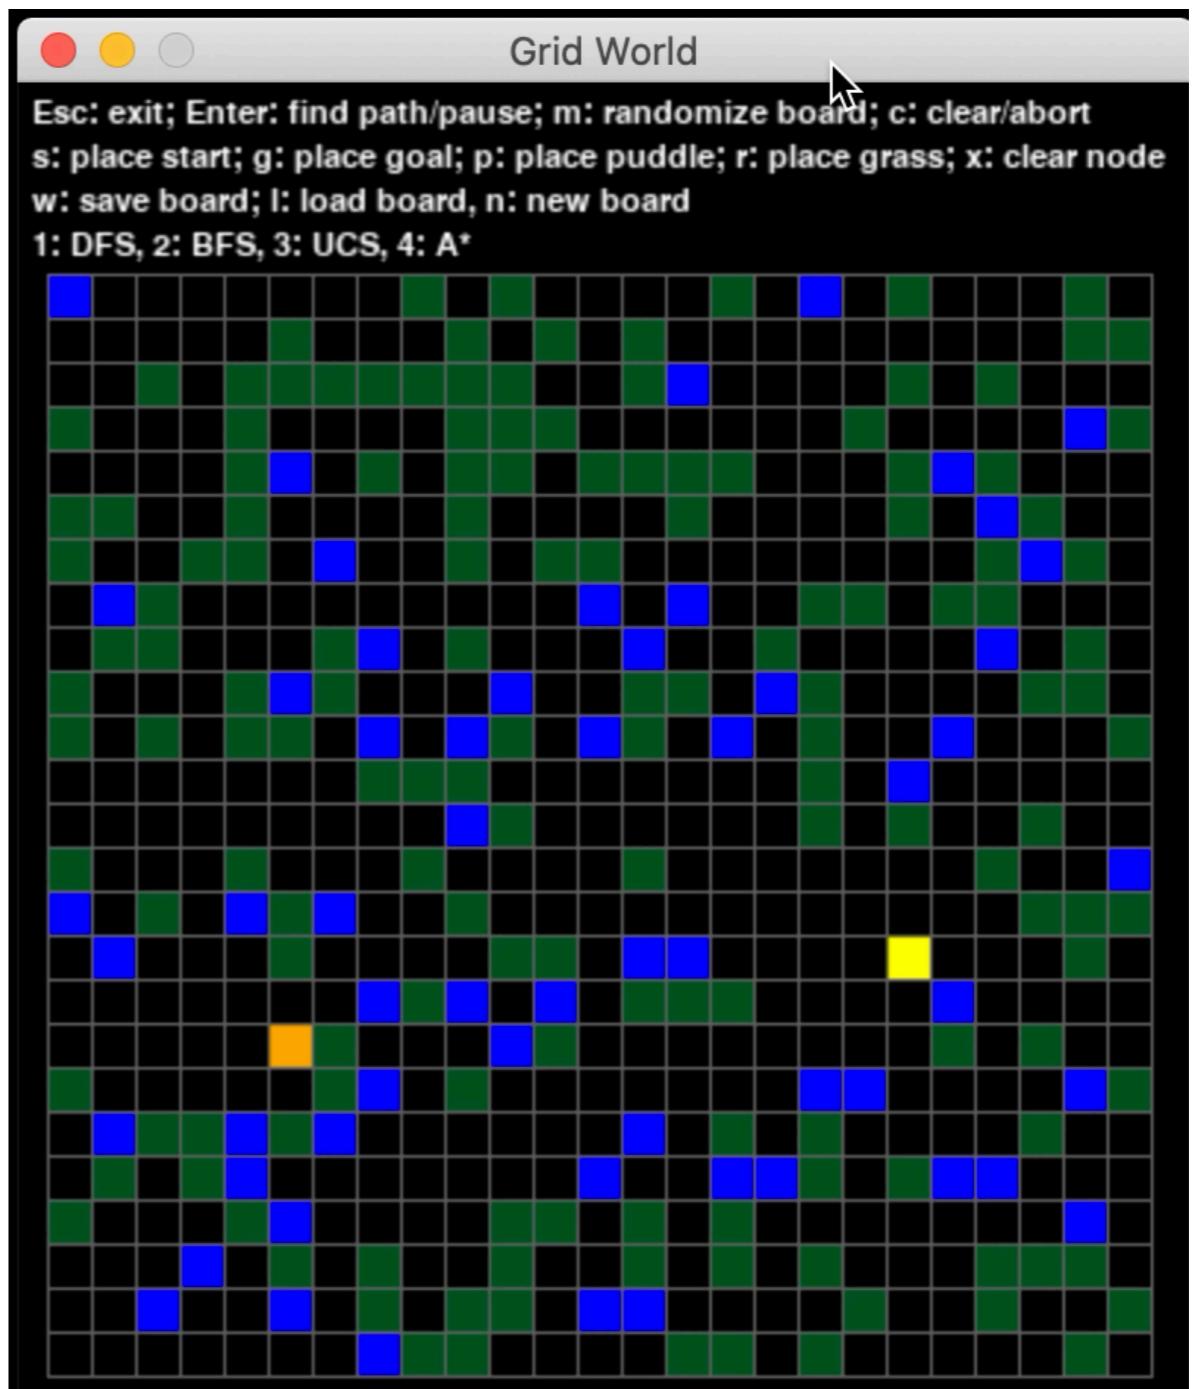
# Search in Graphs



# Search in Graphs

Danger: turning finite graphs into infinite trees

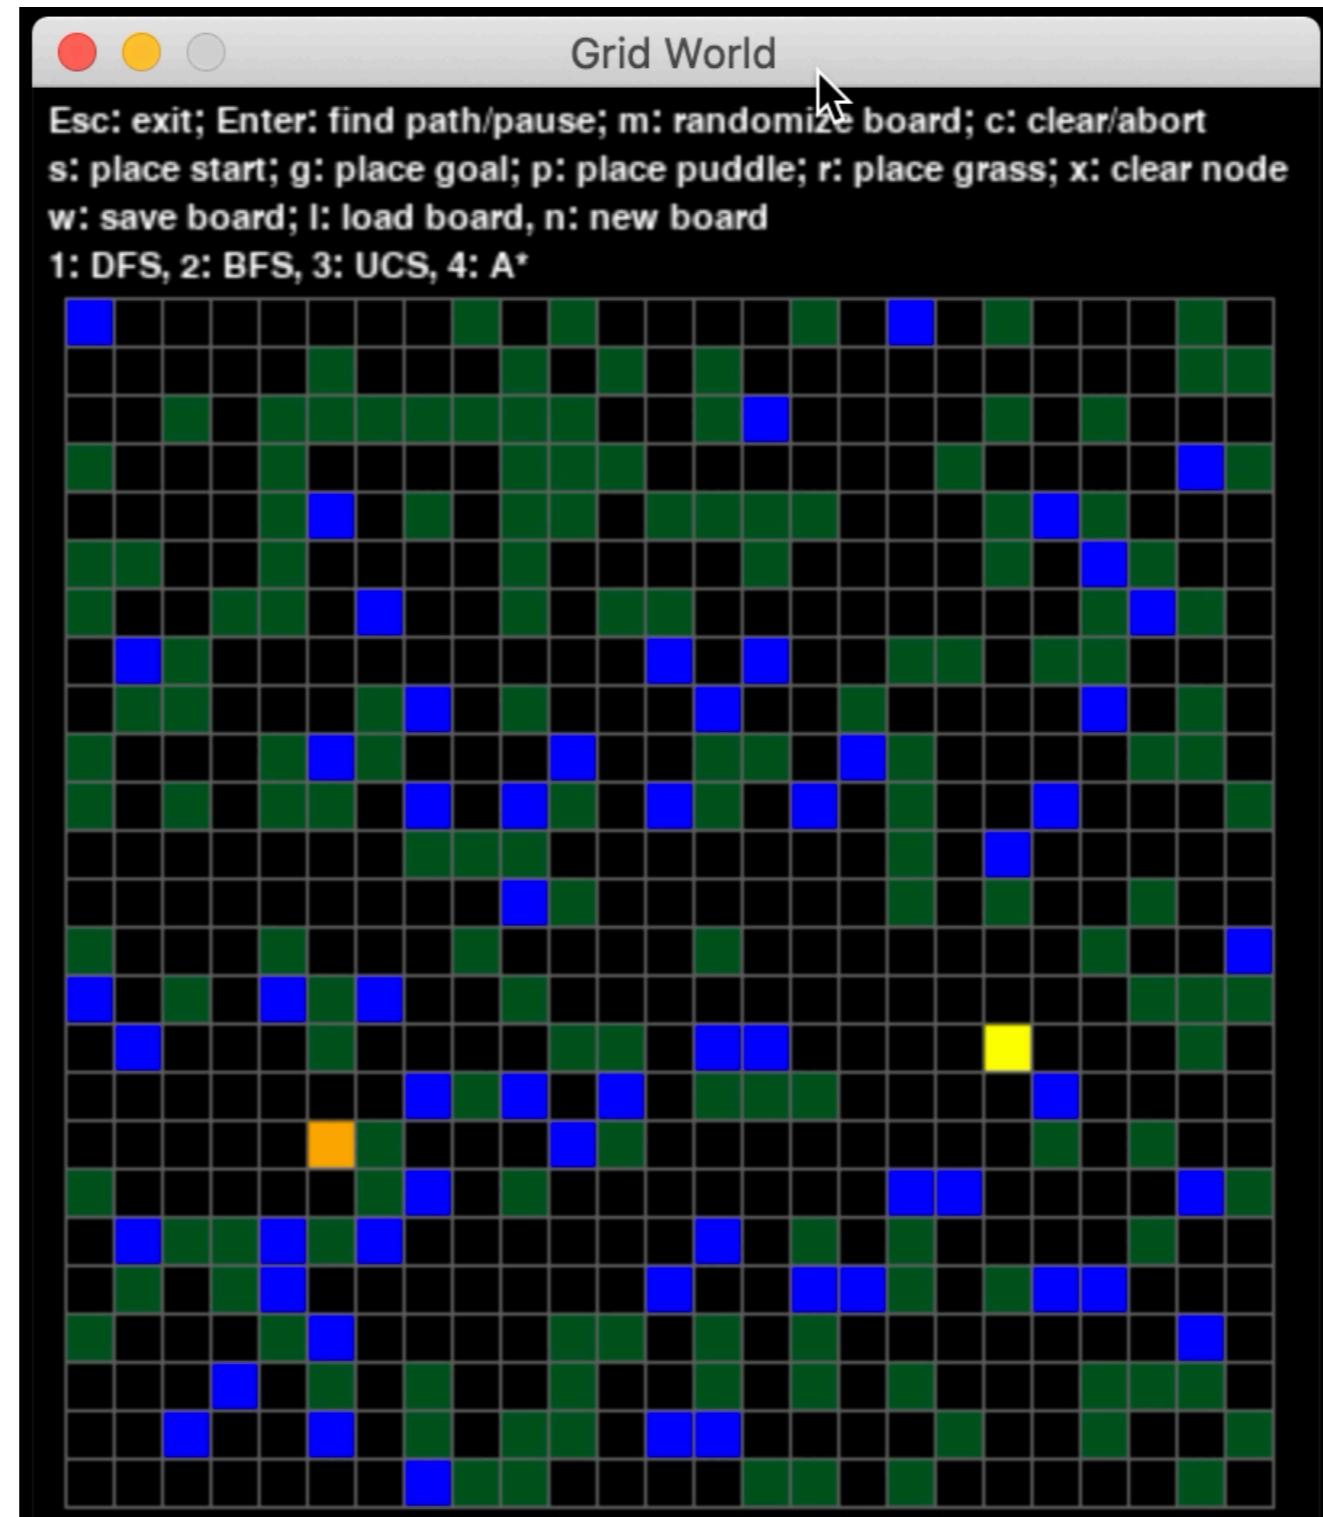
Make sure that you  
don't revisit nodes that  
are already explored



# DFS on graphs (Complete? Optimal?)

```
frontier = empty_stack.push(start)
explored = empty_stack
while frontier:
    current = frontier.pop()
    explored.push(current)
    for node in current.children():
        if node == goal:
            return path '''need bookkeeping'''
        if node in explored or node in frontier:
            continue
        else:
            frontier.push(node)
return failed
```

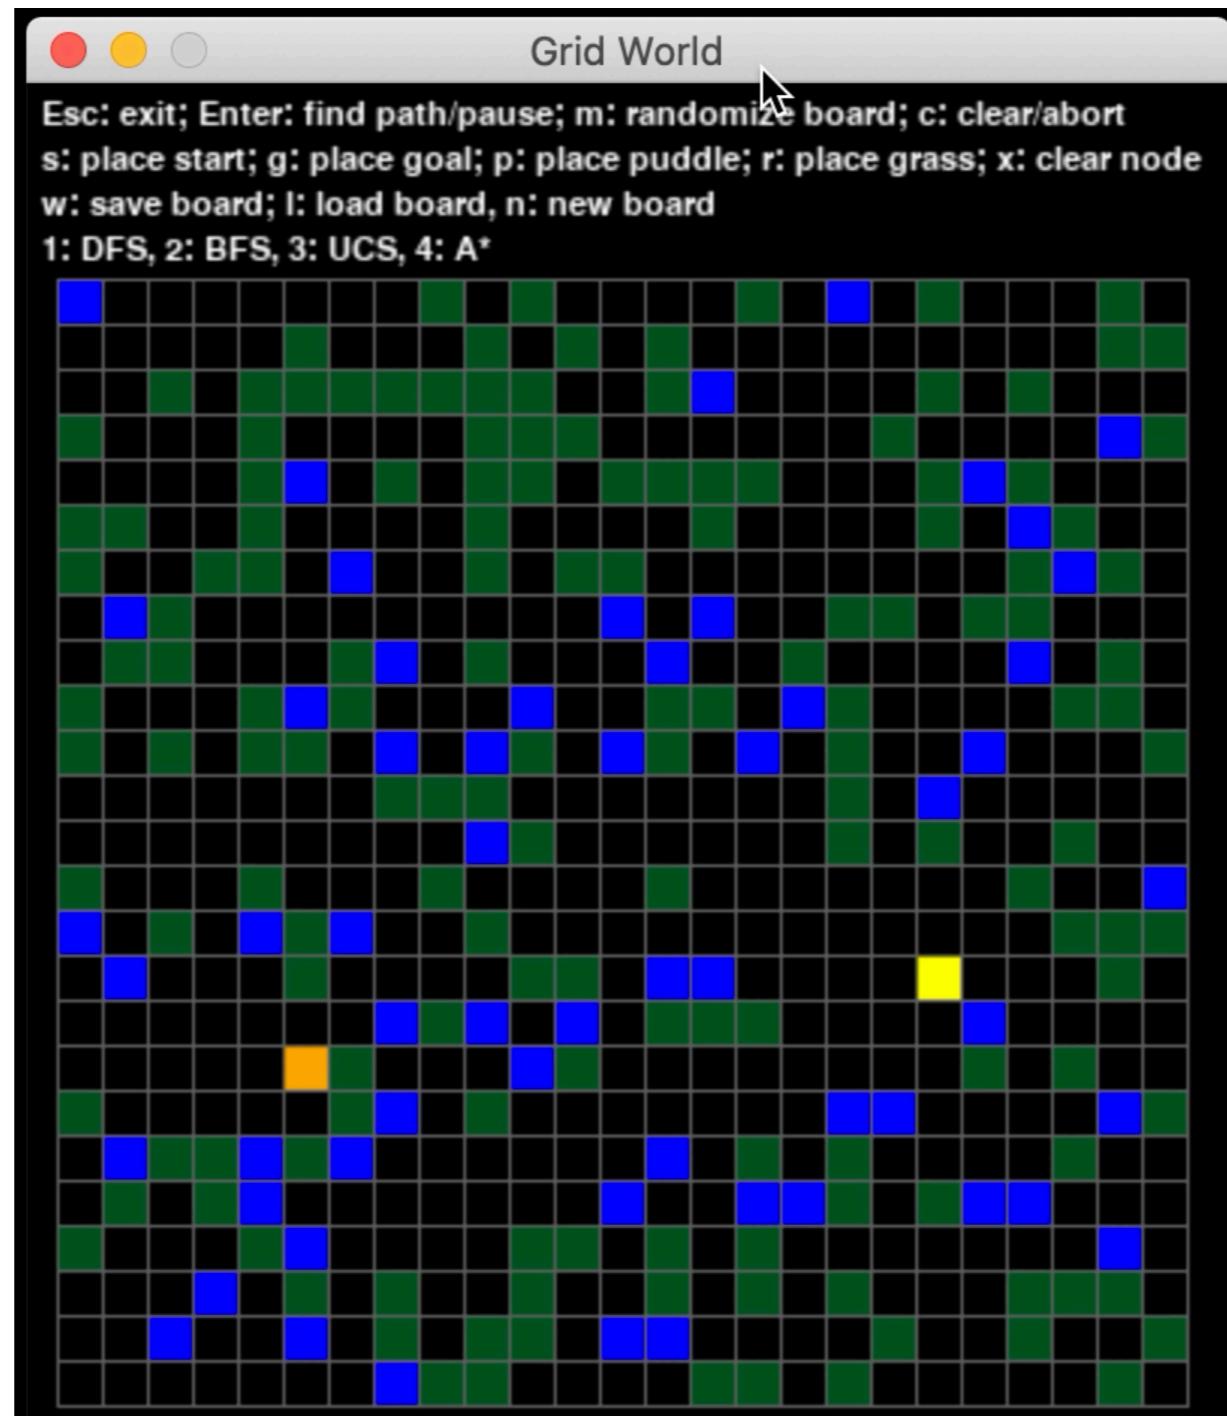
# DFS on graphs (Complete? Optimal?)



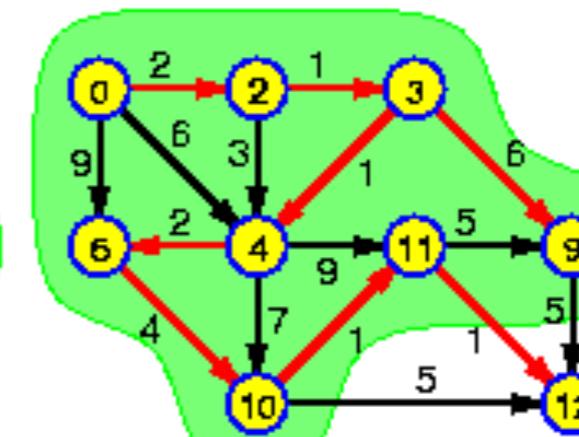
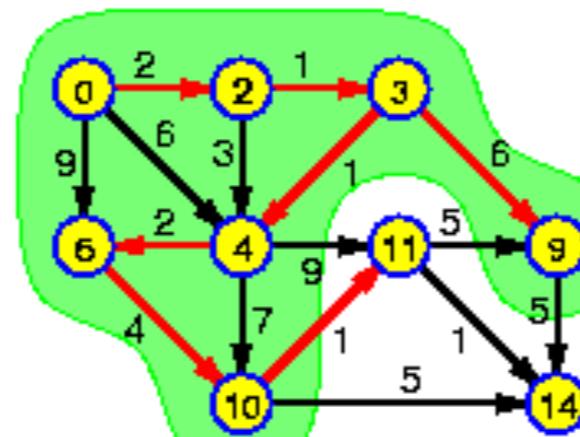
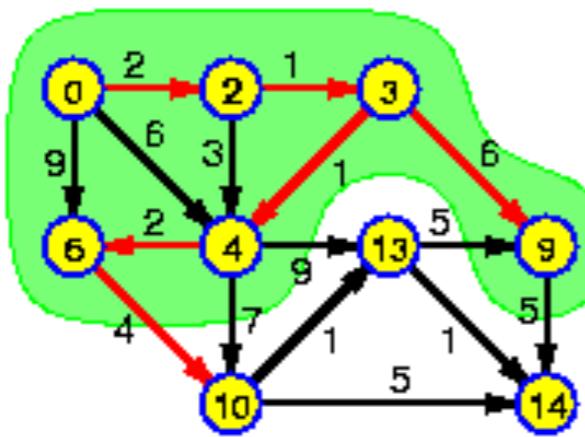
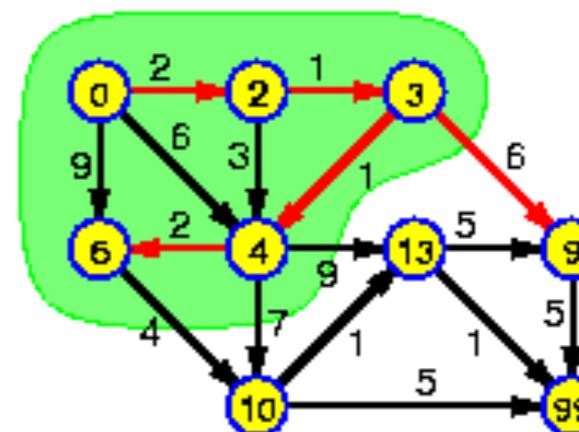
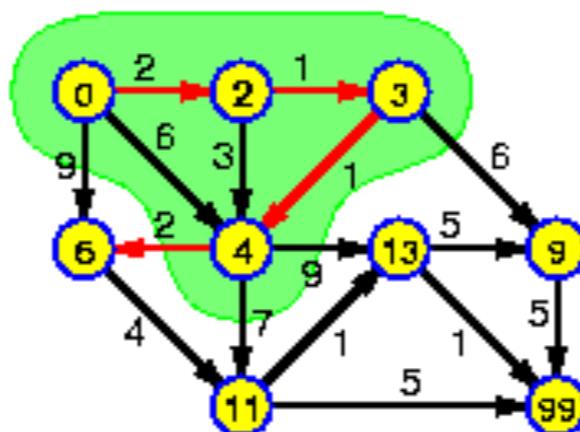
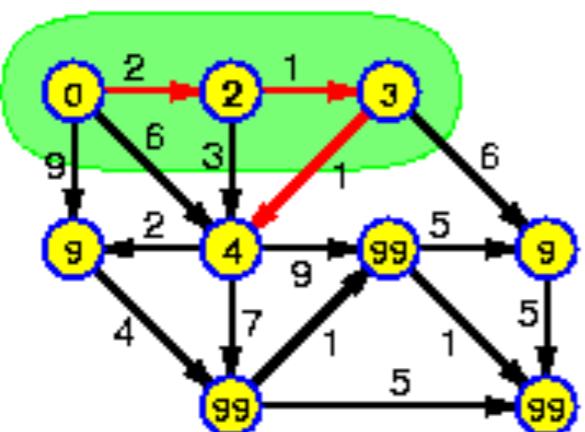
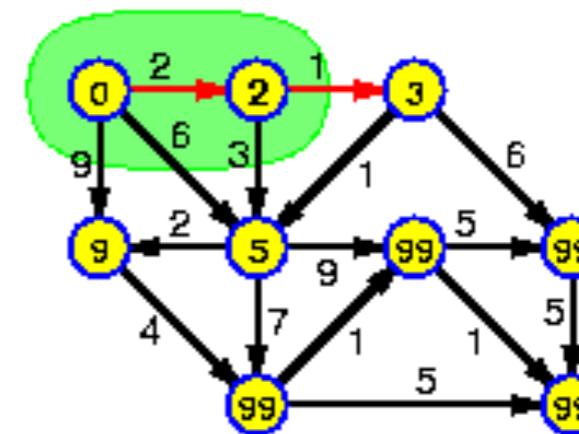
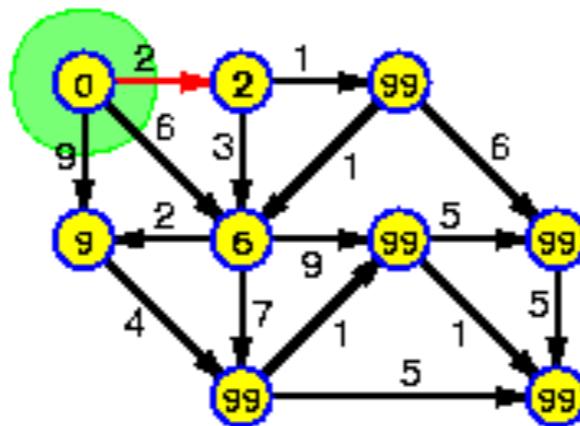
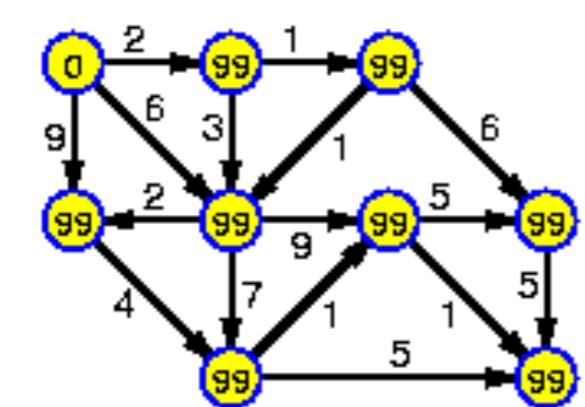
# BFS on graphs (Complete? Optimal?)

```
frontier = empty_queue.enqueue(start)
explored = empty_stack
while frontier:
    current = frontier.dequeue()
    explored.push(current)
    for node in current.children():
        if node == goal:
            return path '''need bookkeeping'''
        if node in explored or node in frontier:
            continue
        else:
            frontier.enqueue(node)
return failed
```

# BFS on graphs (Complete? Optimal?)



# Dijkstra/Uniform Cost Search



# Dijkstra/Uniform Cost Search

**function** UNIFORM-COST-SEARCH(*problem*) **returns** a solution, or failure

*node*  $\leftarrow$  a node with STATE = *problem*.INITIAL-STATE, PATH-COST = 0

*frontier*  $\leftarrow$  a priority queue ordered by PATH-COST, with *node* as the only element

*explored*  $\leftarrow$  an empty set

**loop do**

↑      ↑      ↓

**if** EMPTY?(*frontier*) **then return** failure

*node*  $\leftarrow$  POP(*frontier*) /\* chooses the lowest-cost node in *frontier* \*/

**if** *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)

    add *node*.STATE to *explored*

**for each** *action* **in** *problem*.ACTIONS(*node*.STATE) **do**

*child*  $\leftarrow$  CHILD-NODE(*problem*, *node*, *action*)

**if** *child*.STATE is not in *explored* or *frontier* **then**

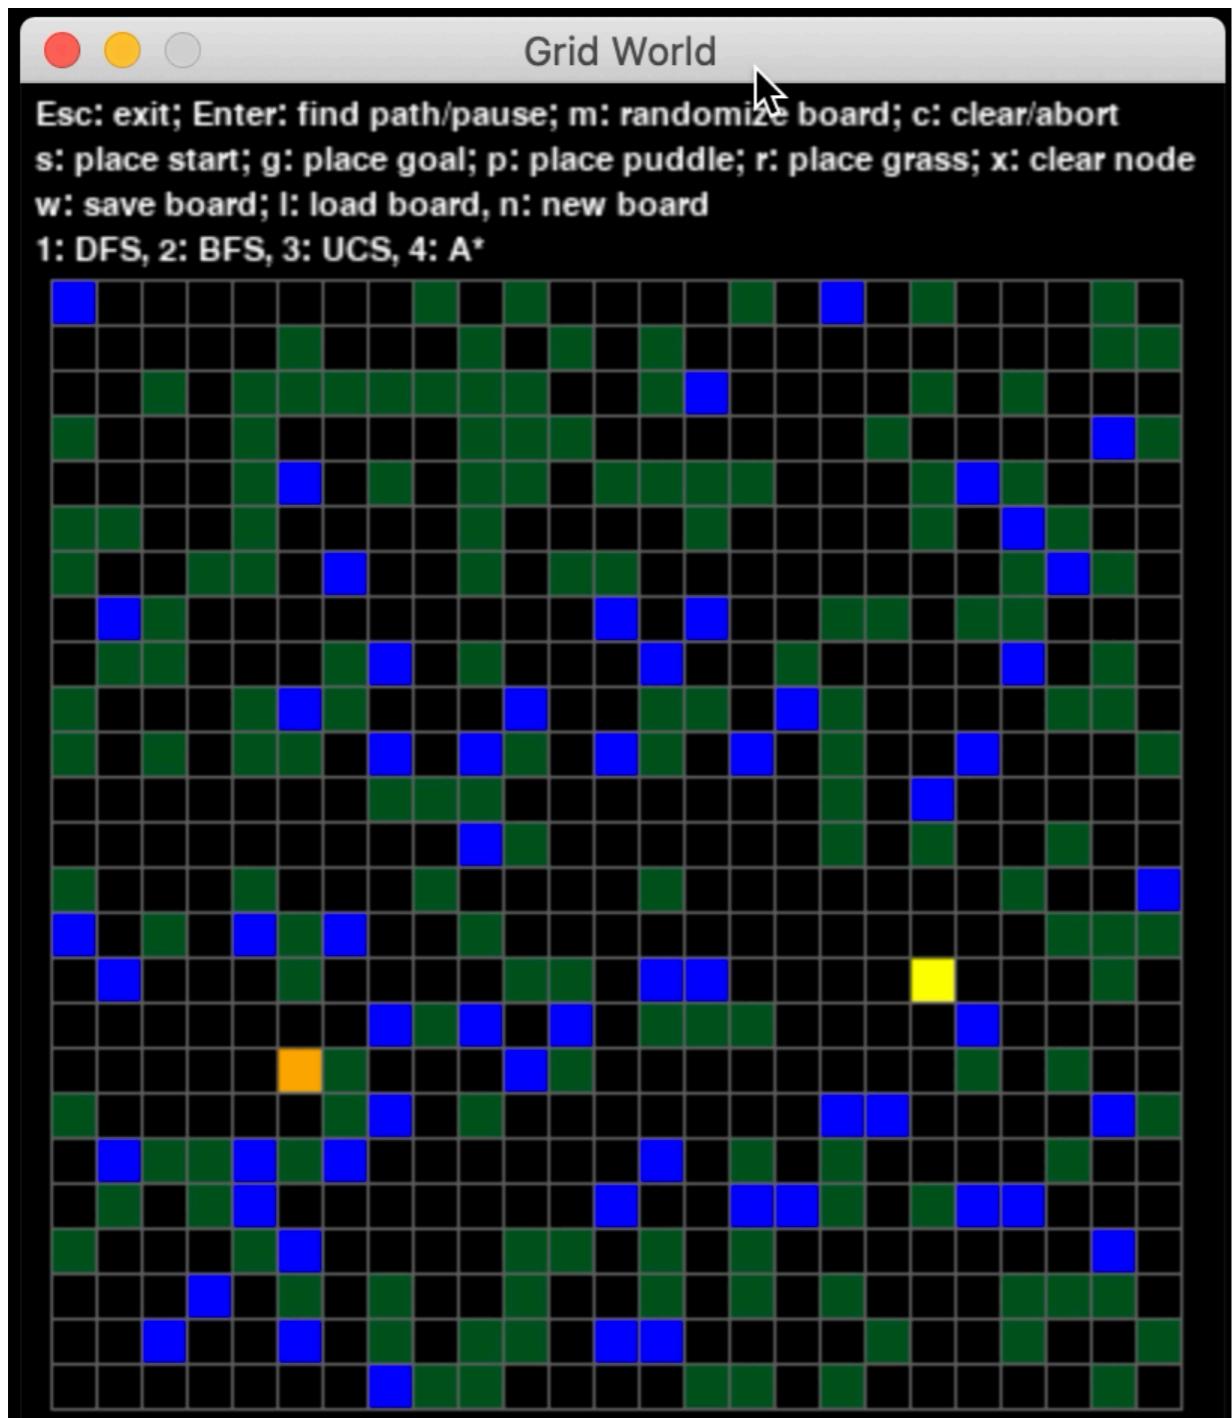
*frontier*  $\leftarrow$  INSERT(*child*, *frontier*)

**else if** *child*.STATE is in *frontier* with higher PATH-COST **then**

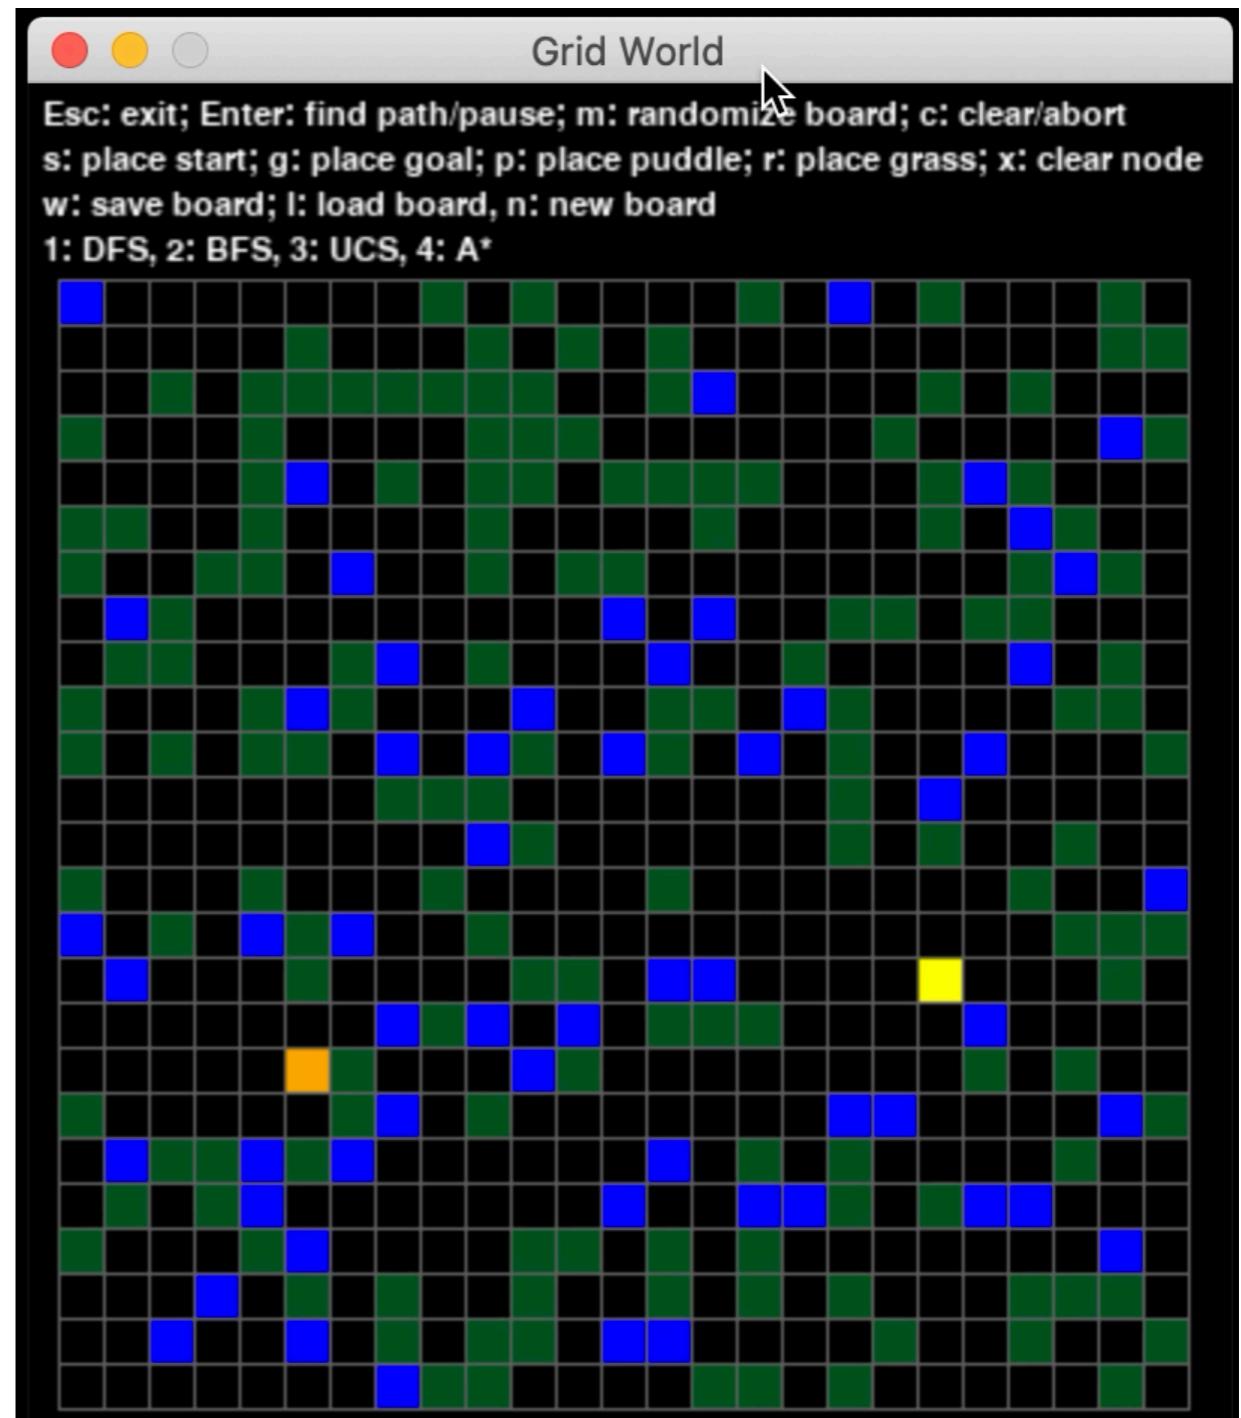
            replace that *frontier* node with *child*

pop a frontier node, maintain explored

push children, maintain frontier



BFS



UCS

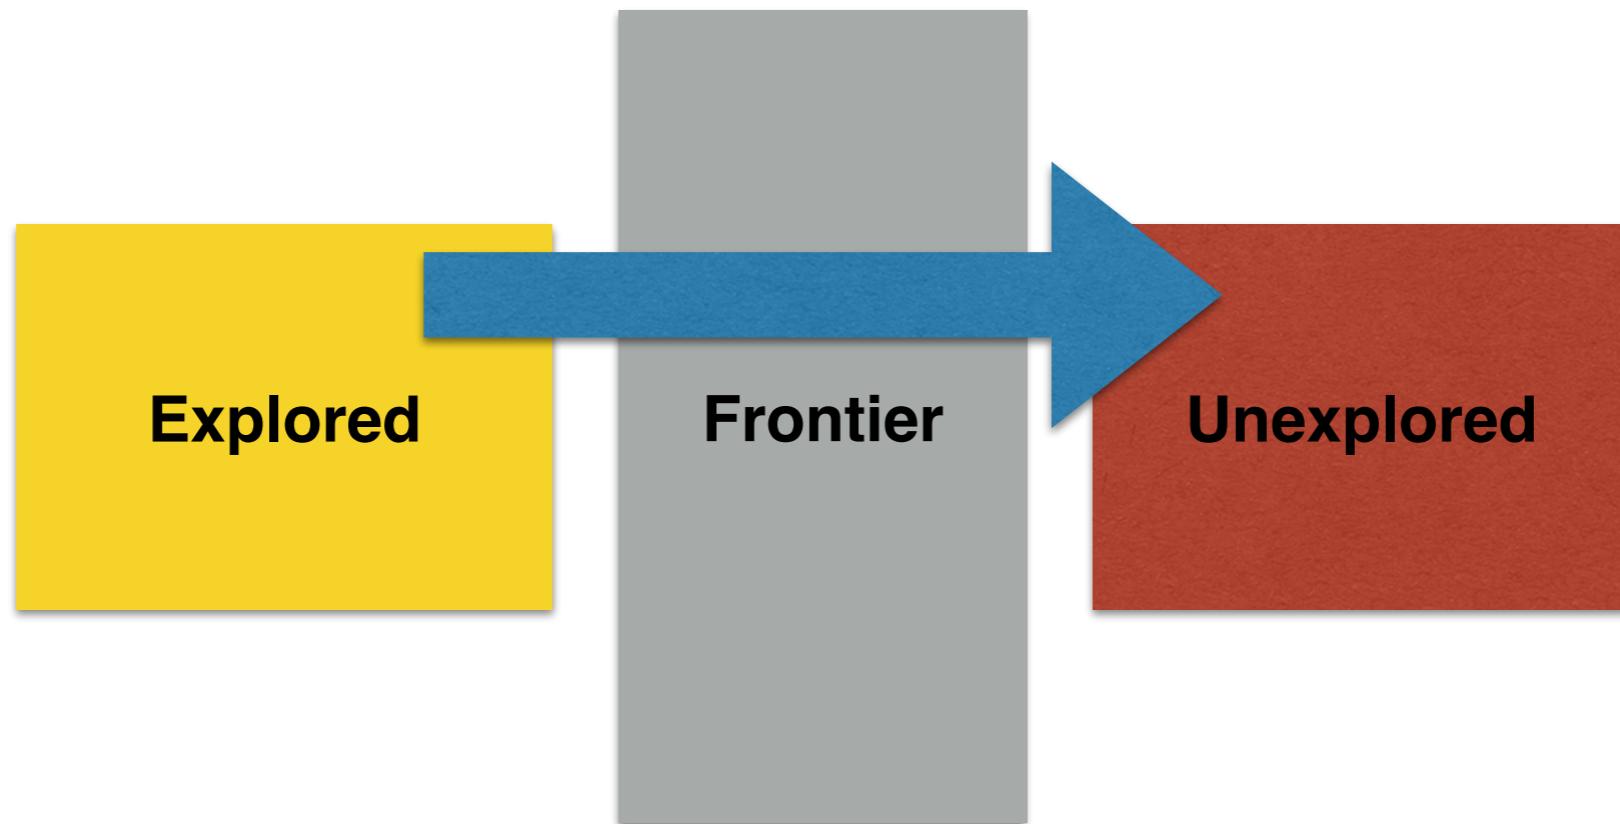
# Analysis of UCS: Optimality

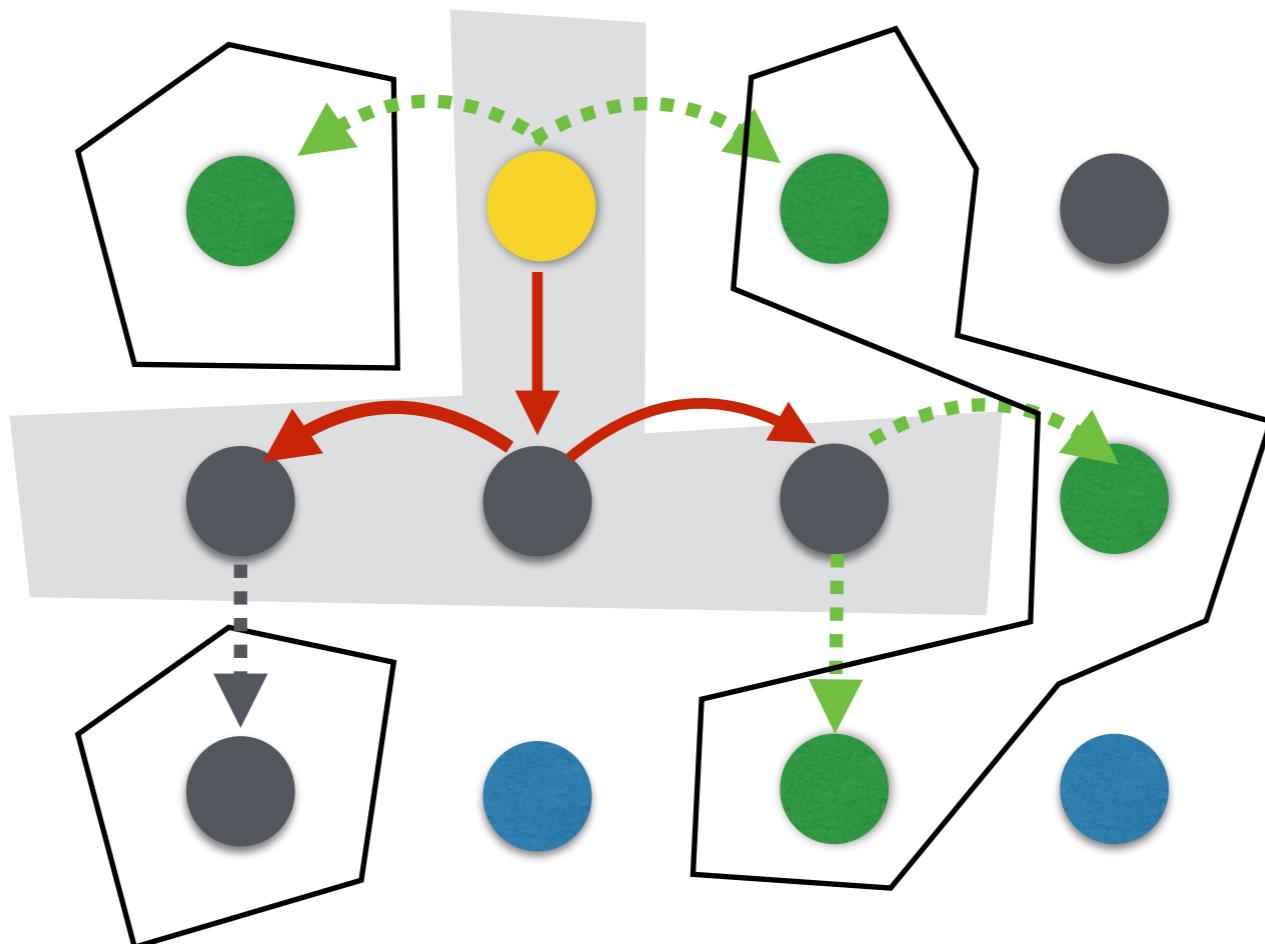
- When any node is expanded, UCS has found the optimal path to it (i.e., with minimum cost).

How to prove this?

# Analysis of UCS: Separation

- Lemma: At the end of each iteration of UCS, every acyclic path from any explored node to any unexplored node passes through at least one node in the frontier.

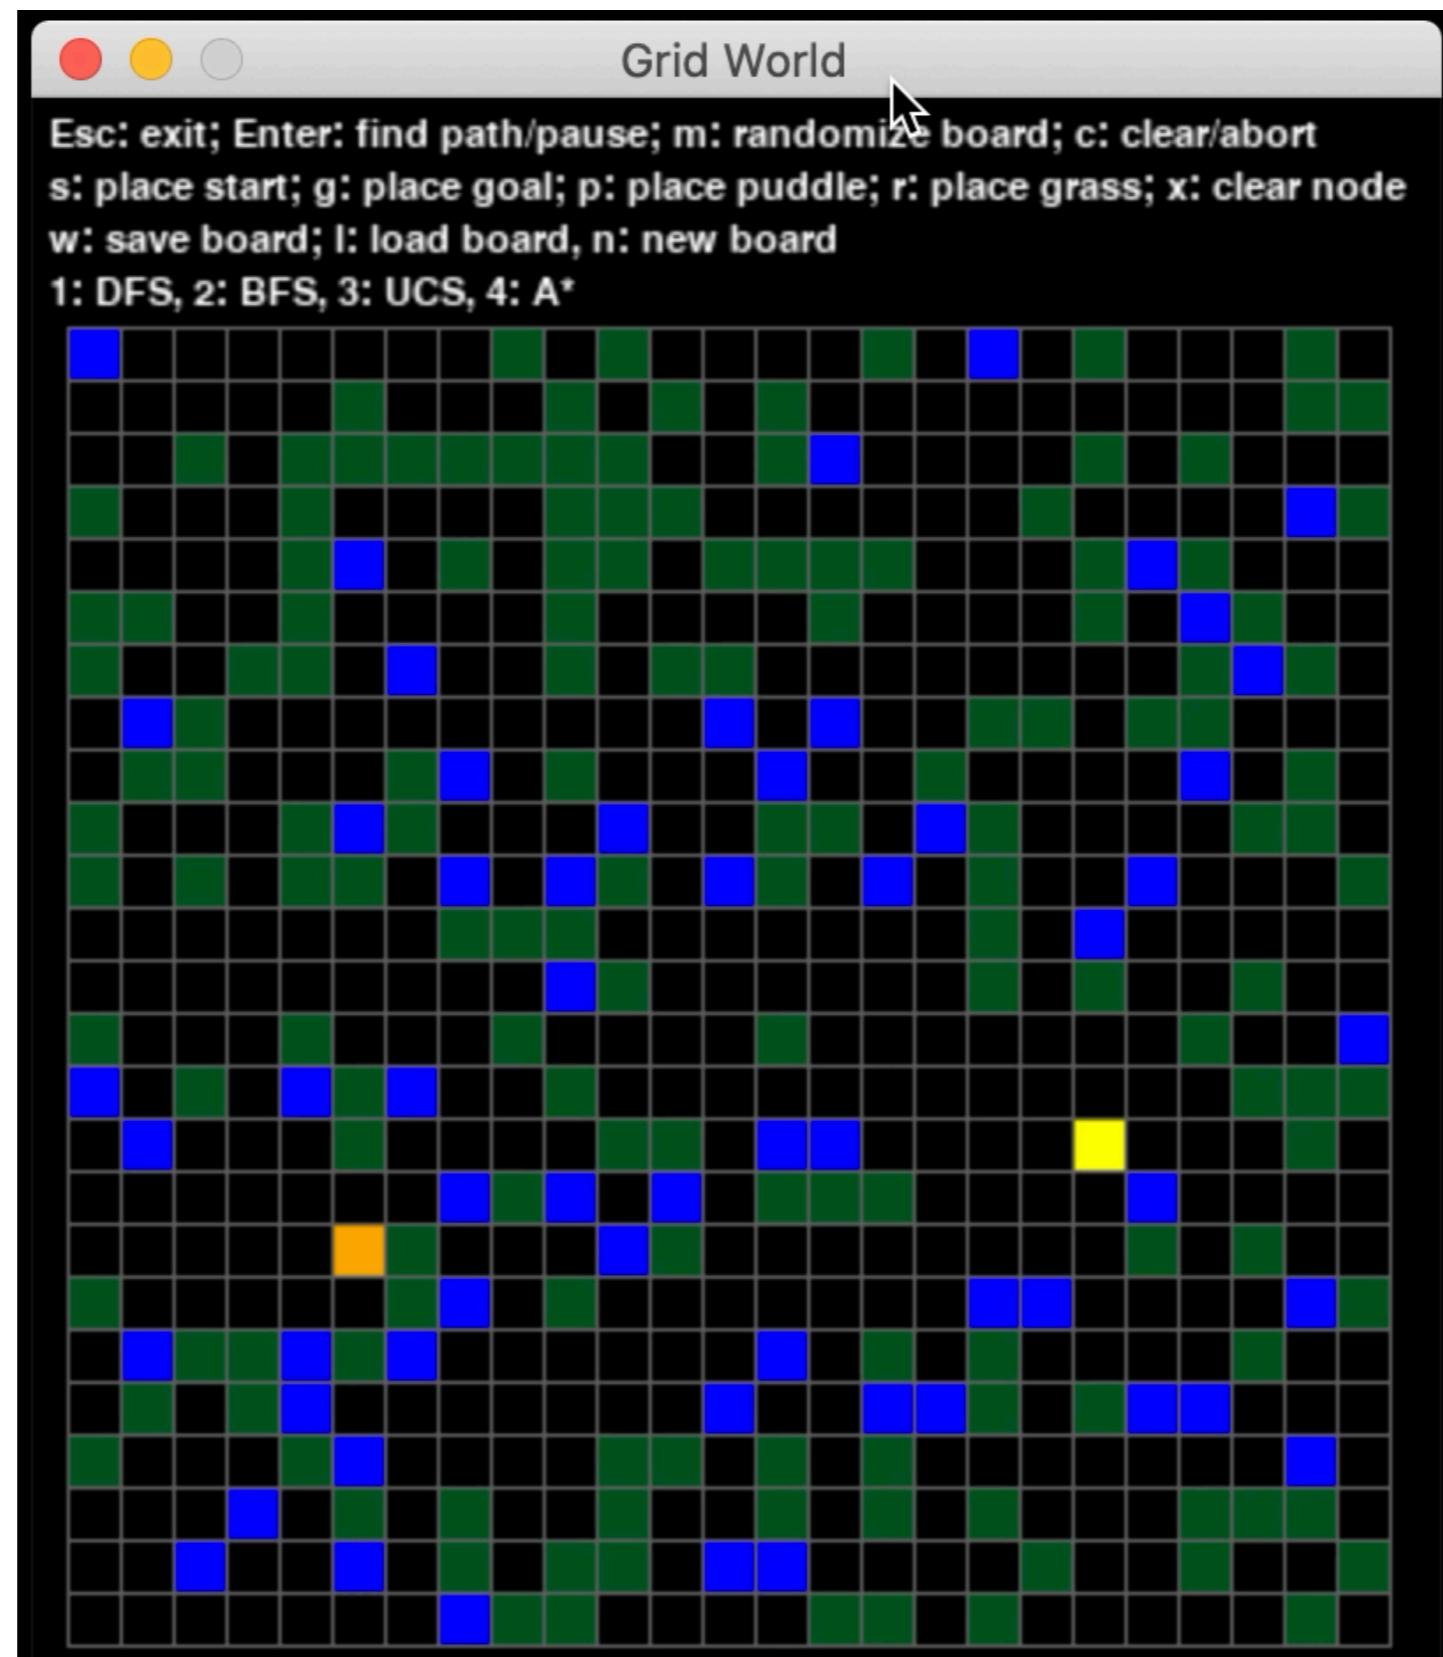




$+5$

$+\infty$

$+1$



# Analysis of UCS: Separation

- Lemma: At the end of each iteration of UCS, every acyclic path from any explored node to any unexplored node passes through at least one node in the frontier.
- How to prove this?
  - Induction over the steps of the algorithm.

# Analysis of UCS: Optimality

- When any node is expanded, UCS has found the optimal path to it (i.e., with minimum cost).

Proof: Suppose when node n is expanded, UCS has not found the optimal path to it. Then there exists another path to n that has lower cost. But by the Lemma, that other path must pass through at least one node in the current frontier. So…

# Properties of UCS

- Completeness: Yes, except for pathological cases (actions that do nothing and no cost)
- Optimality: uniform-cost search expands nodes in order of their optimal path cost.
- Complexity: UCS does not care about number of steps, only cost! So it may need to explore more than the depth of the goal.

# Uninformed vs. Heuristic

DFS, BFS, UCS are all uninformed search:  
you don't know where you are unless you  
happen to hit the goal.

Uninformed search is not the  
way to go for any slightly  
interesting problem.

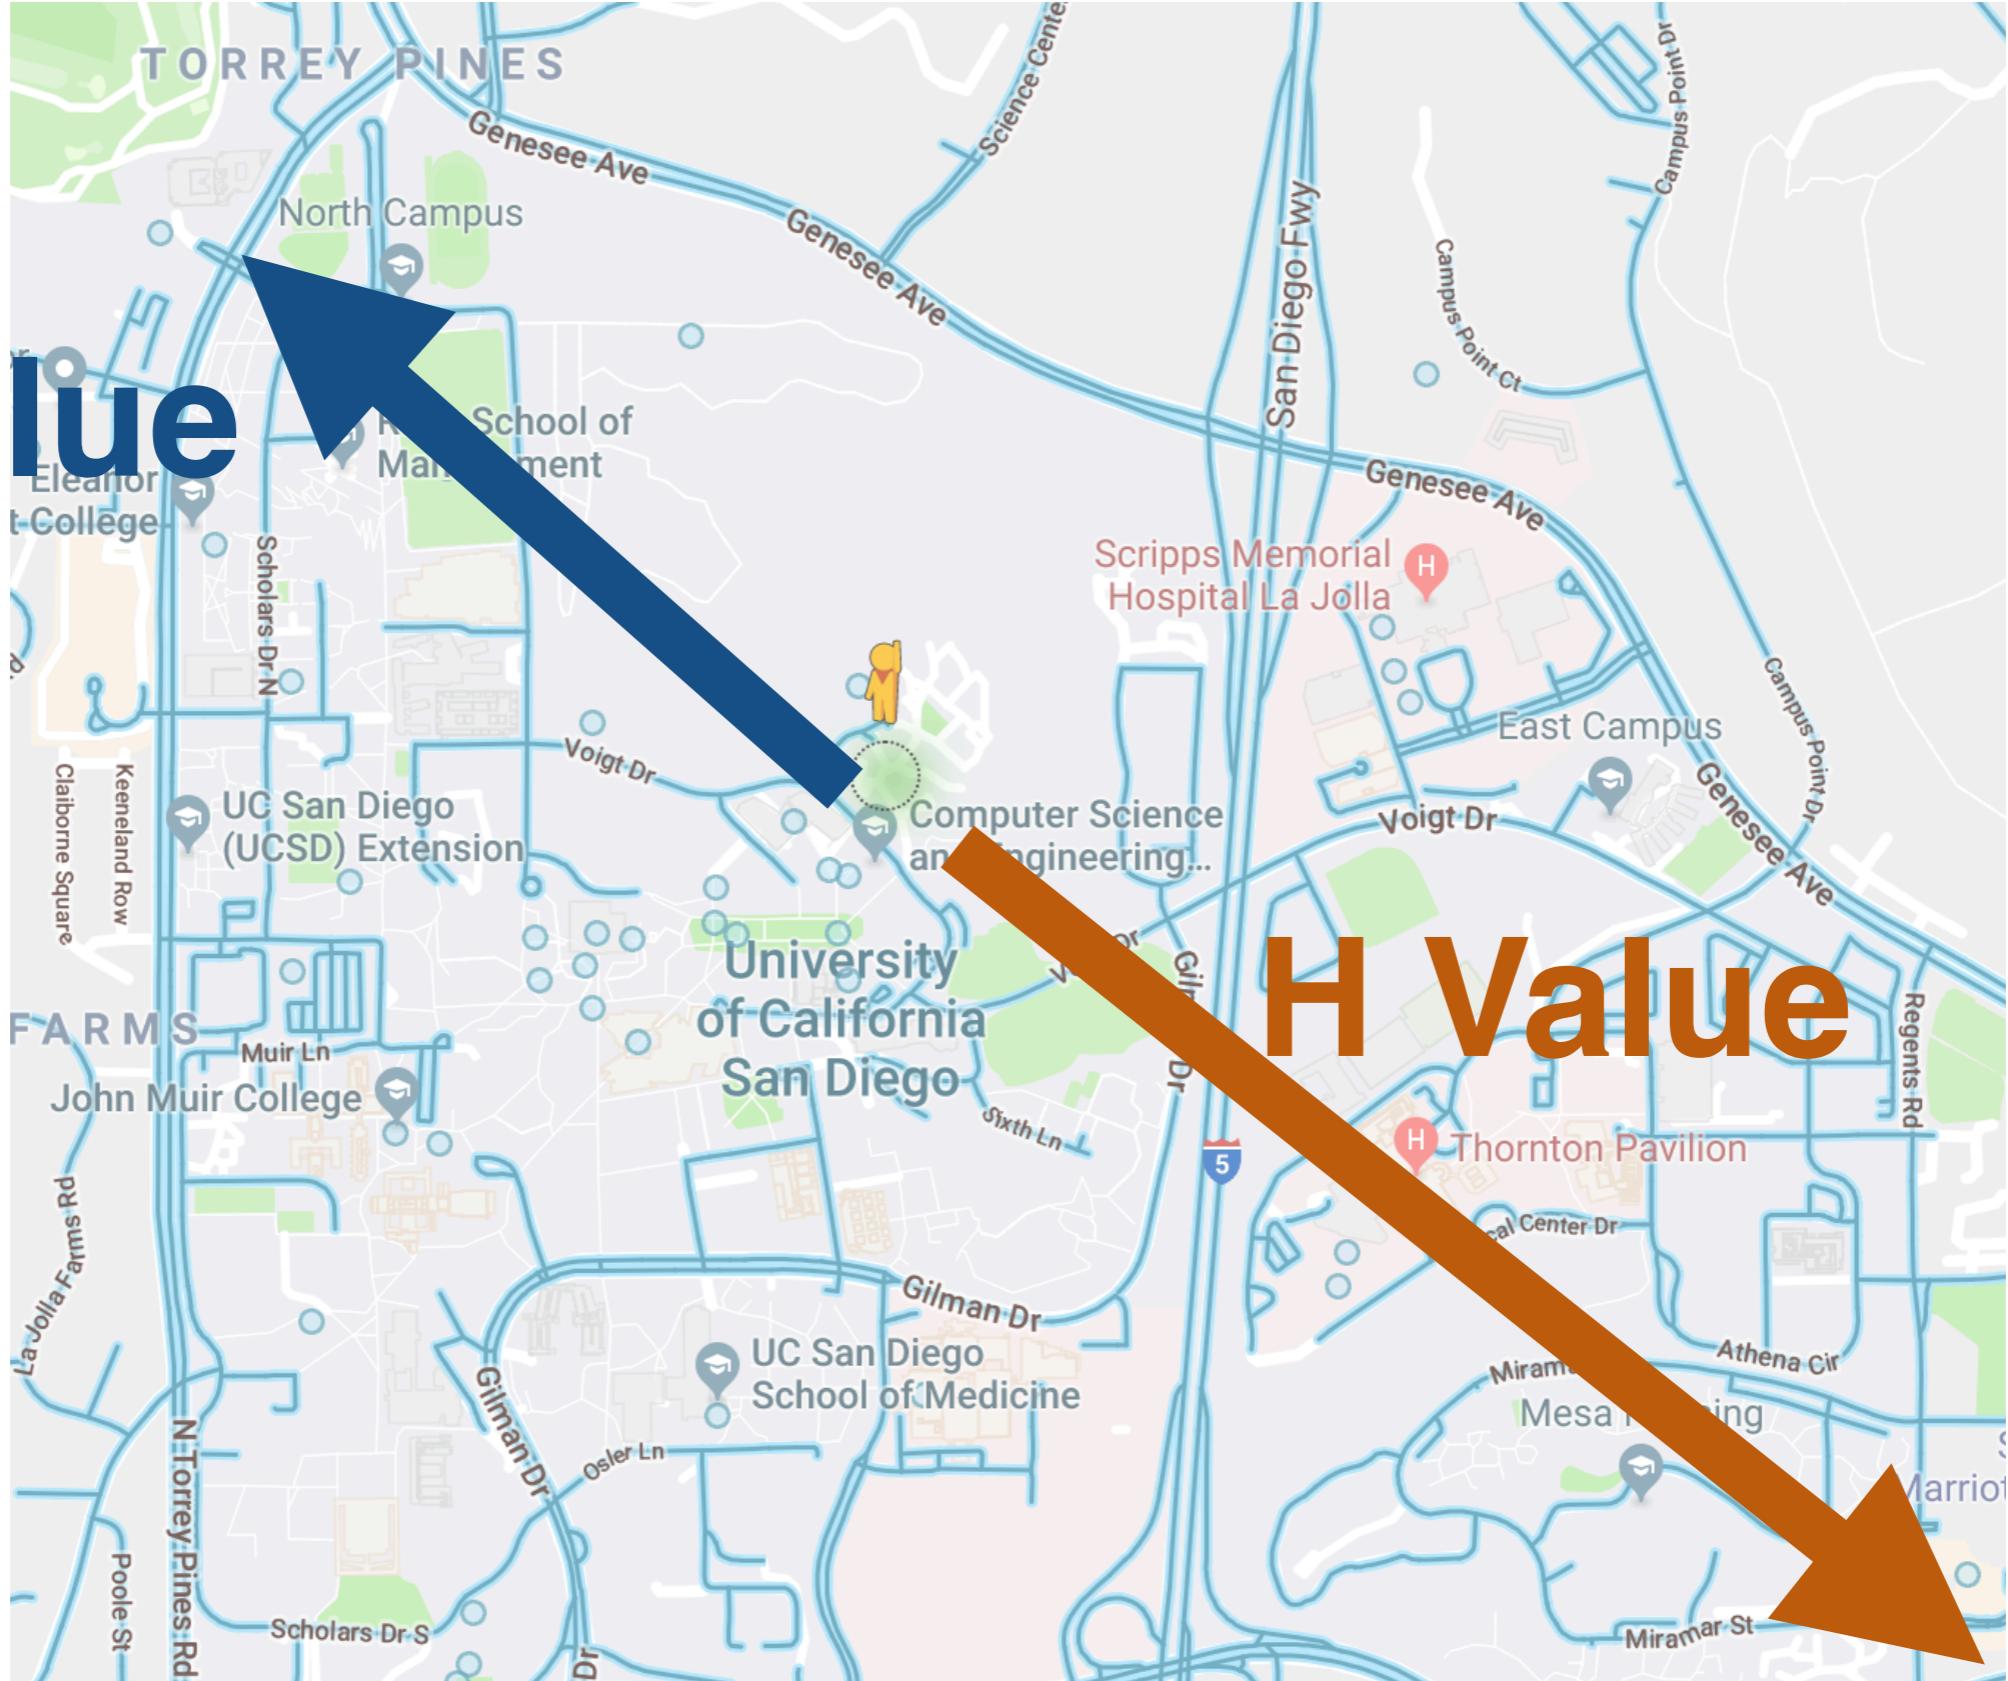
# Informed Search: A\*

- What if you can measure how close you are to the goal?
  - This measure is always nonnegative and is zero at the goal.
  - Euclidean distance. Manhattan distance.
- Do exactly like UCS, but use a new cost:

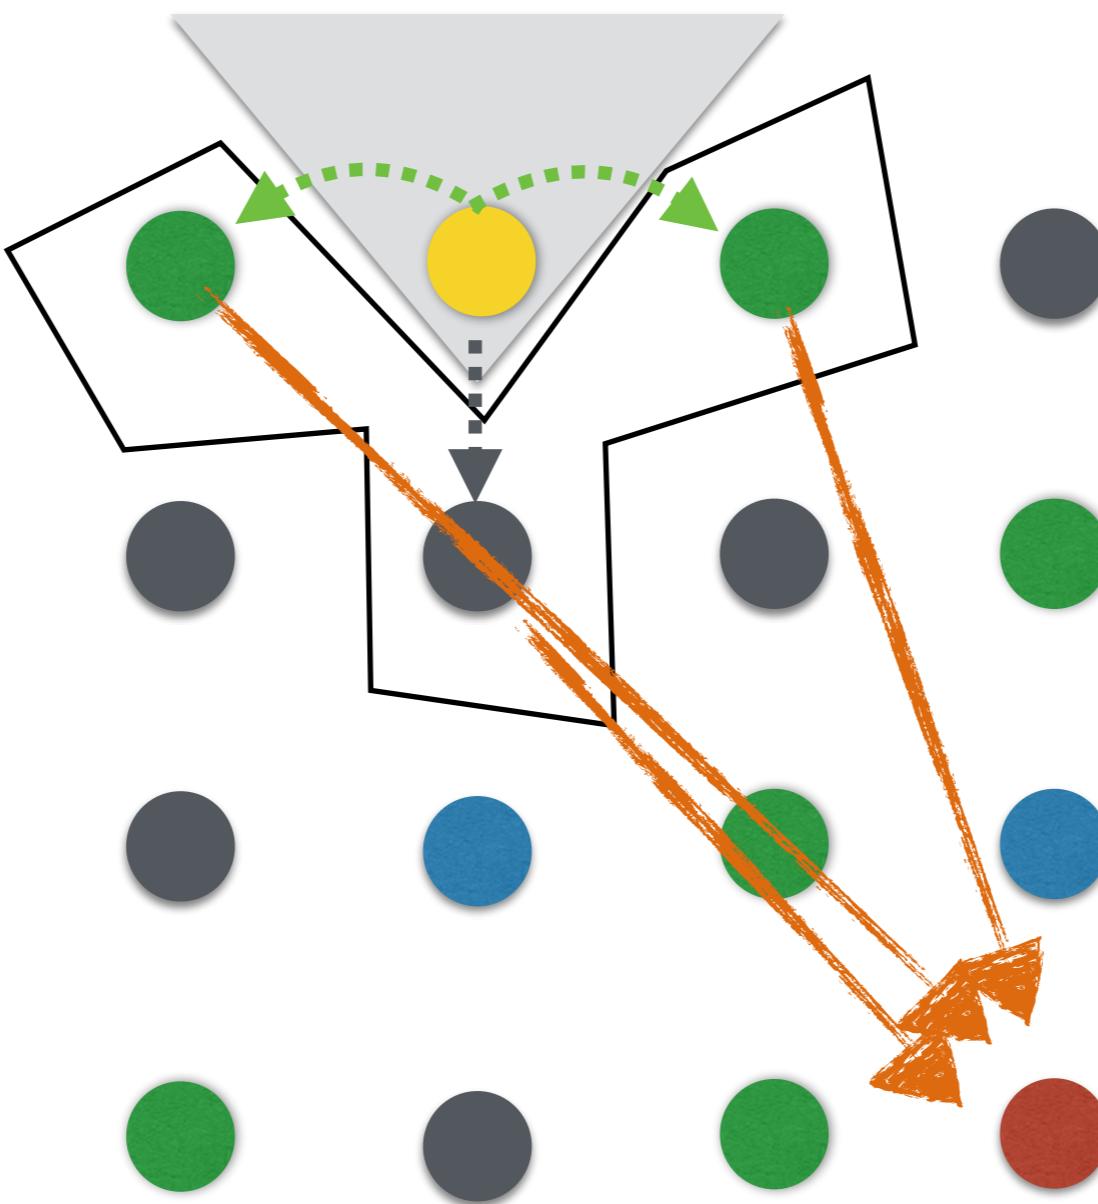
$$F = G + H$$

Real cost + Heuristic cost

# G Value



# H Value



+5



+inf

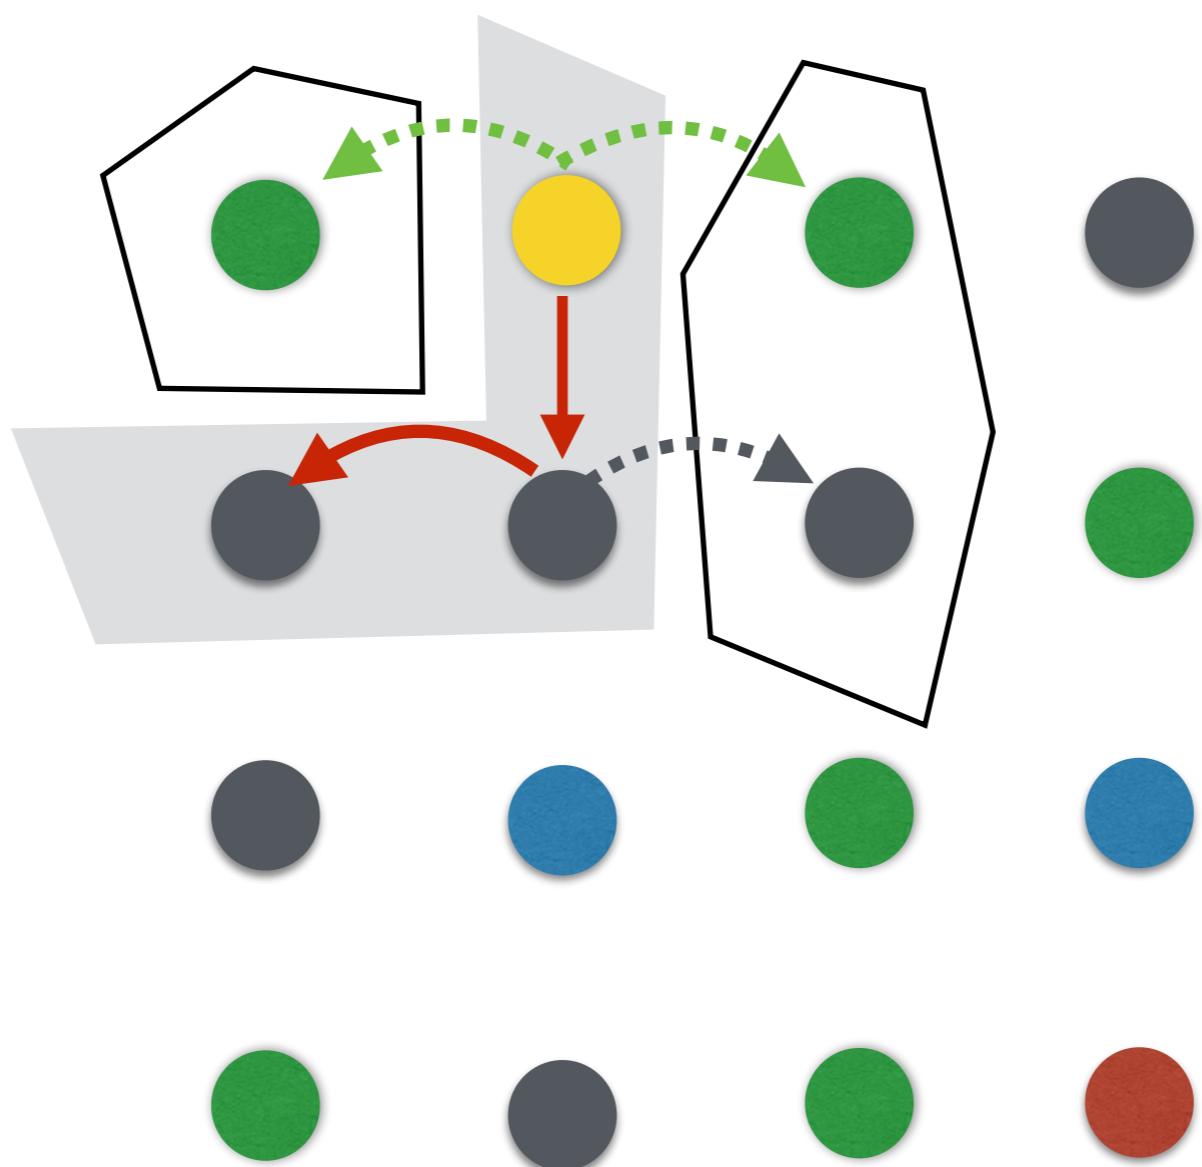


+1

# A\* is Almost Just UCS

```
function UNIFORM-COST-SEARCH(problem) returns a solution, or failure
  node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0 G+H
  frontier  $\leftarrow$  a priority queue ordered by PATH COST, with node as the only element
  explored  $\leftarrow$  an empty set G+H
  loop do
    if EMPTY?(frontier) then return failure
    node  $\leftarrow$  POP(frontier) /* chooses the lowest-cost node in frontier */
    if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        frontier  $\leftarrow$  INSERT(child, frontier) G+H
      else if child.STATE is in frontier with higher PATH COST then
        replace that frontier node with child
```

In UCS

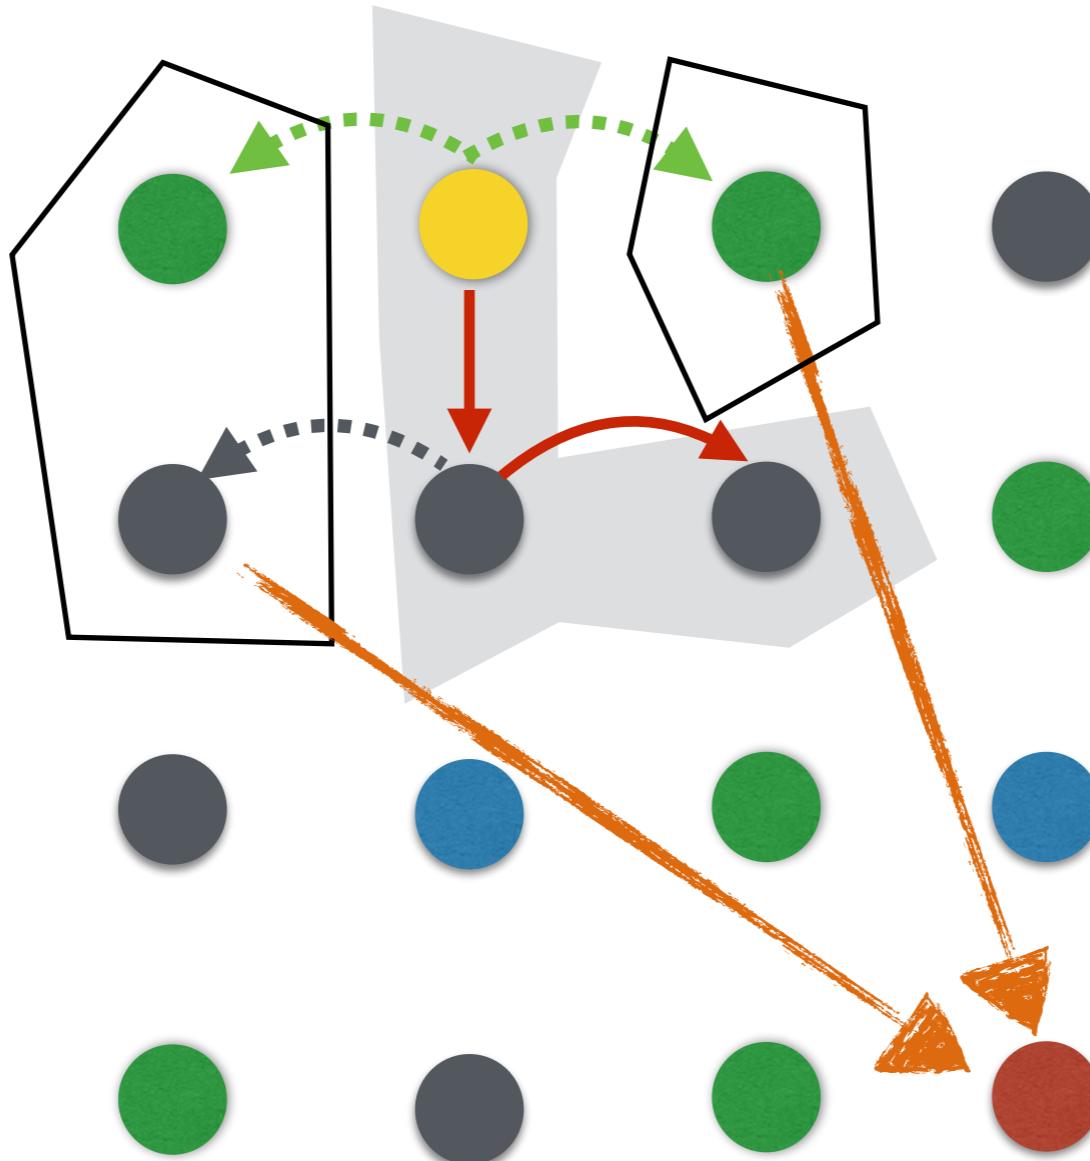


+5

+inf

+1

In A\*

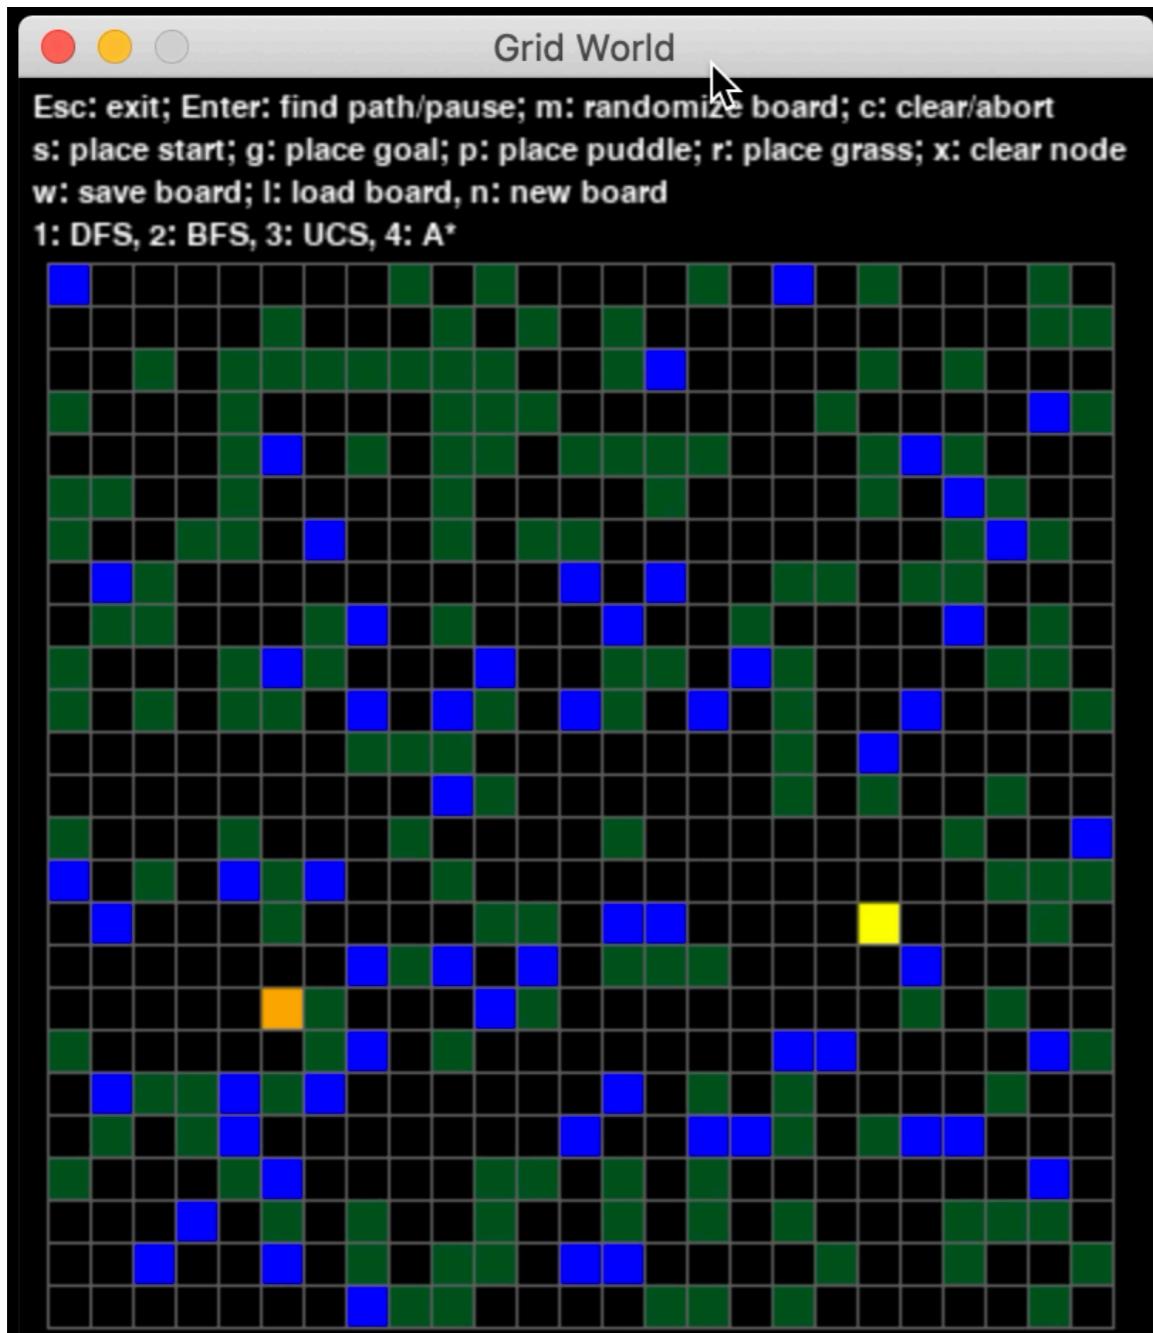


Does H affect optimality? How?

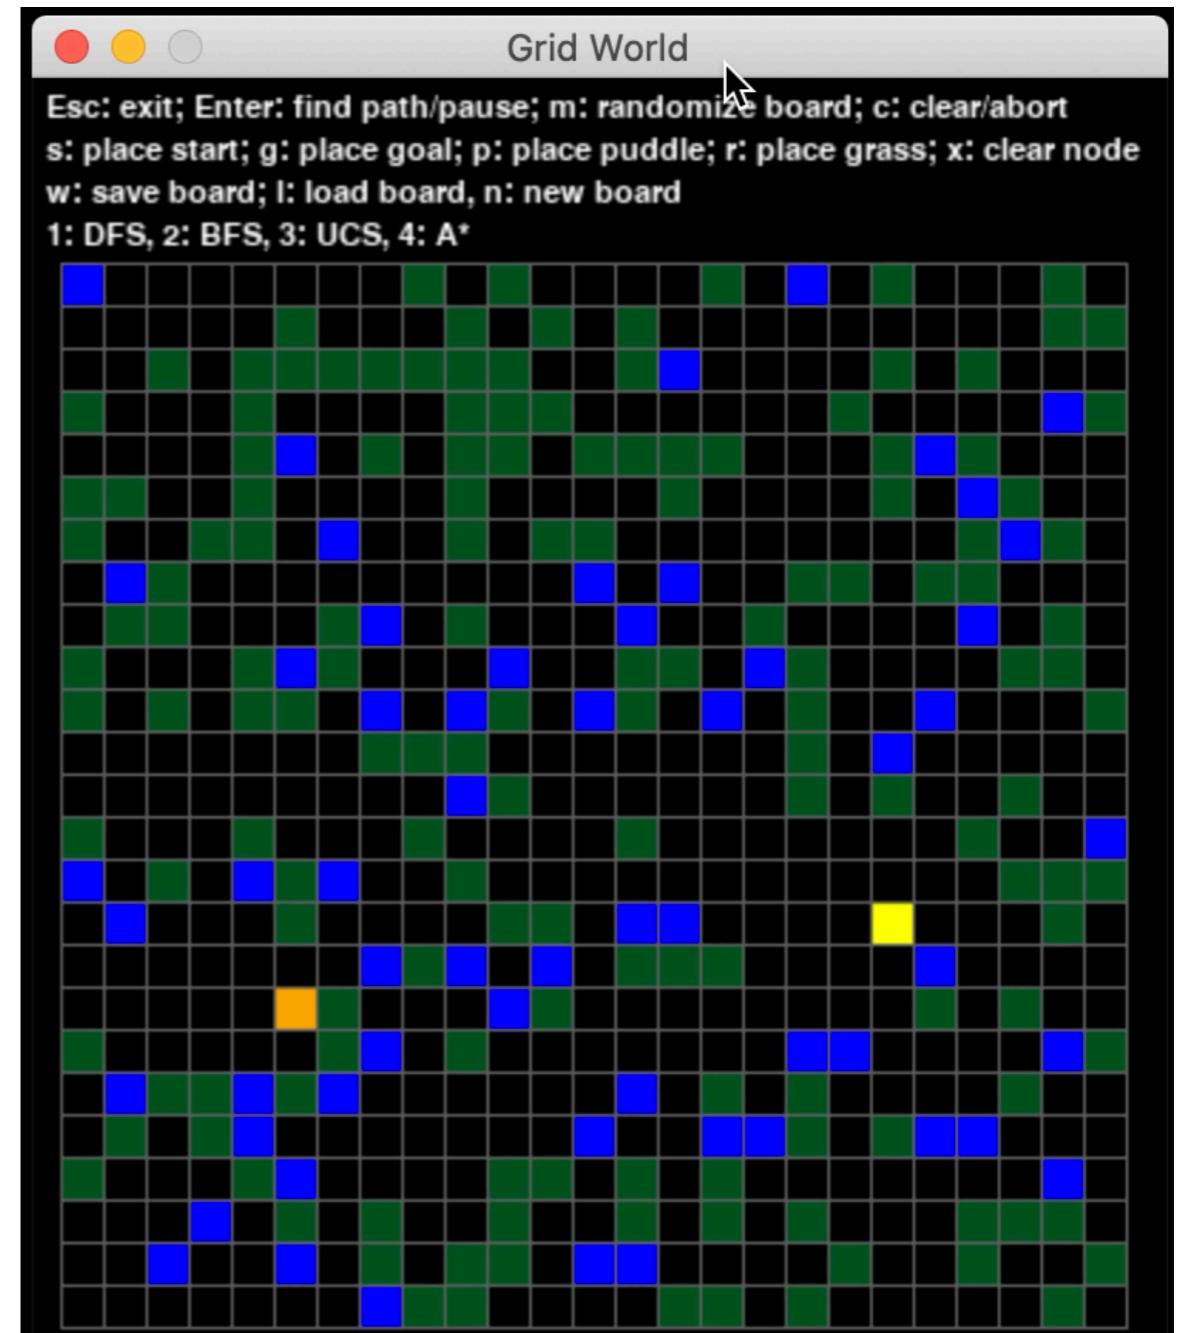
+5

+inf

+1



UCS



A \*

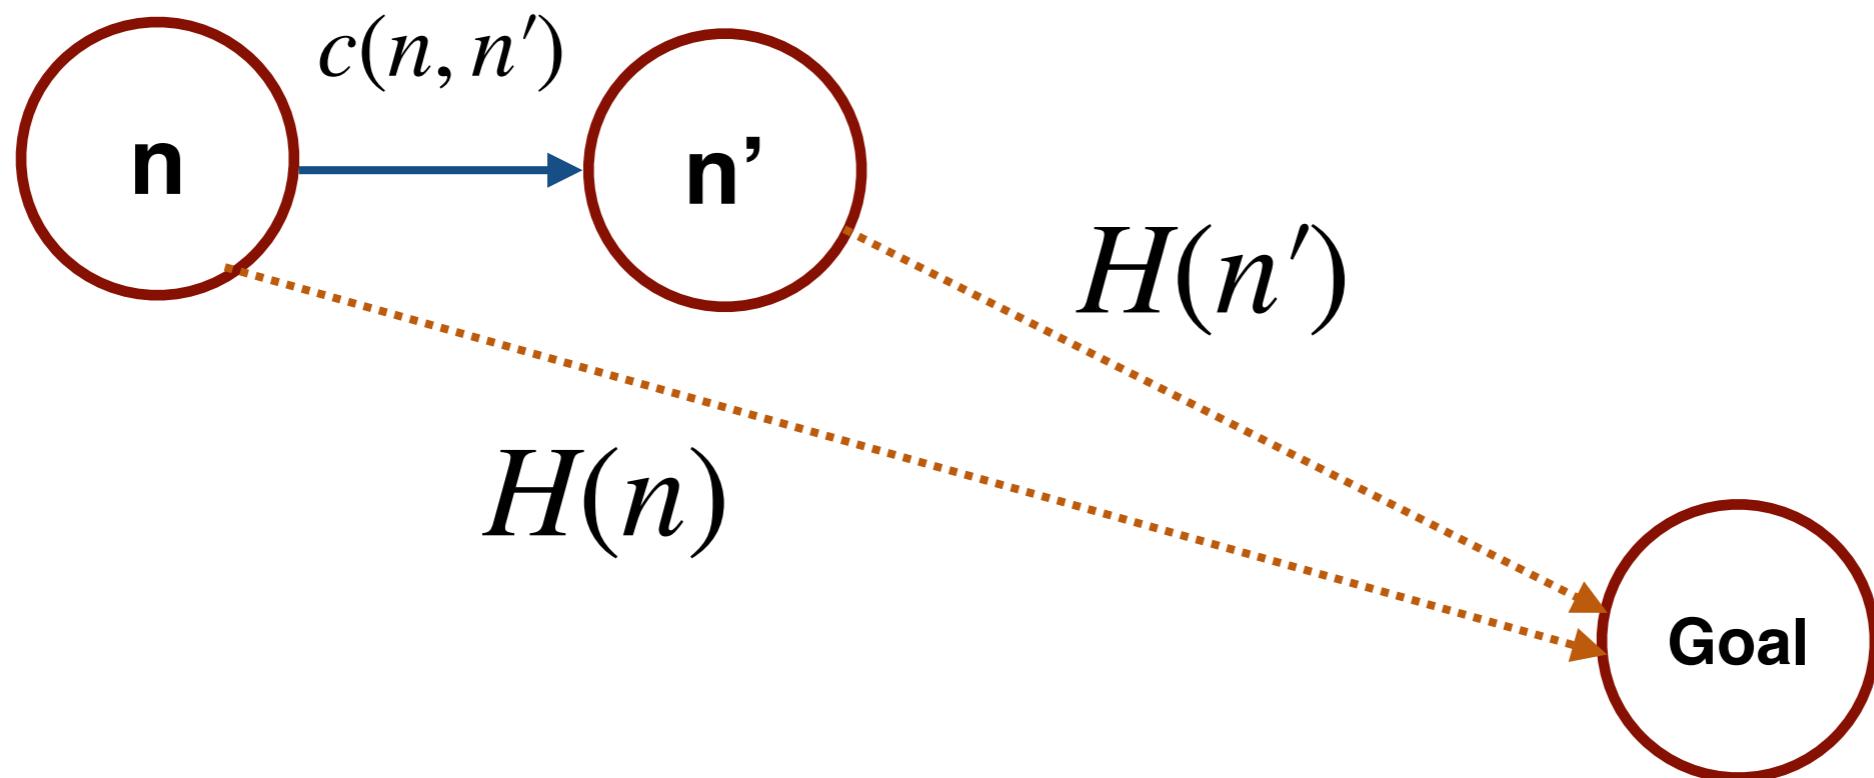
# Key requirements for good heuristics

- Basic:
  - $H(goal) = 0$
  - $H$  is always nonnegative
- When can A\* find the optimal solution?
  - What if  $H$  is super small?
  - What if  $H$  is super large?

# Analysis of A\*: Optimality

- Consistency: Whenever  $n'$  is a child of  $n$ ,

$$H(n) \leq H(n') + c(n, n')$$

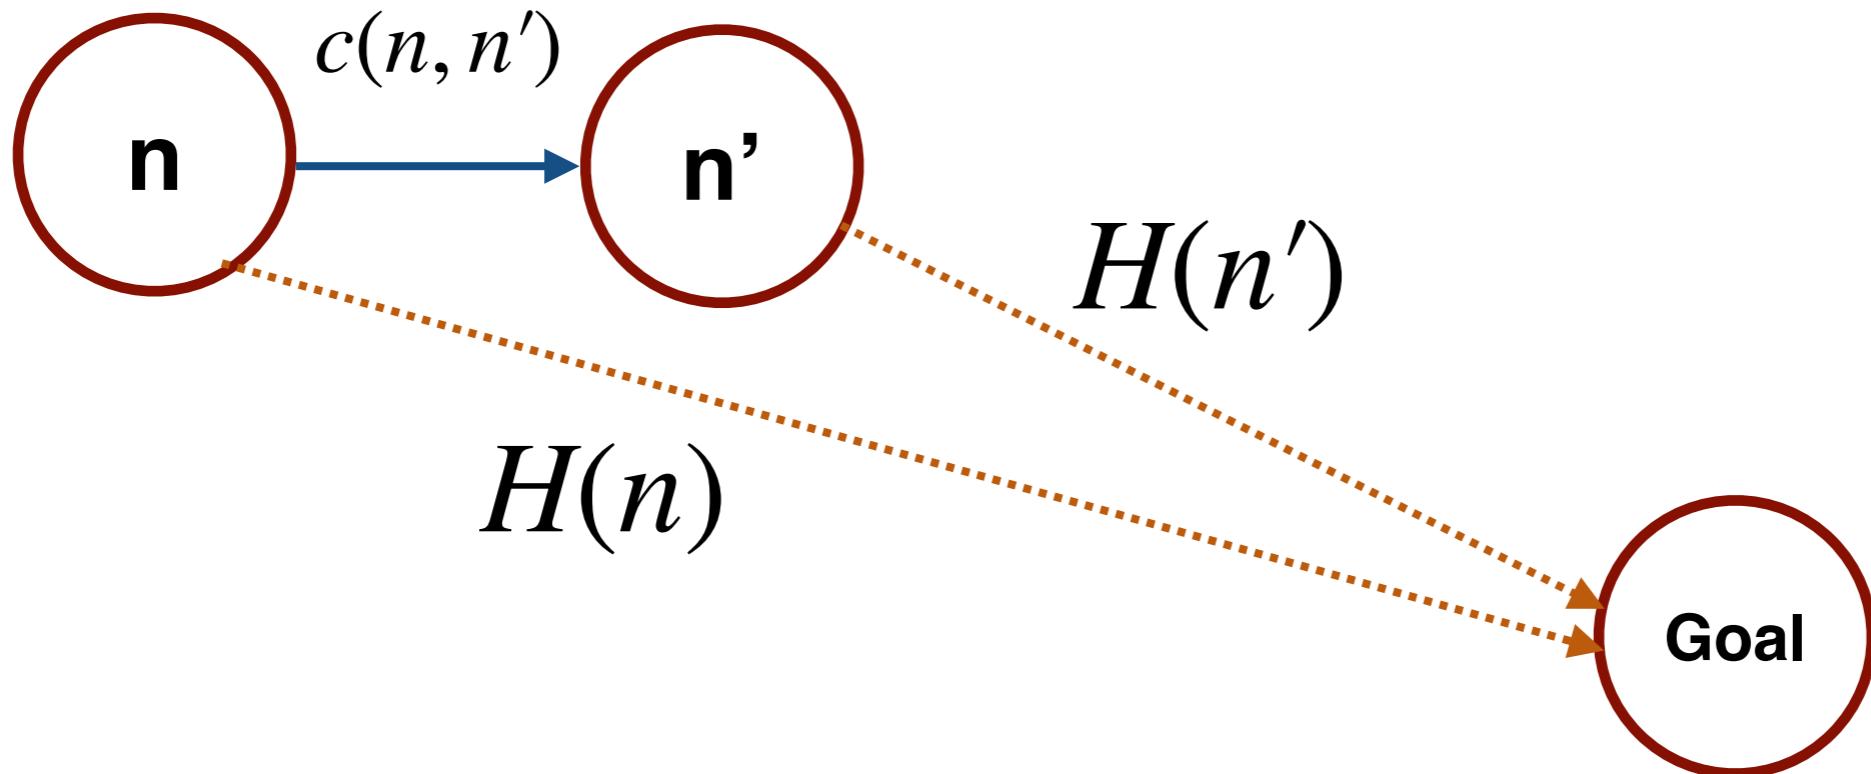


# Analysis of A\*: Optimality

- Consistency:  $H(n) \leq H(n') + c(n, n')$

- ✓ Euclidean distance and Manhattan distance

(when each step costs no less than 1)



# Analysis of A\*: Optimality

- Consistency (never give contradictory predictions): Whenever  $n'$  is a child of  $n$ ,

$$G(n) + H(n) \leq H(n') + c(n, n') + G(n)$$

# Analysis of A\*: Optimality

- Consistency (never give contradictory predictions): Whenever  $n'$  is a child of  $n$ ,

$$\underline{G(n) + H(n)} \leq H(n') + c(n, n') + G(n)$$

$$F(n)$$

# Analysis of A\*: Optimality

- Consistency (never give contradictory predictions): Whenever  $n'$  is a child of  $n$ ,

$$G(n) + H(n) \leq \underline{H(n') + c(n, n') + G(n)}$$

$$F(n')$$

# Analysis of A\*: Optimality

- Consistency (never give contradictory predictions): Whenever  $n'$  is a child of  $n$ ,

$$\frac{G(n) + H(n)}{\text{---}} \leq \frac{H(n') + c(n, n') + G(n)}{\text{---}}$$

$$F(n) \leq F(n')$$

# Analysis of A\*: Optimality

$$F(n) \leq F(n')$$

- Now recall that A\* is exactly the same algorithm as UCS, using F as the cost function.
- So from optimality of UCS, we know A\* always finds the path with lowest F cost.
- Now, what is F at the goal node??

# Analysis of A\*: Optimality

## Theorem

A\* always finds the optimal path (minimum real cost) when the heuristic function is consistent.

# Analysis of A<sup>\*</sup>:

- Consistency (never give contradictory predictions): Whenever  $n'$  is a child of  $n$ ,

$$H(n) \leq H(n') + c(n, n')$$

- Consistency implies Admissibility:

$$H(n) \leq \min_{\pi} \text{Cost}(\pi(n, \text{goal}))$$

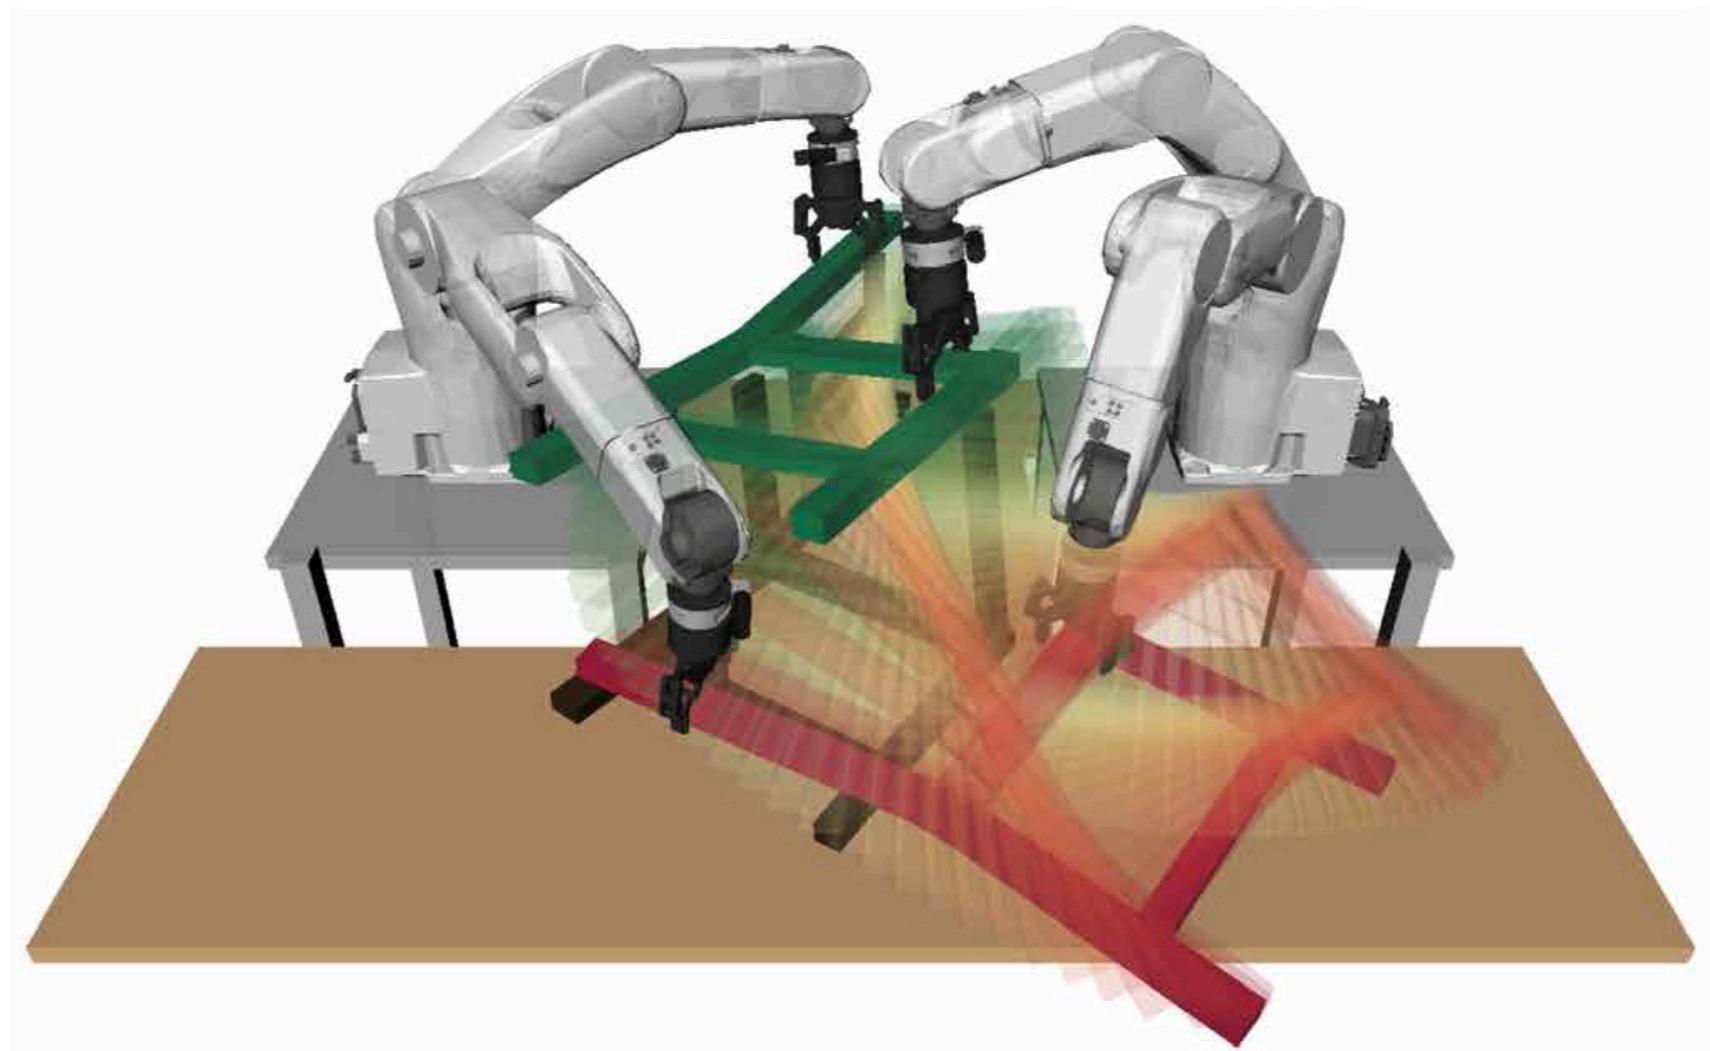
# But in the end, what is good?

Trade off between optimality and speed.

- Big H values: faster search
- Small H values: more optimal

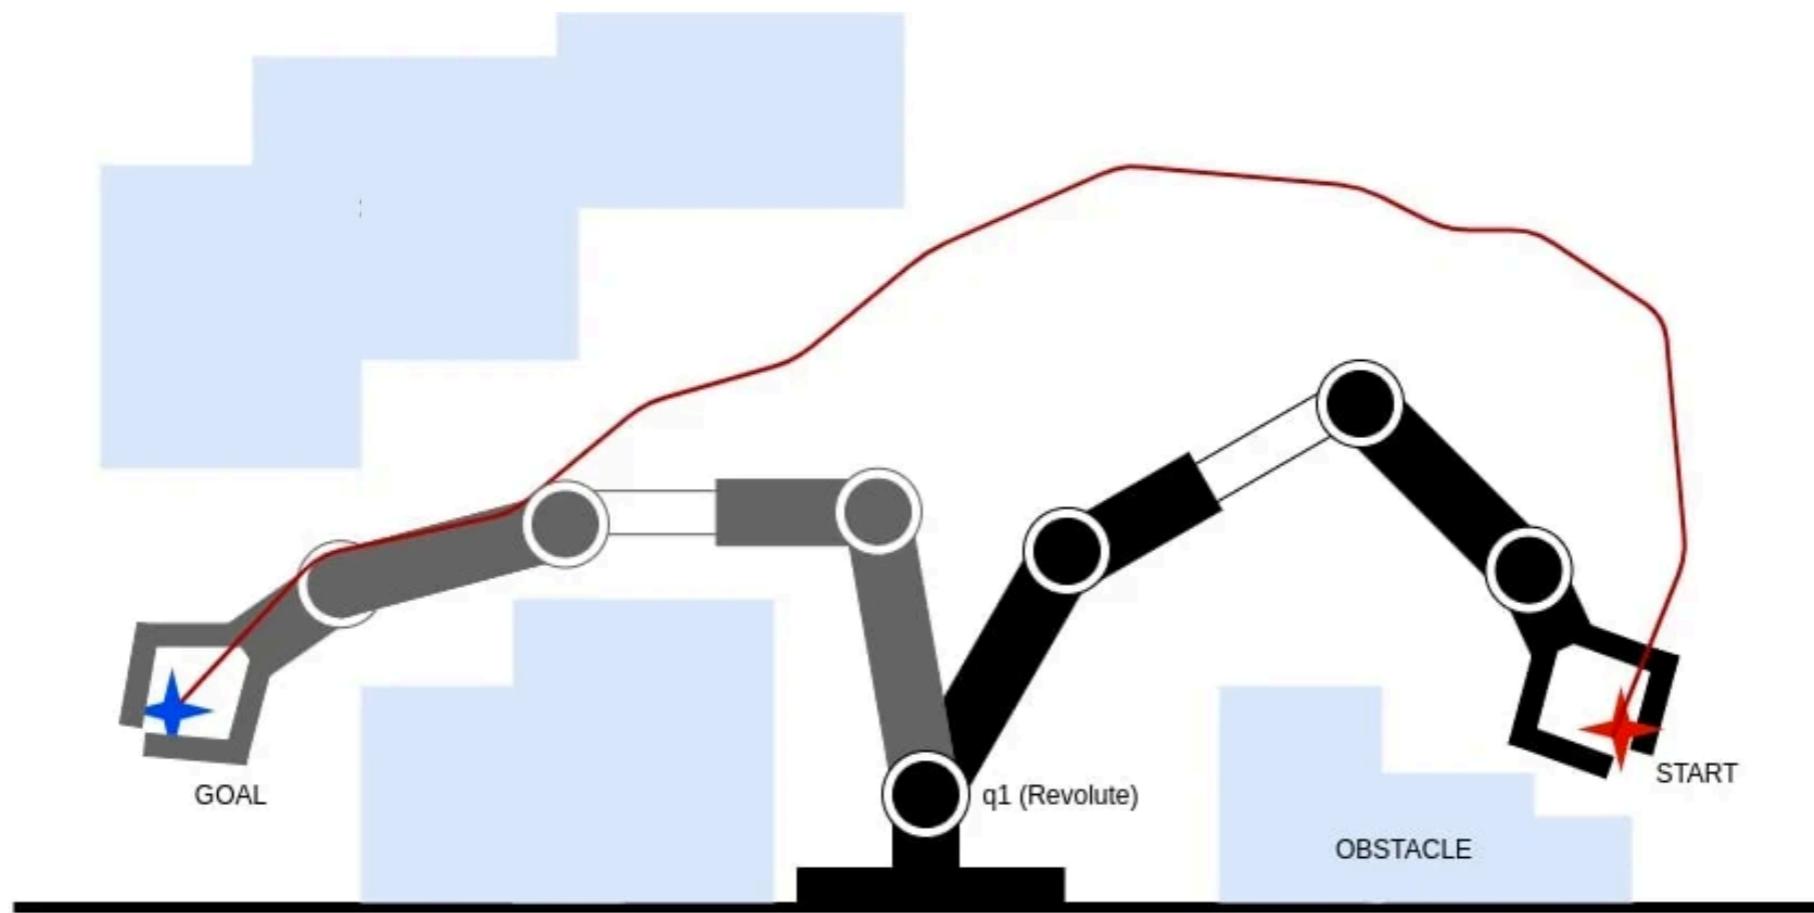
This is where science becomes art. Learning can often help you.

# Motion Planning (possible project)



# Motion Planning

- What if there is no explicit graph?



# General Path Finding

- Configuration space: high-dimensional
- Work space: at most 3D
- Find a path (sequence of states in the configuration space) from start to goal such that each segment of the path is collision-free
  - Collision checking can be very expensive

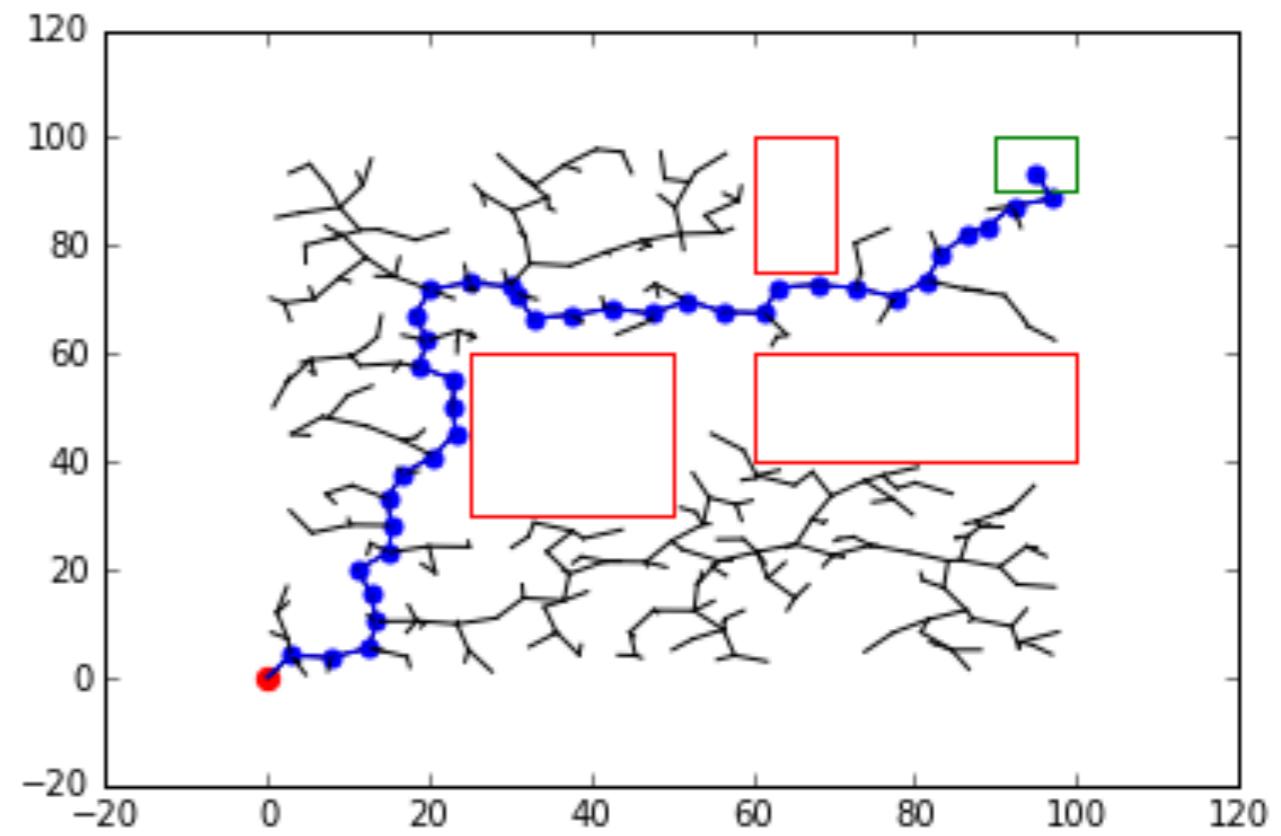
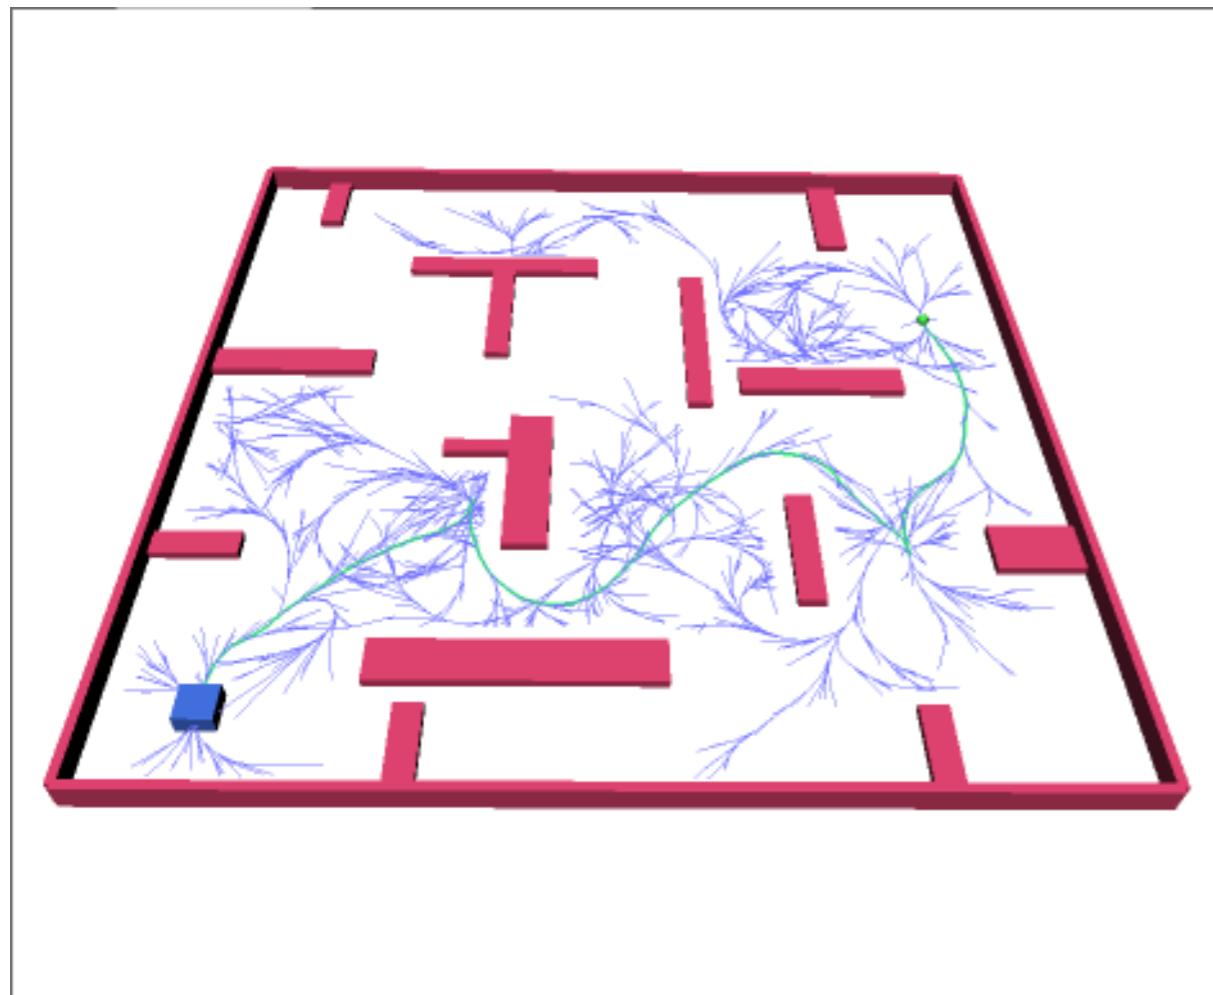
# General Path Finding

- Configuration space  $C \subseteq \mathbb{R}^n$ ,  $C = C_o \cup C_f$ 
  - Obstacles  $C_o$
  - Free space  $C_f$
- Start  $s_0 \in C_f$  and goal  $s_g \in C_f$
- Find state sequence  $s_0, \dots, s_k$  such that  $s_k = s_g$  and the edge  $E(s_i, s_{i+1}) \subseteq C_f$  for all  $0 \leq i < k$

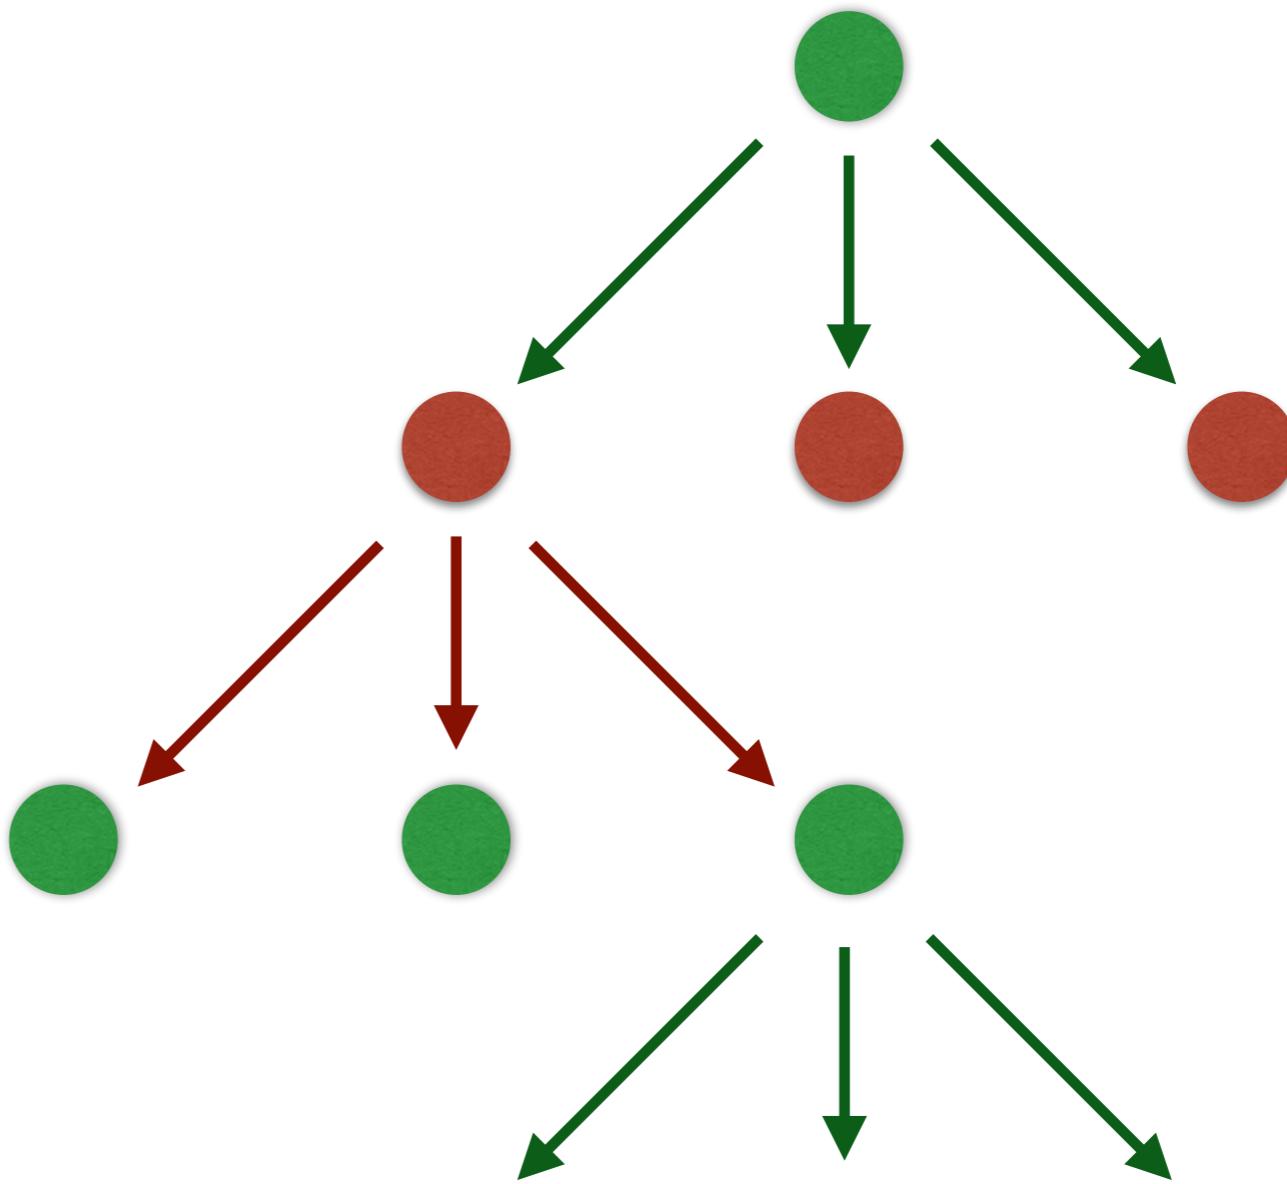
# Rapidly Exploring Random Trees (RRT)

- Grow a tree in the free space
  - In each iteration, randomly sample a state from the free space.
  - Find the node in the current tree nearest to the sampled state, and grow the tree at that node by a small step
  - Periodically check if the goal is a small distance away from any node in the tree

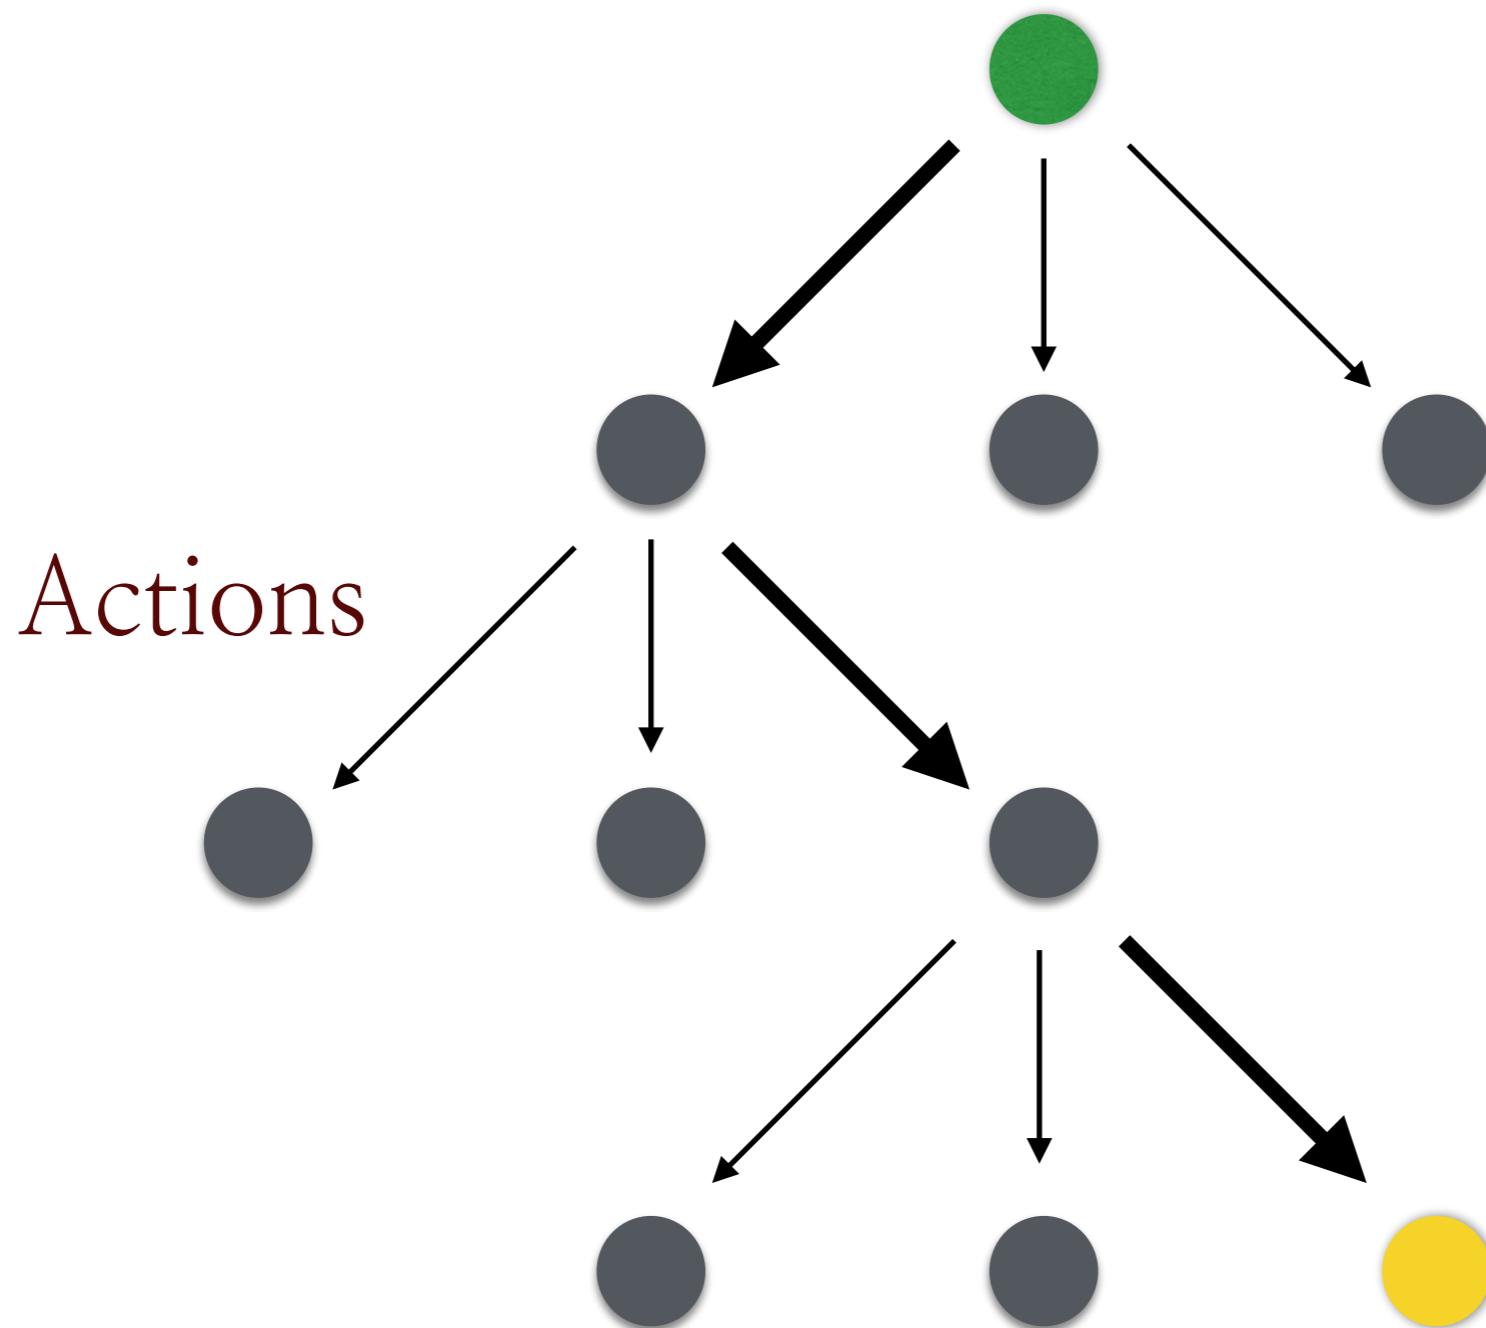
# Rapidly Exploring Random Trees (RRT)



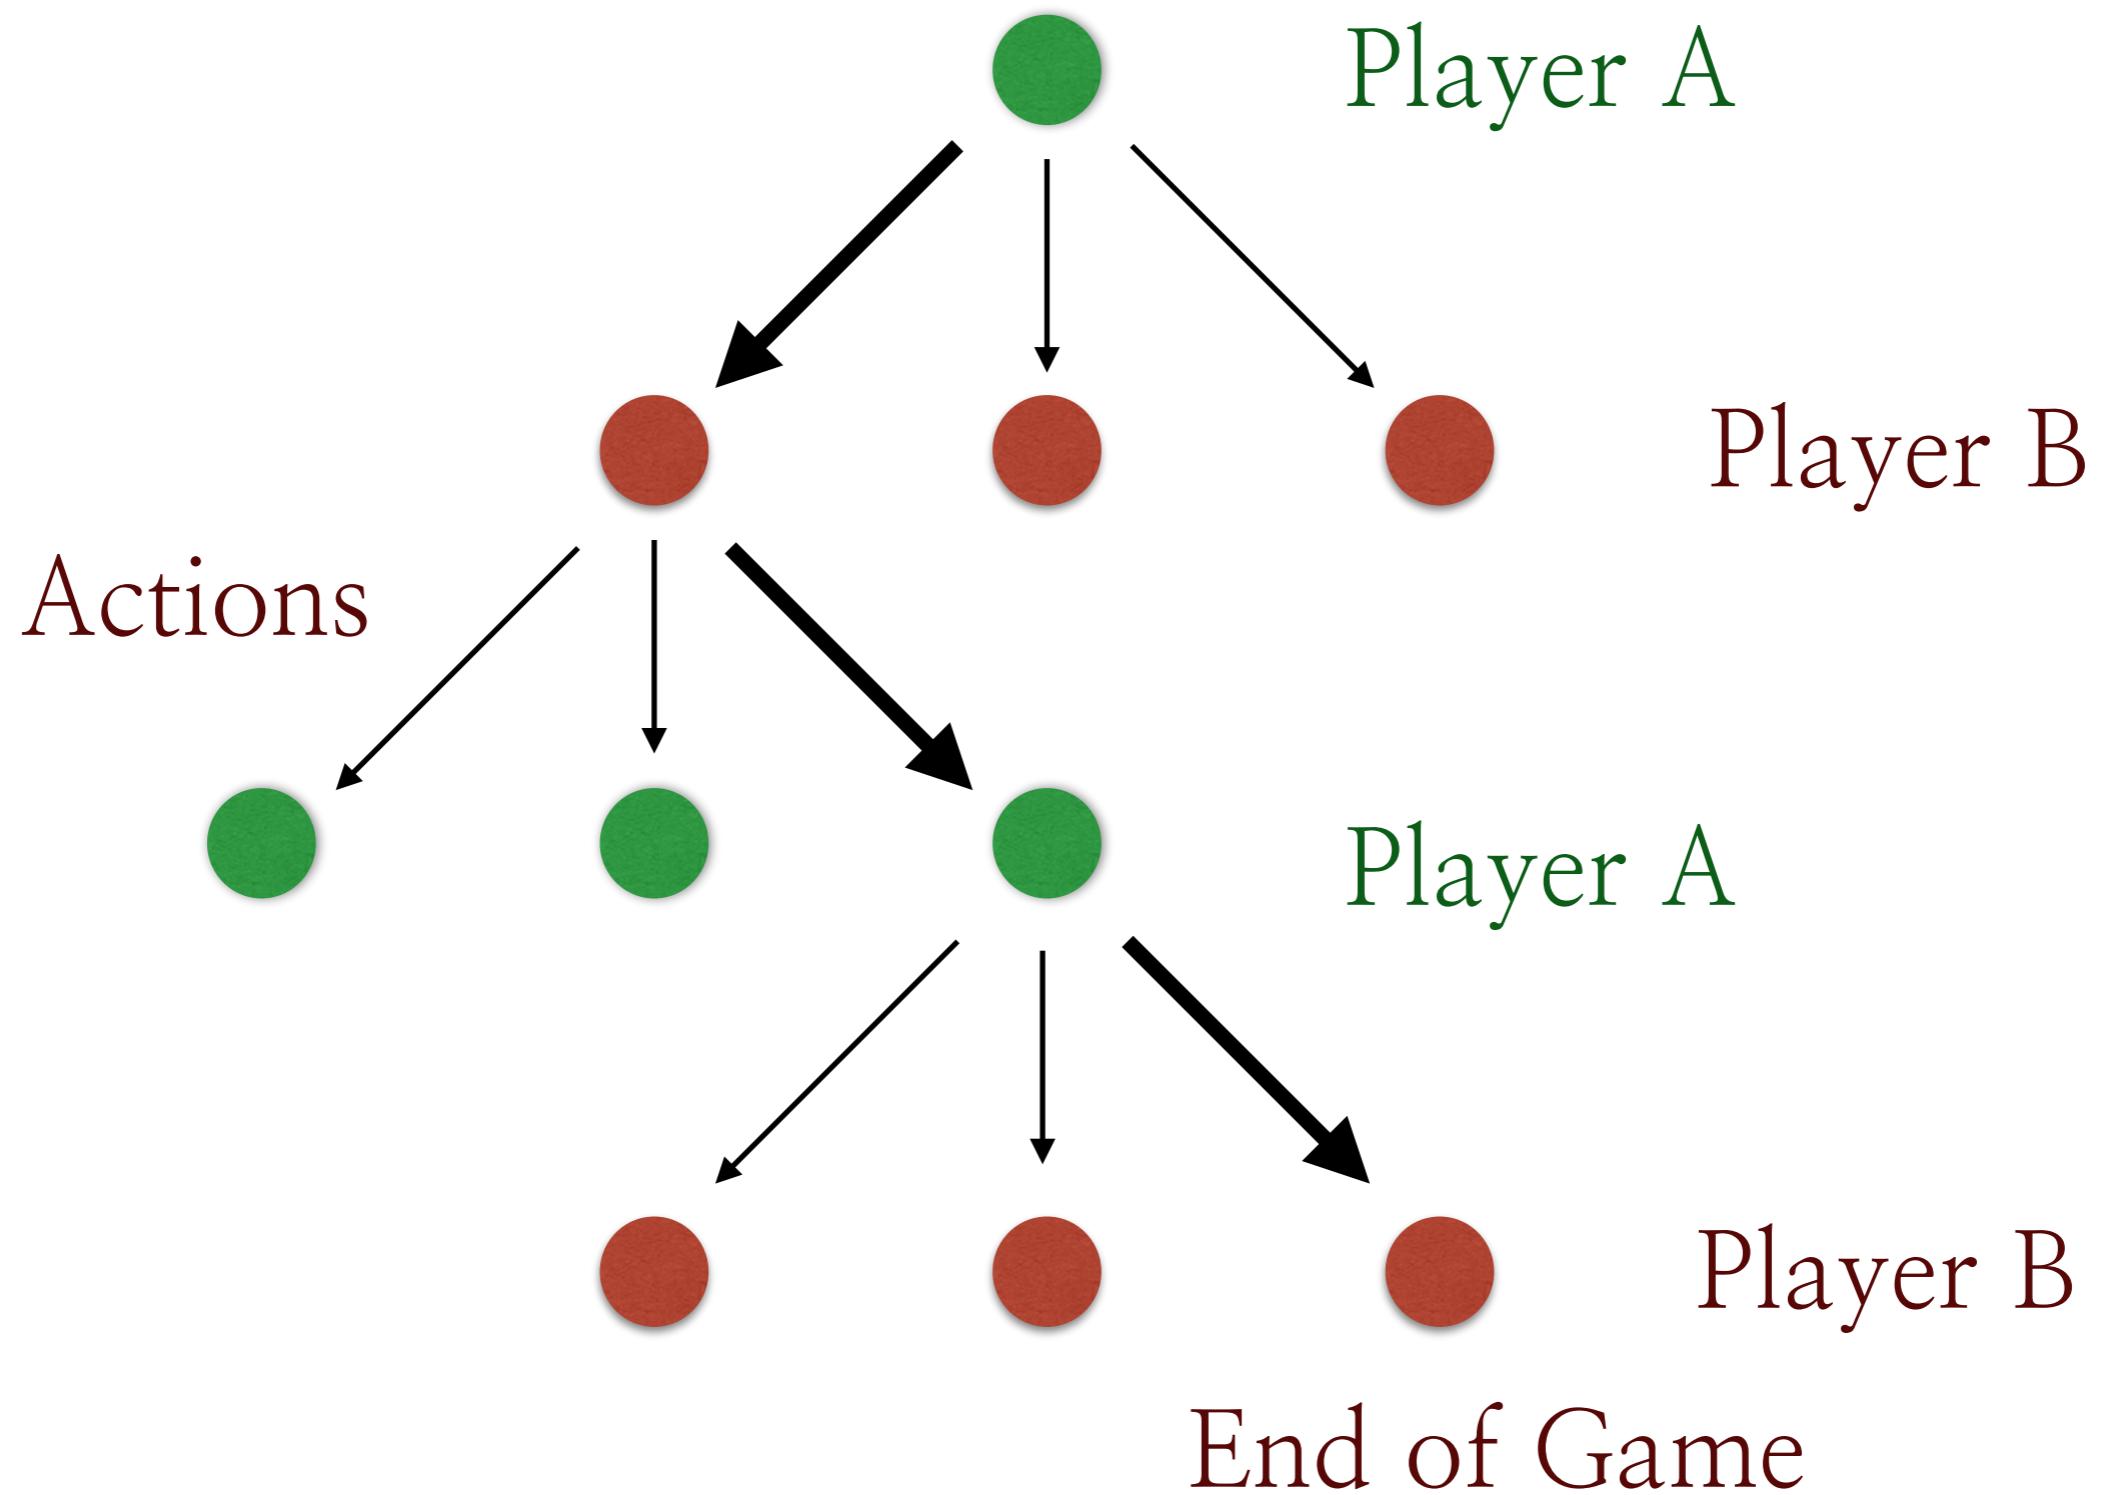
# Adversarial Search



Initial State



Goal State



C wants max, R wants min, C picks first

	C1	C2
R1	7	5
R2	2	10

C wants max, R wants min, C picks first

	C1	C2
R1	7	5
R2	2	10

Moves: (C2, R1)

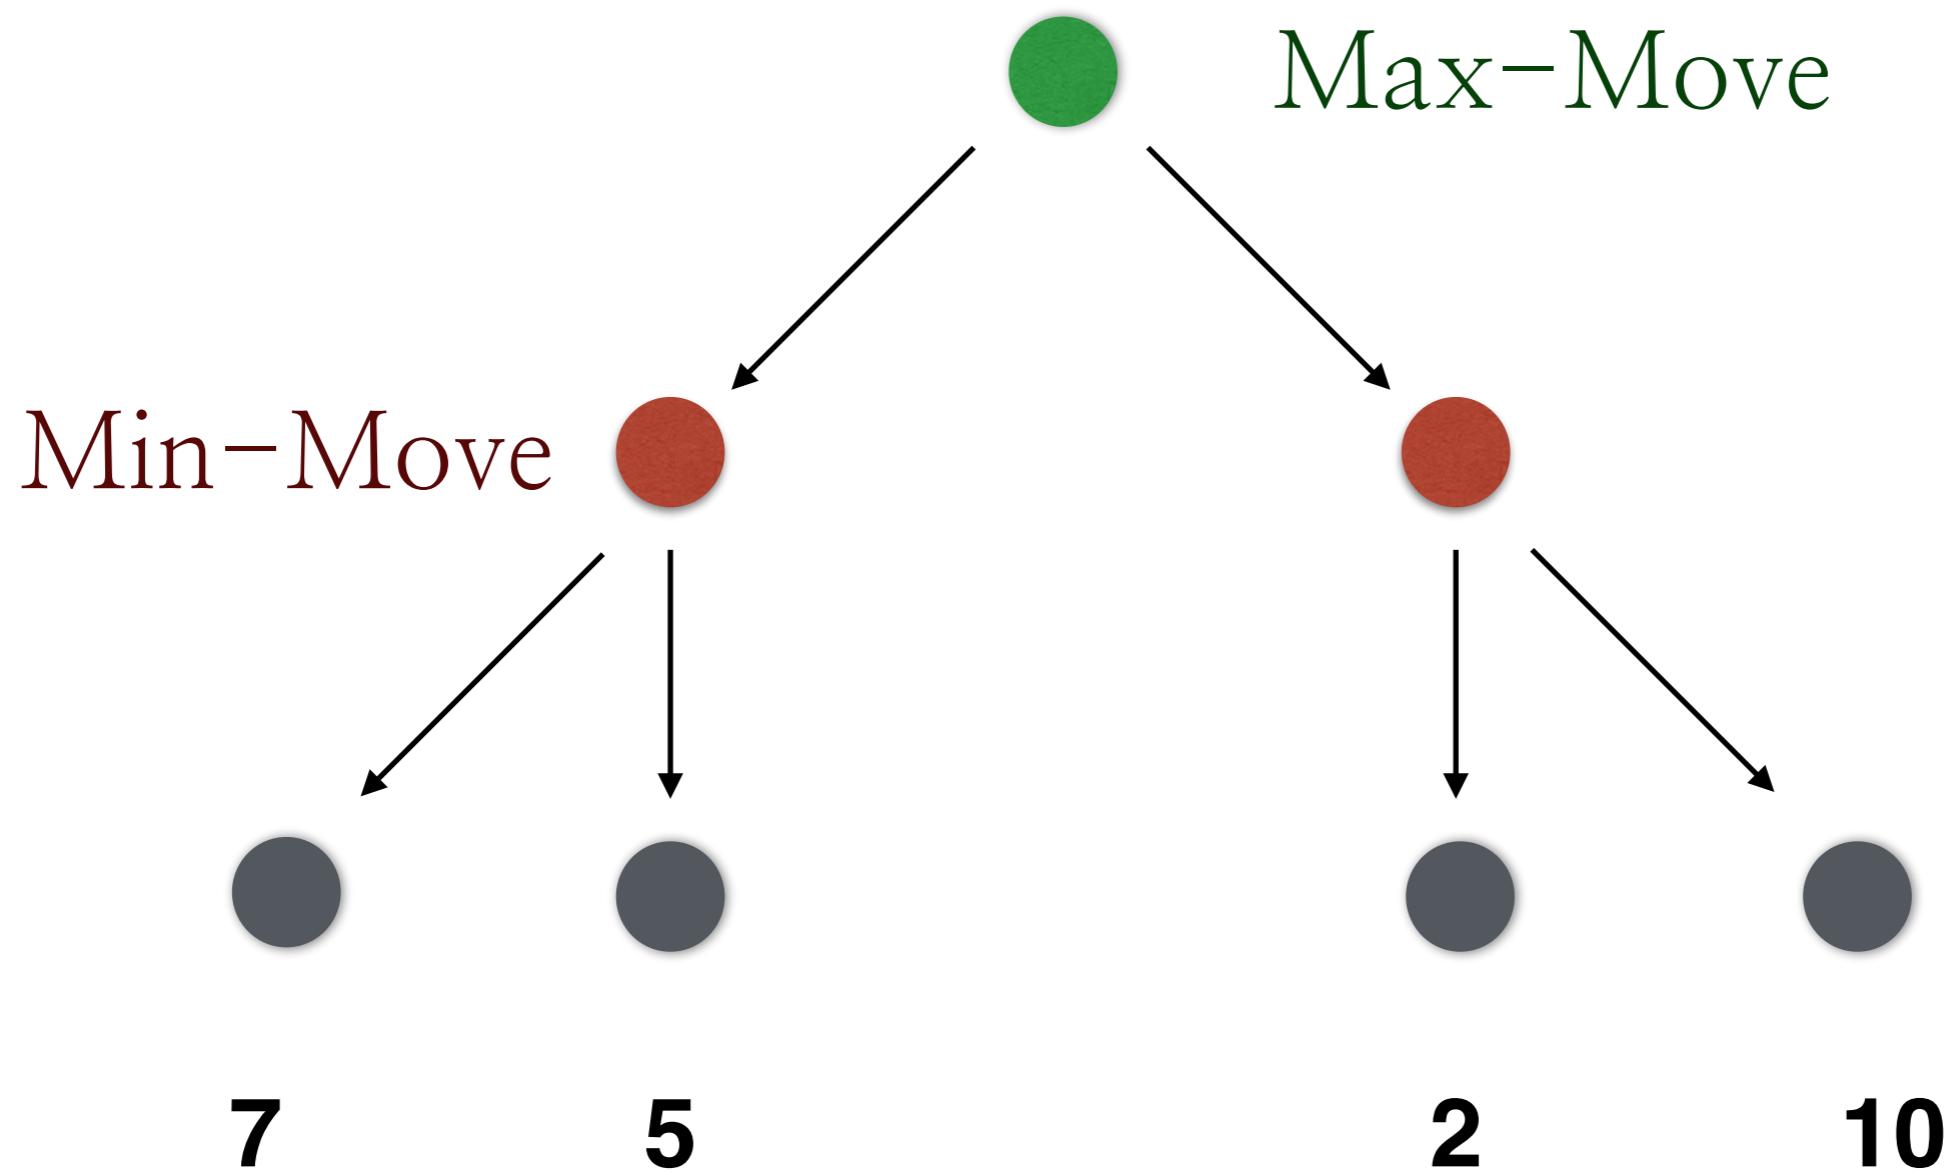
C wants max, R wants min, R picks first

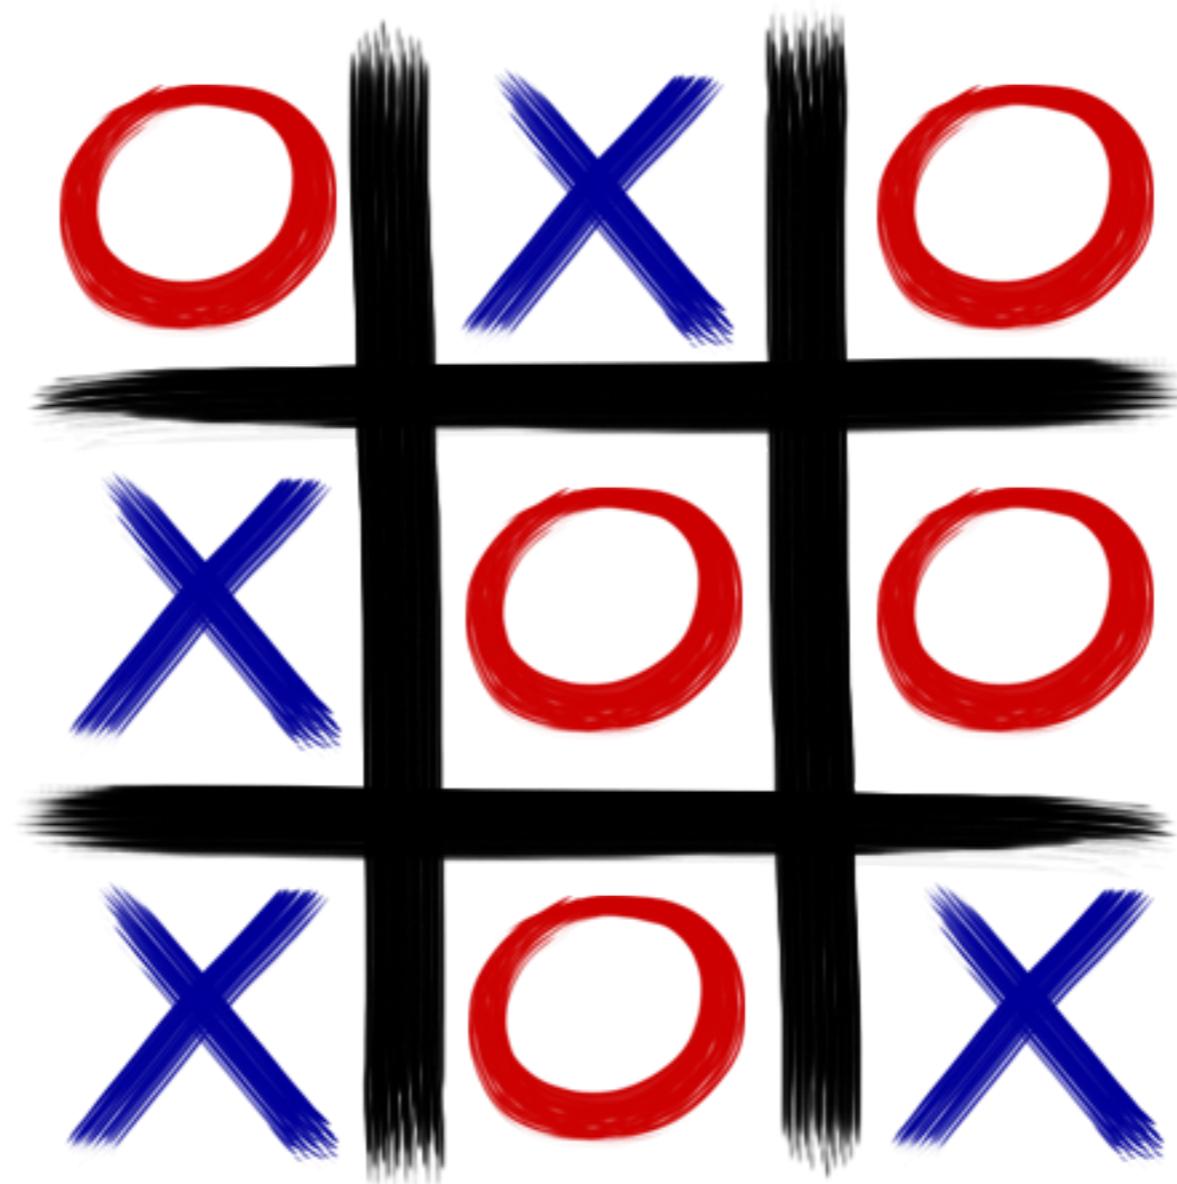
	C1	C2
R1	7	5
R2	2	10

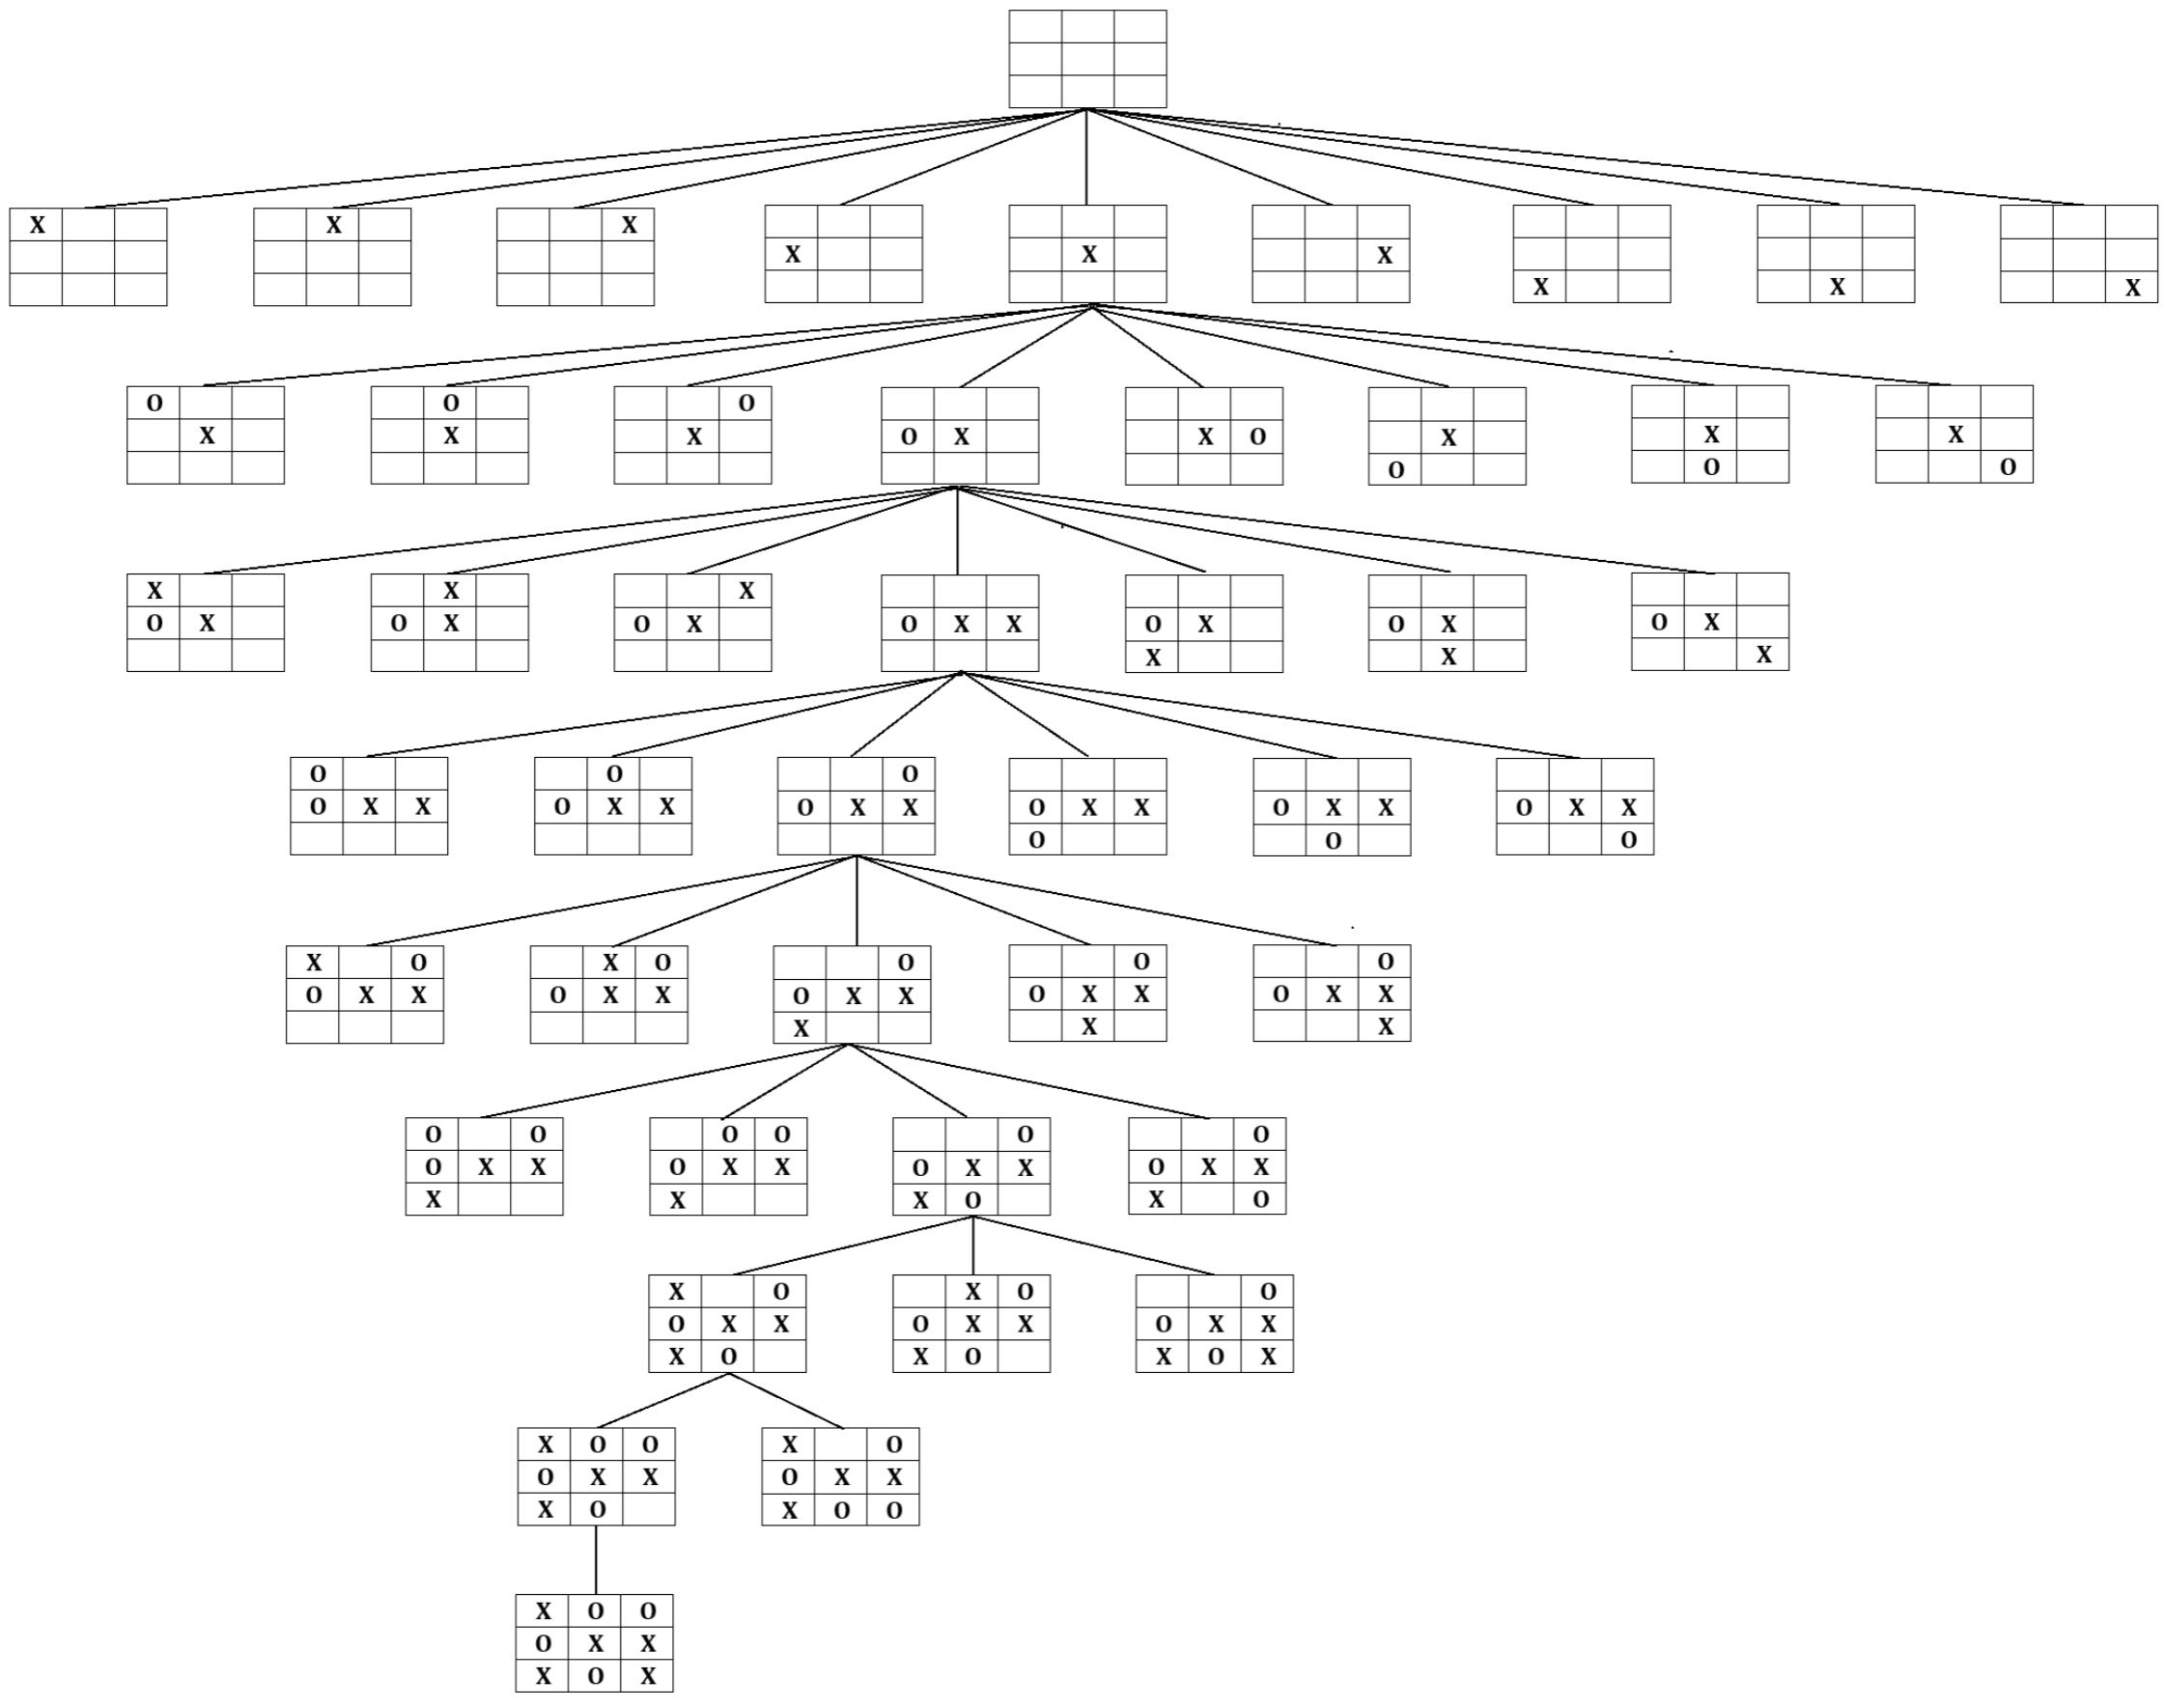
C wants max, R wants min, R picks first

	C1	C2
R1	7	5
R2	2	10

Moves: (R1, C1)

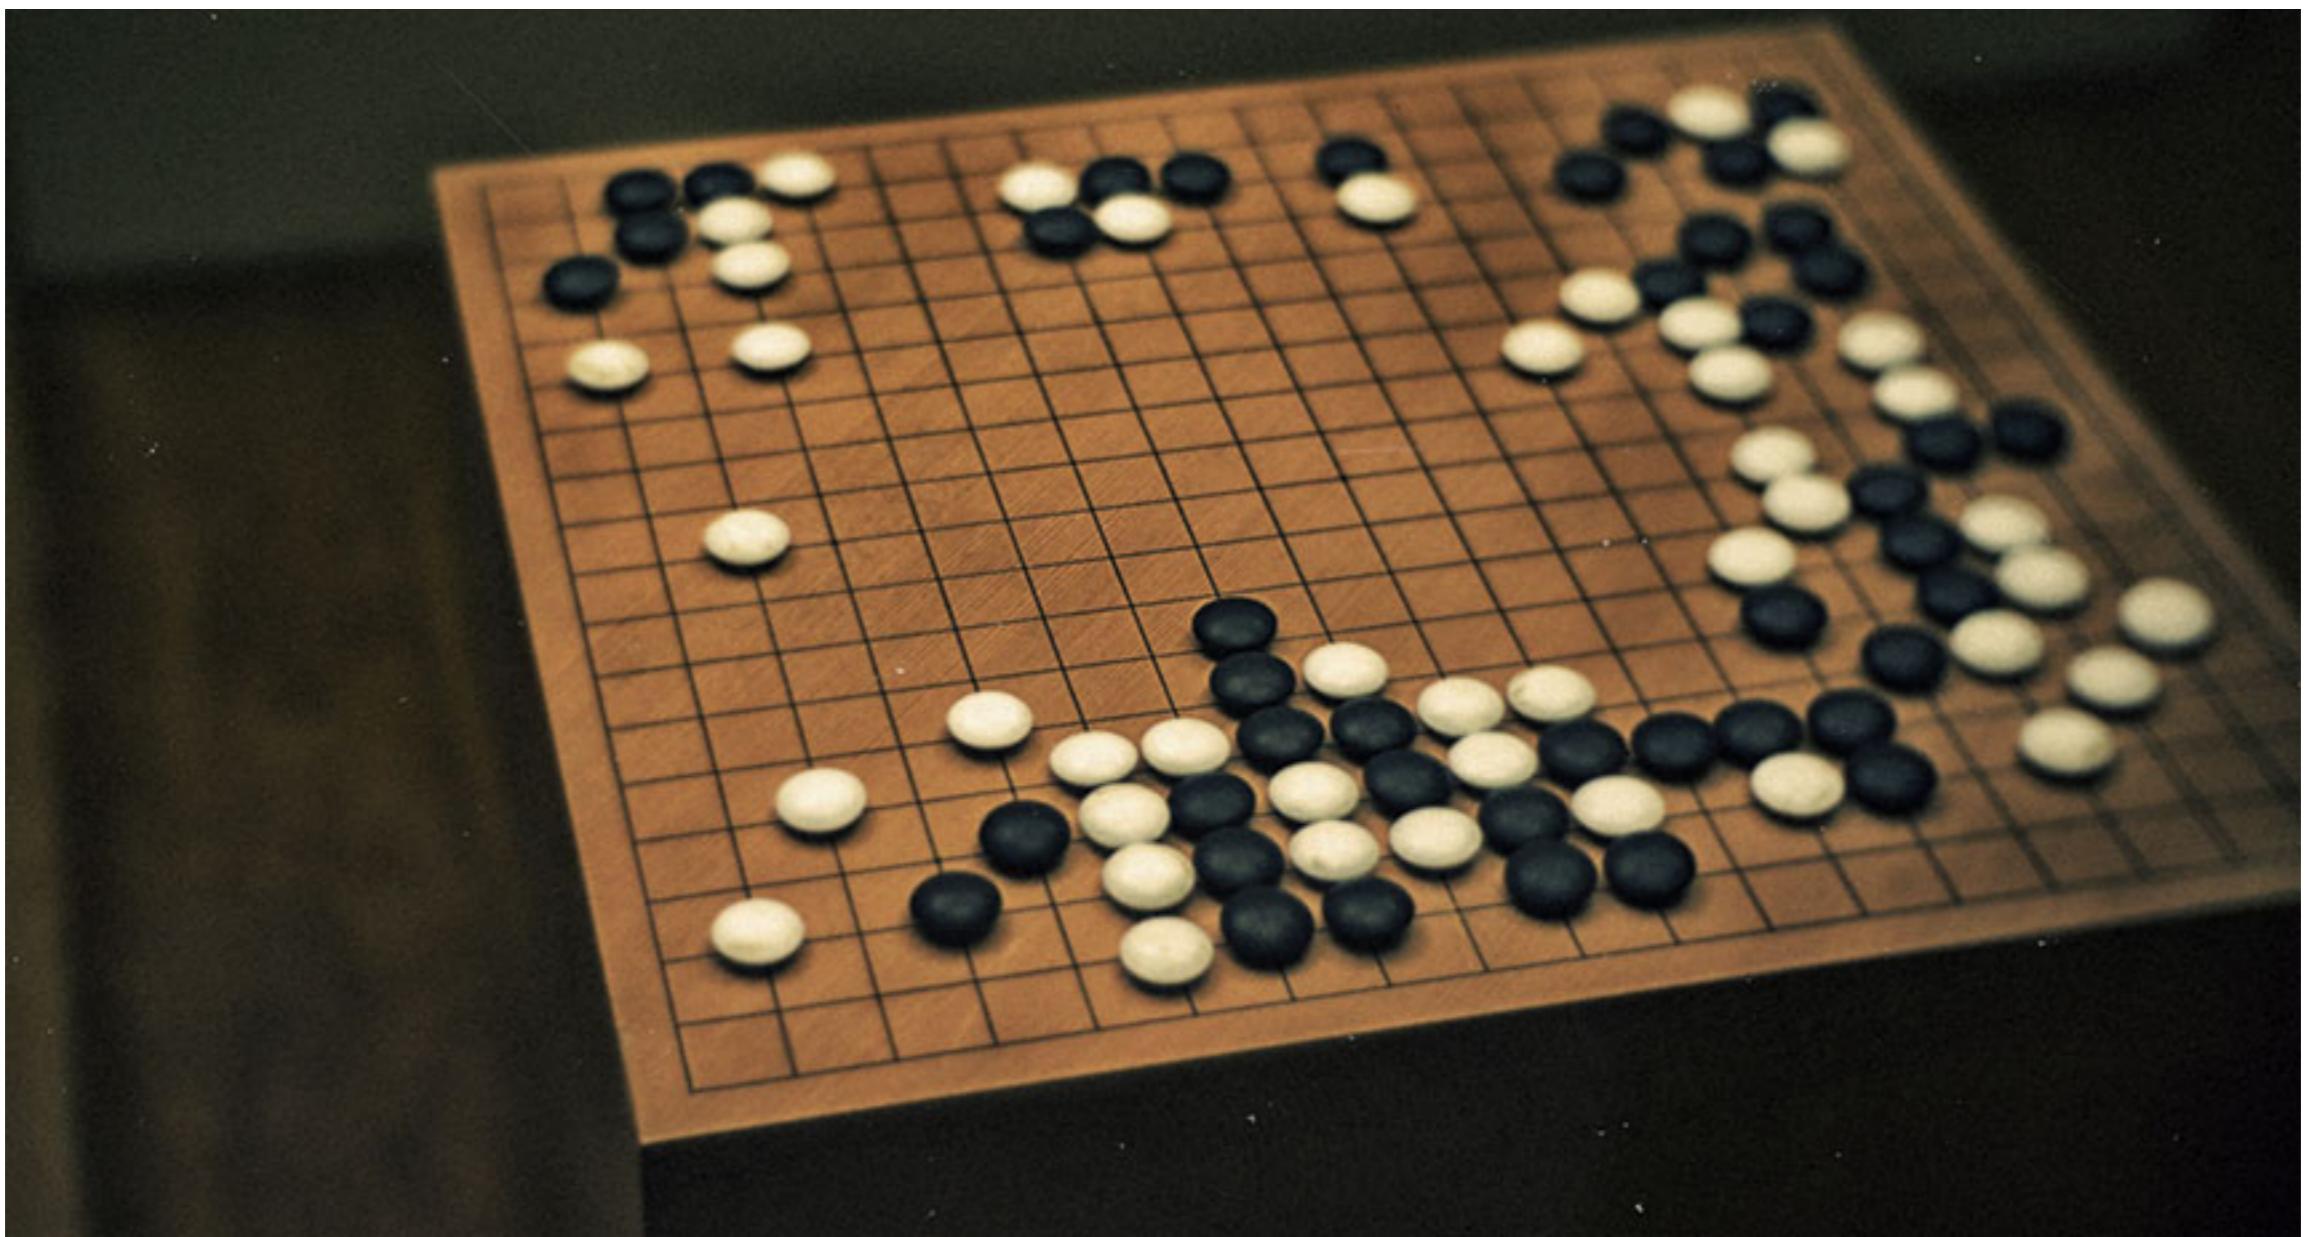






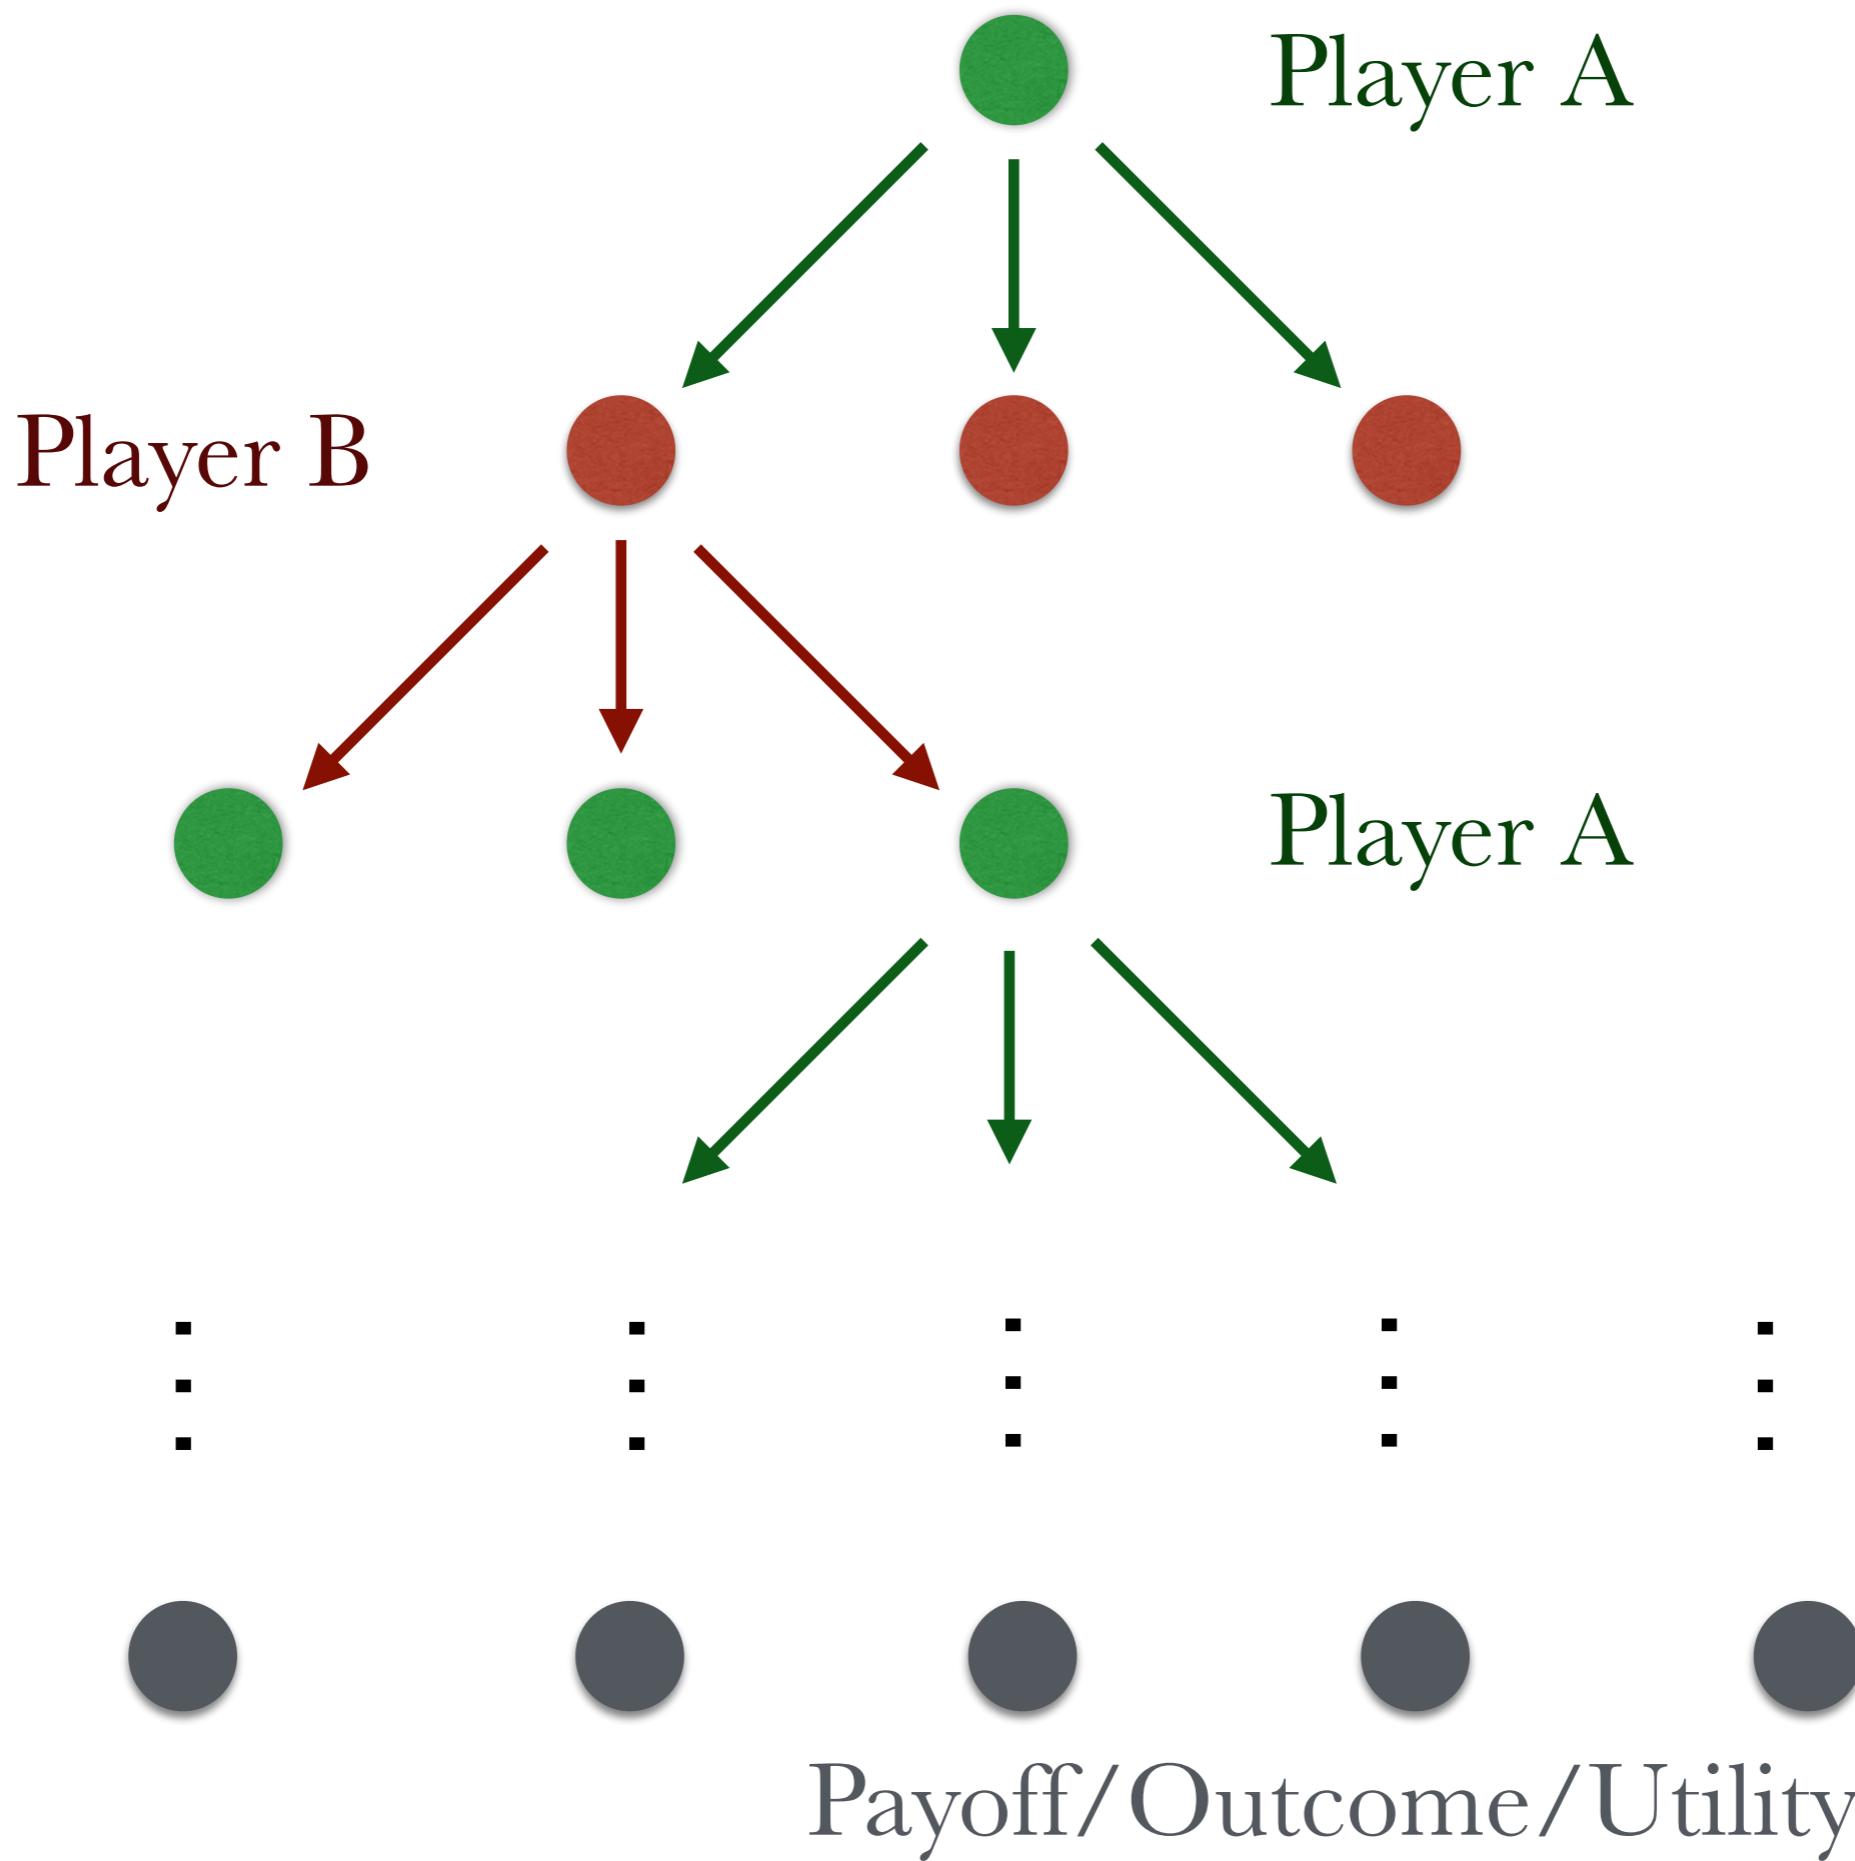
**(Winner - X)**

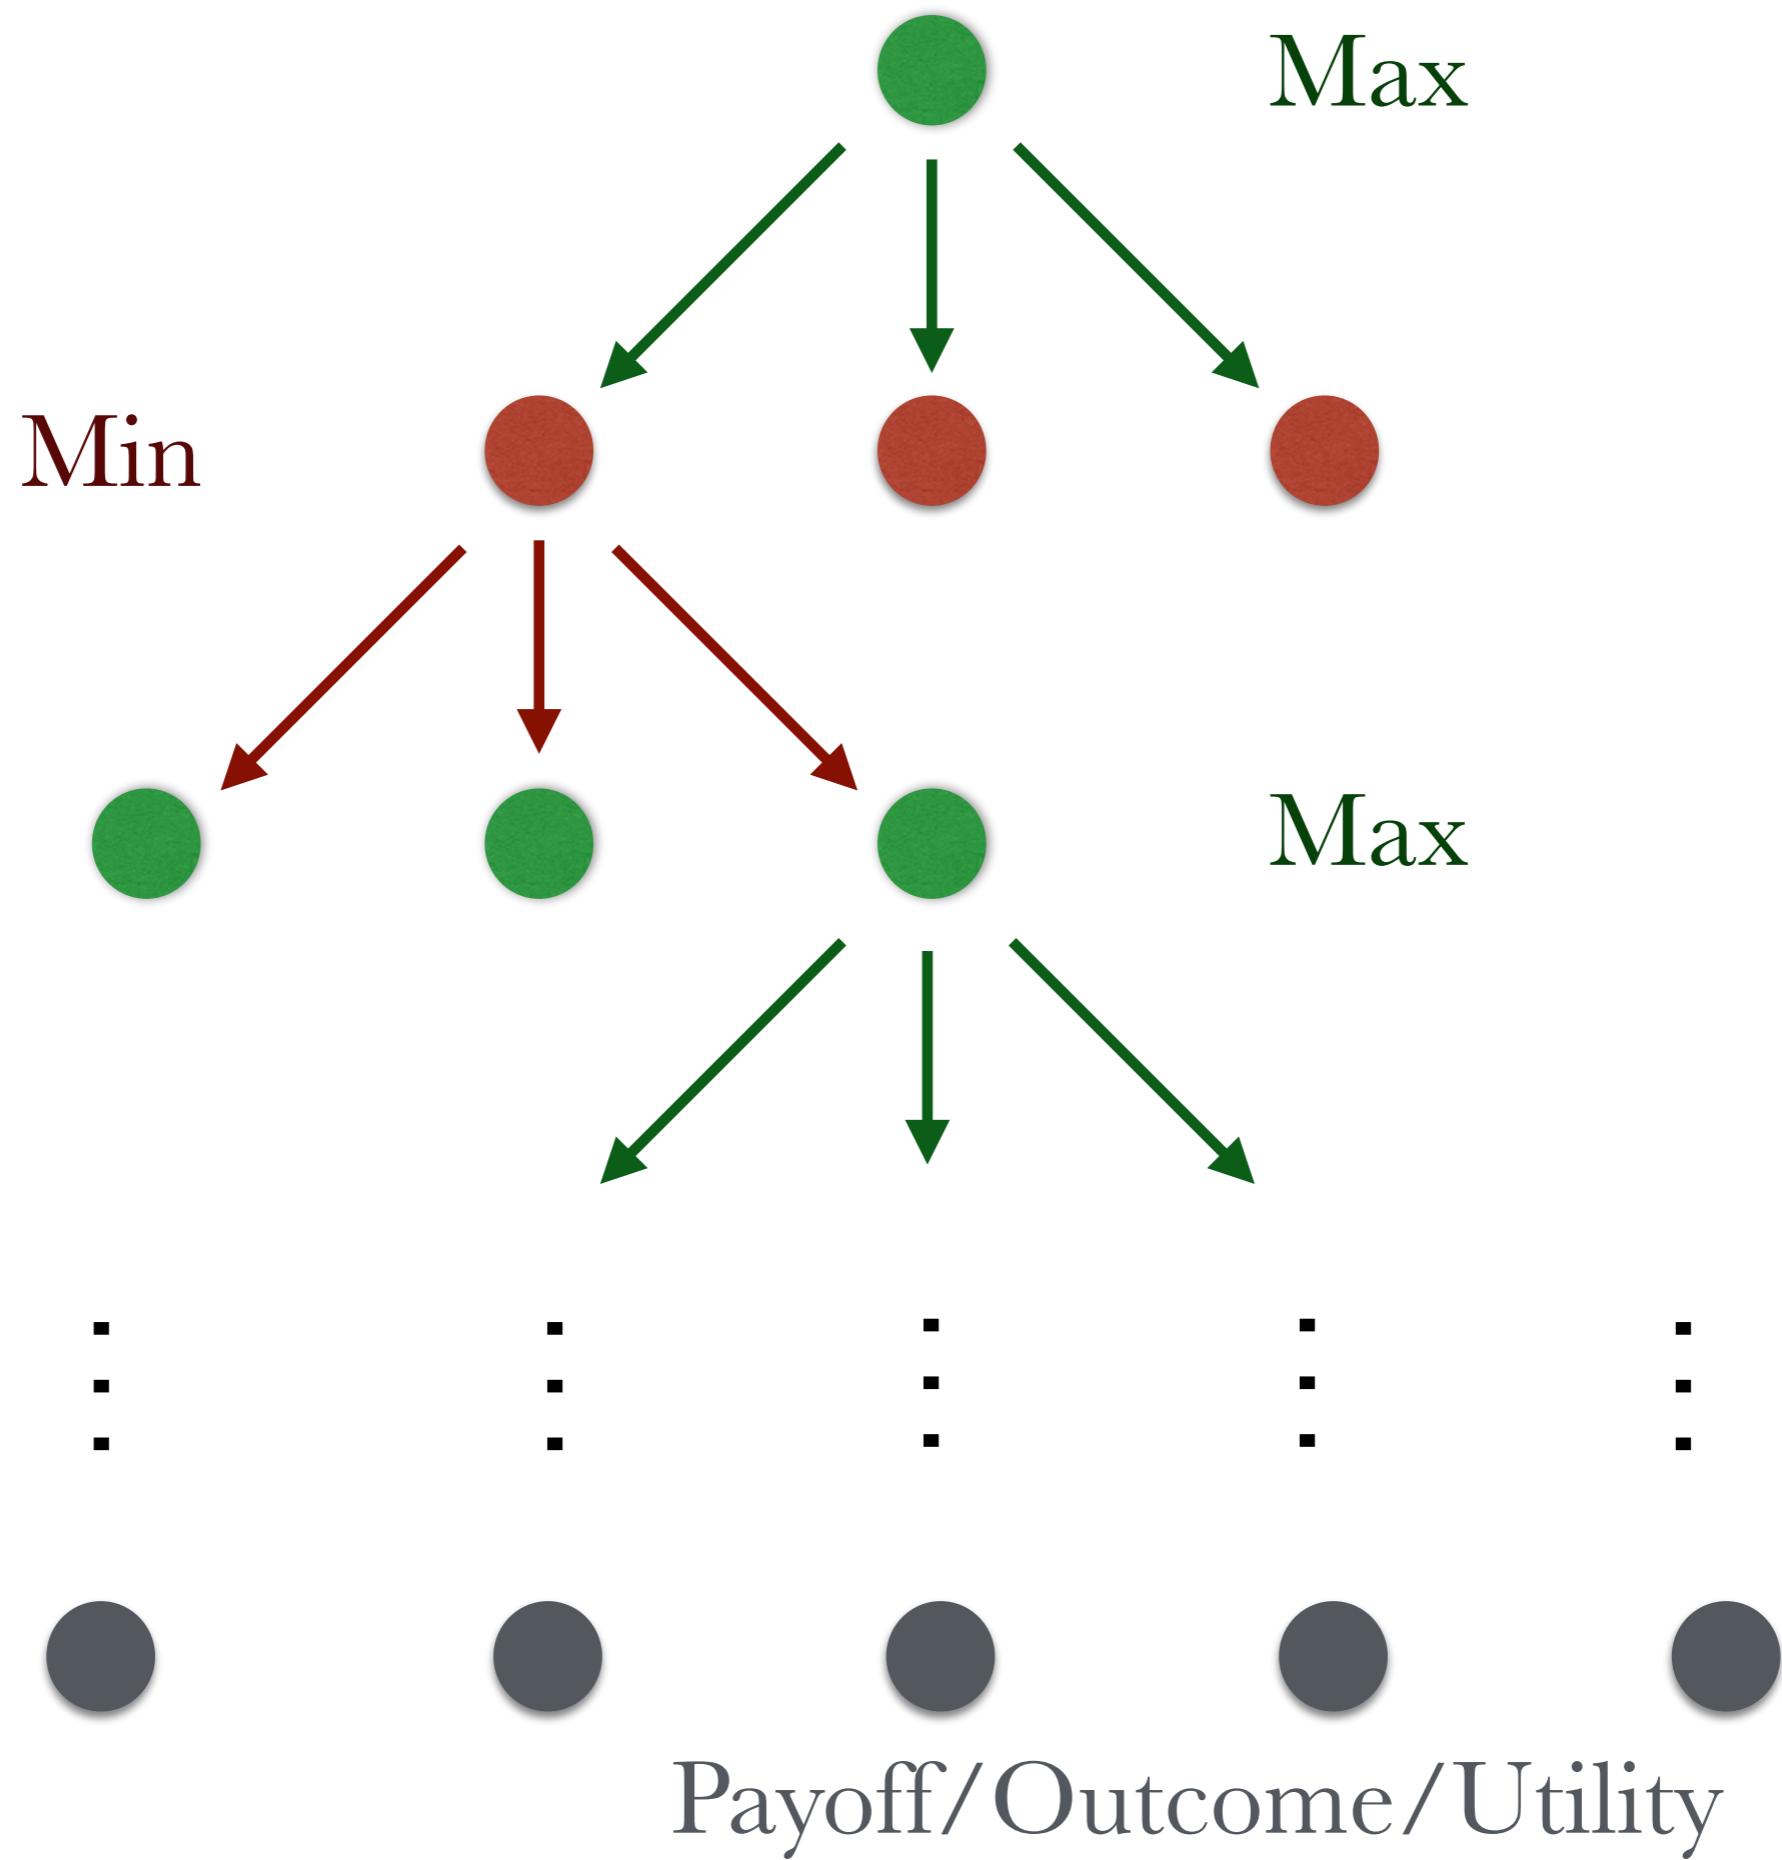


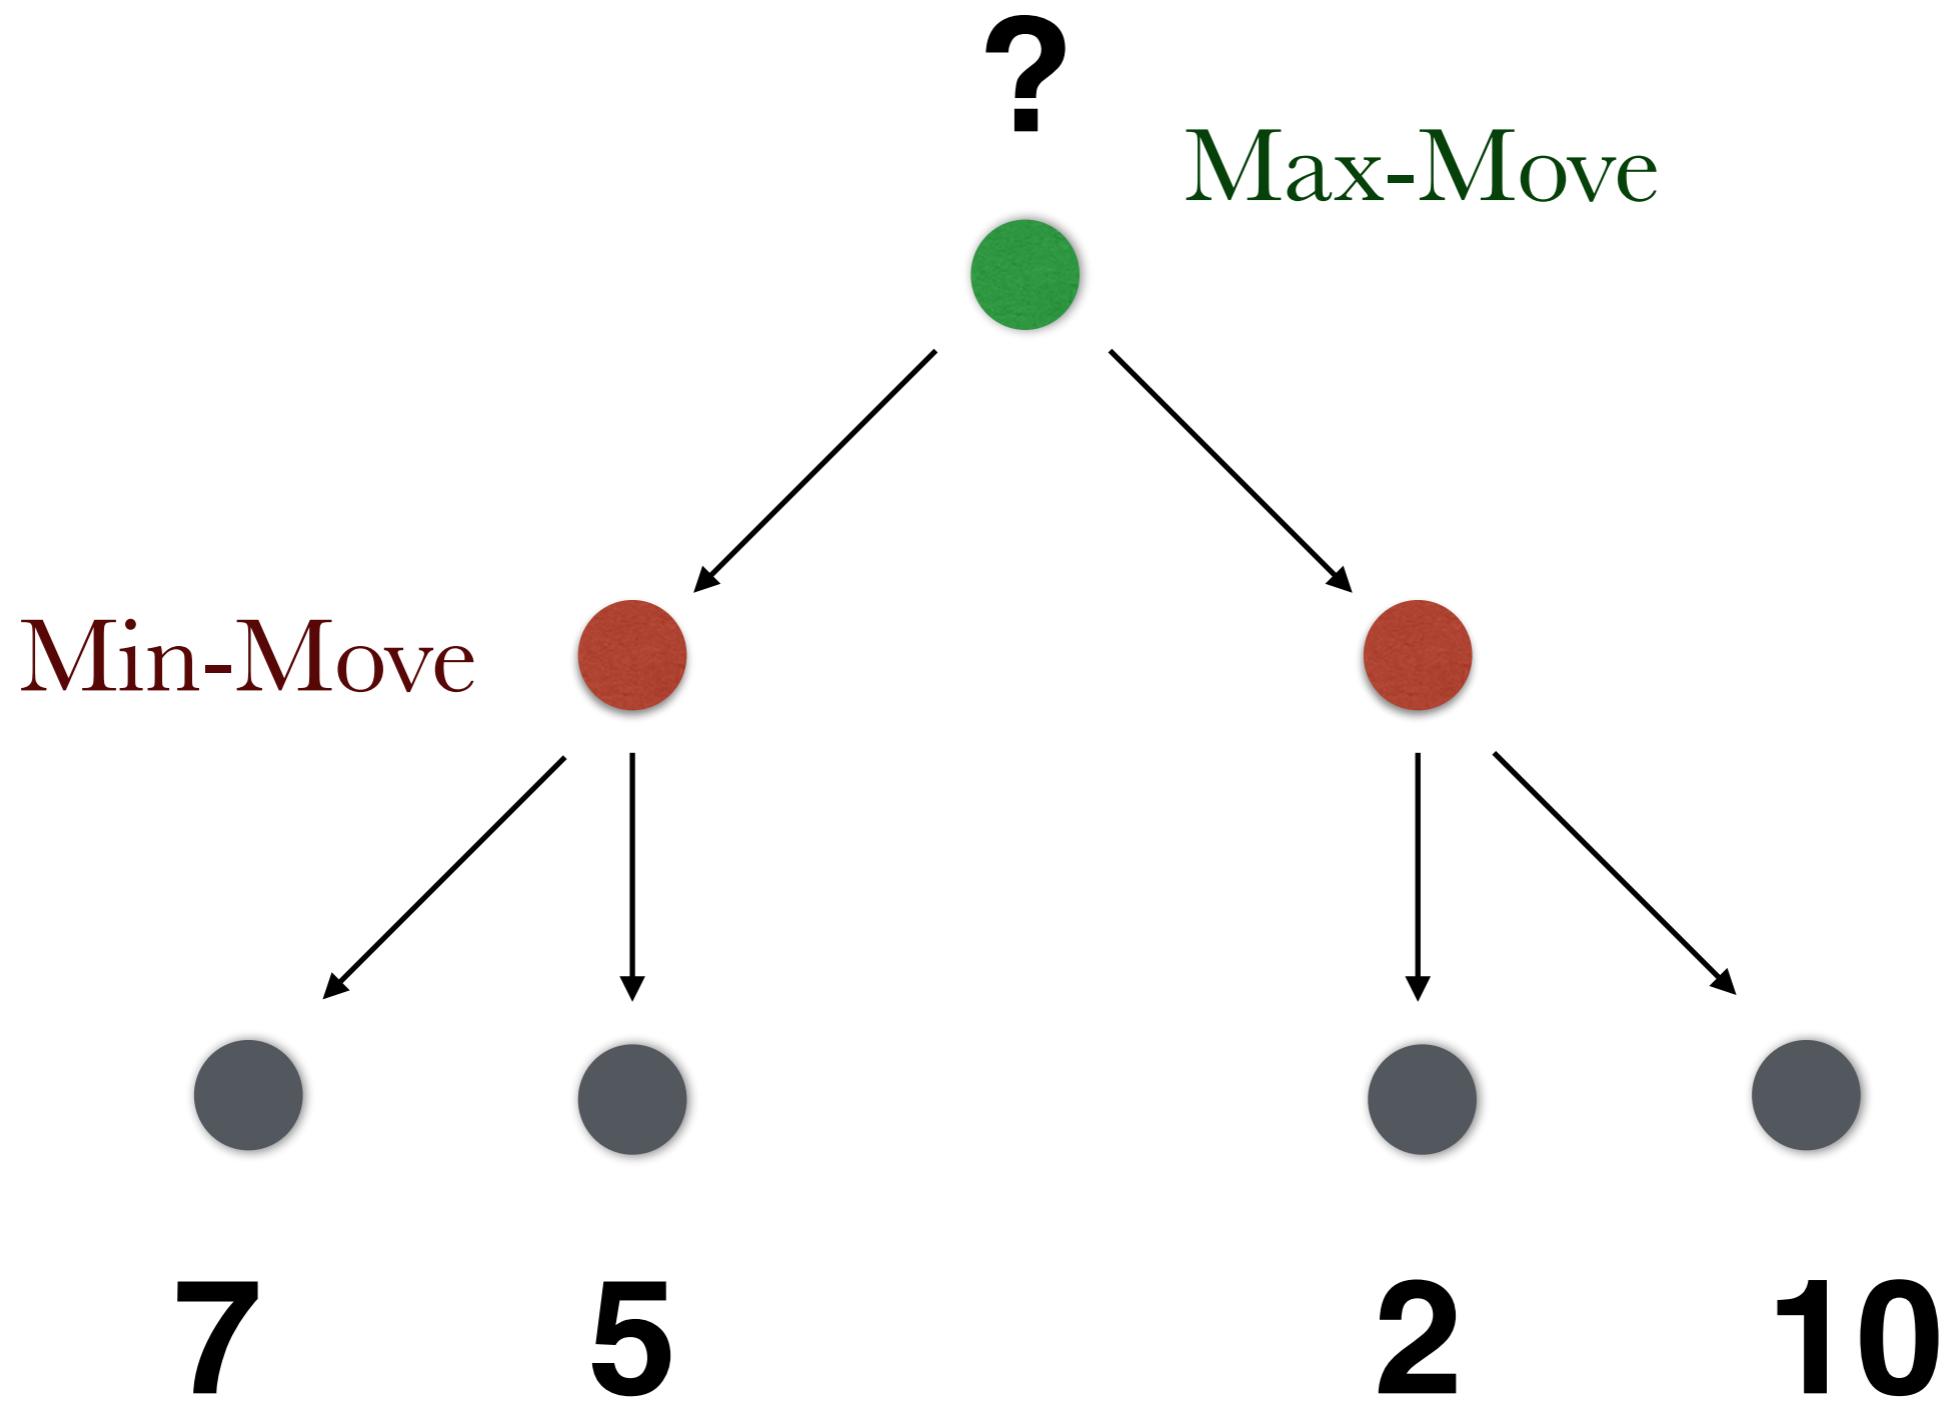


# Basic Two-Player Zero-Sum Games

- Two players, take turns
- Actions: each player has a finite number of actions
- Finite length: all sequences of moves terminate
- Zero-Sum: strictly competitive, the gain of one player at the end is exactly the loss of the other player
- Perfect information







# Minimax Values

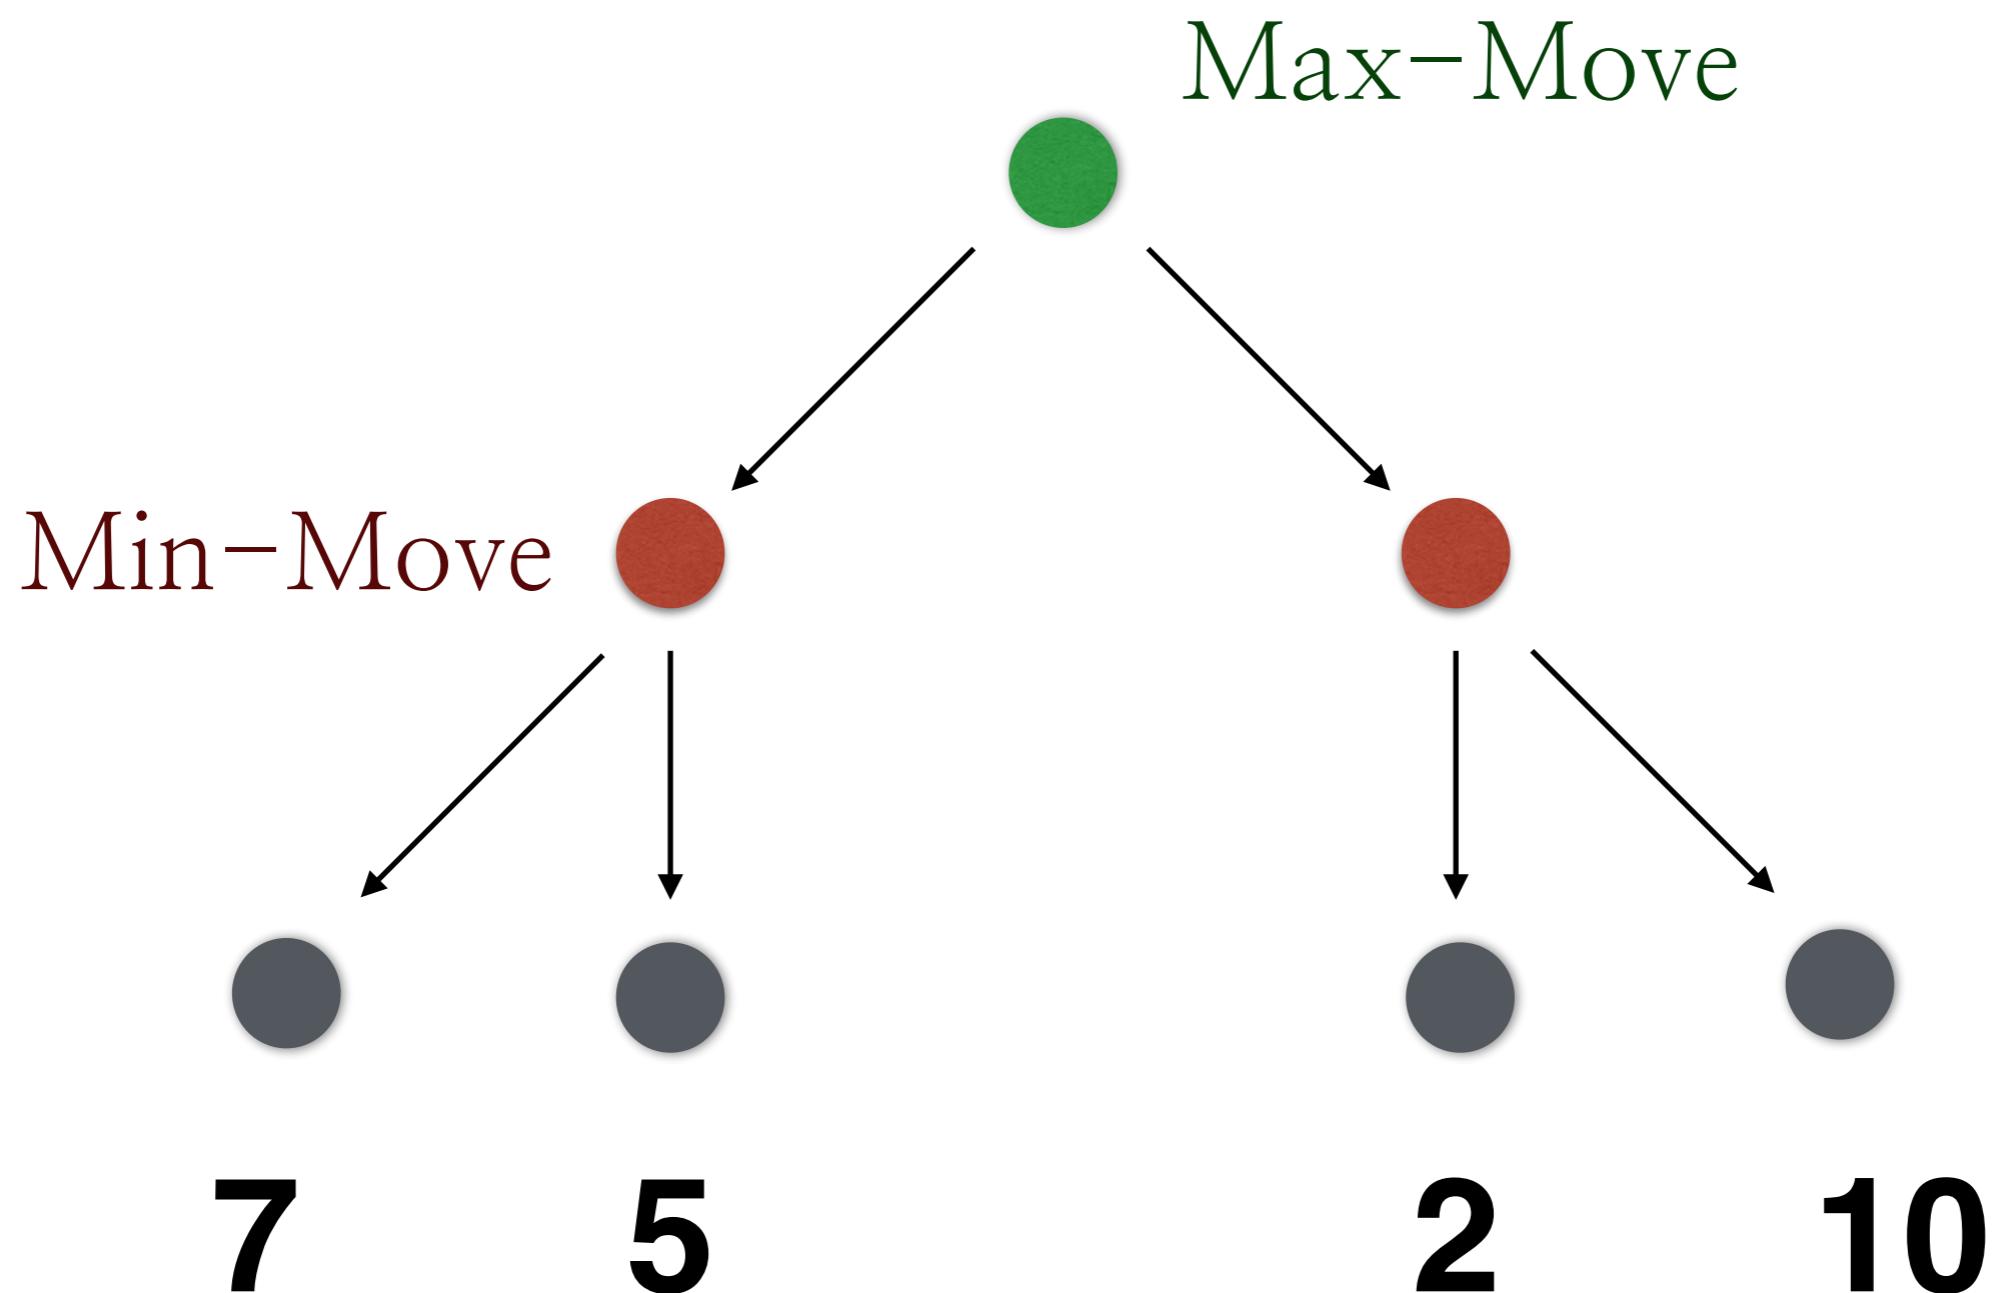
Minimax(node):

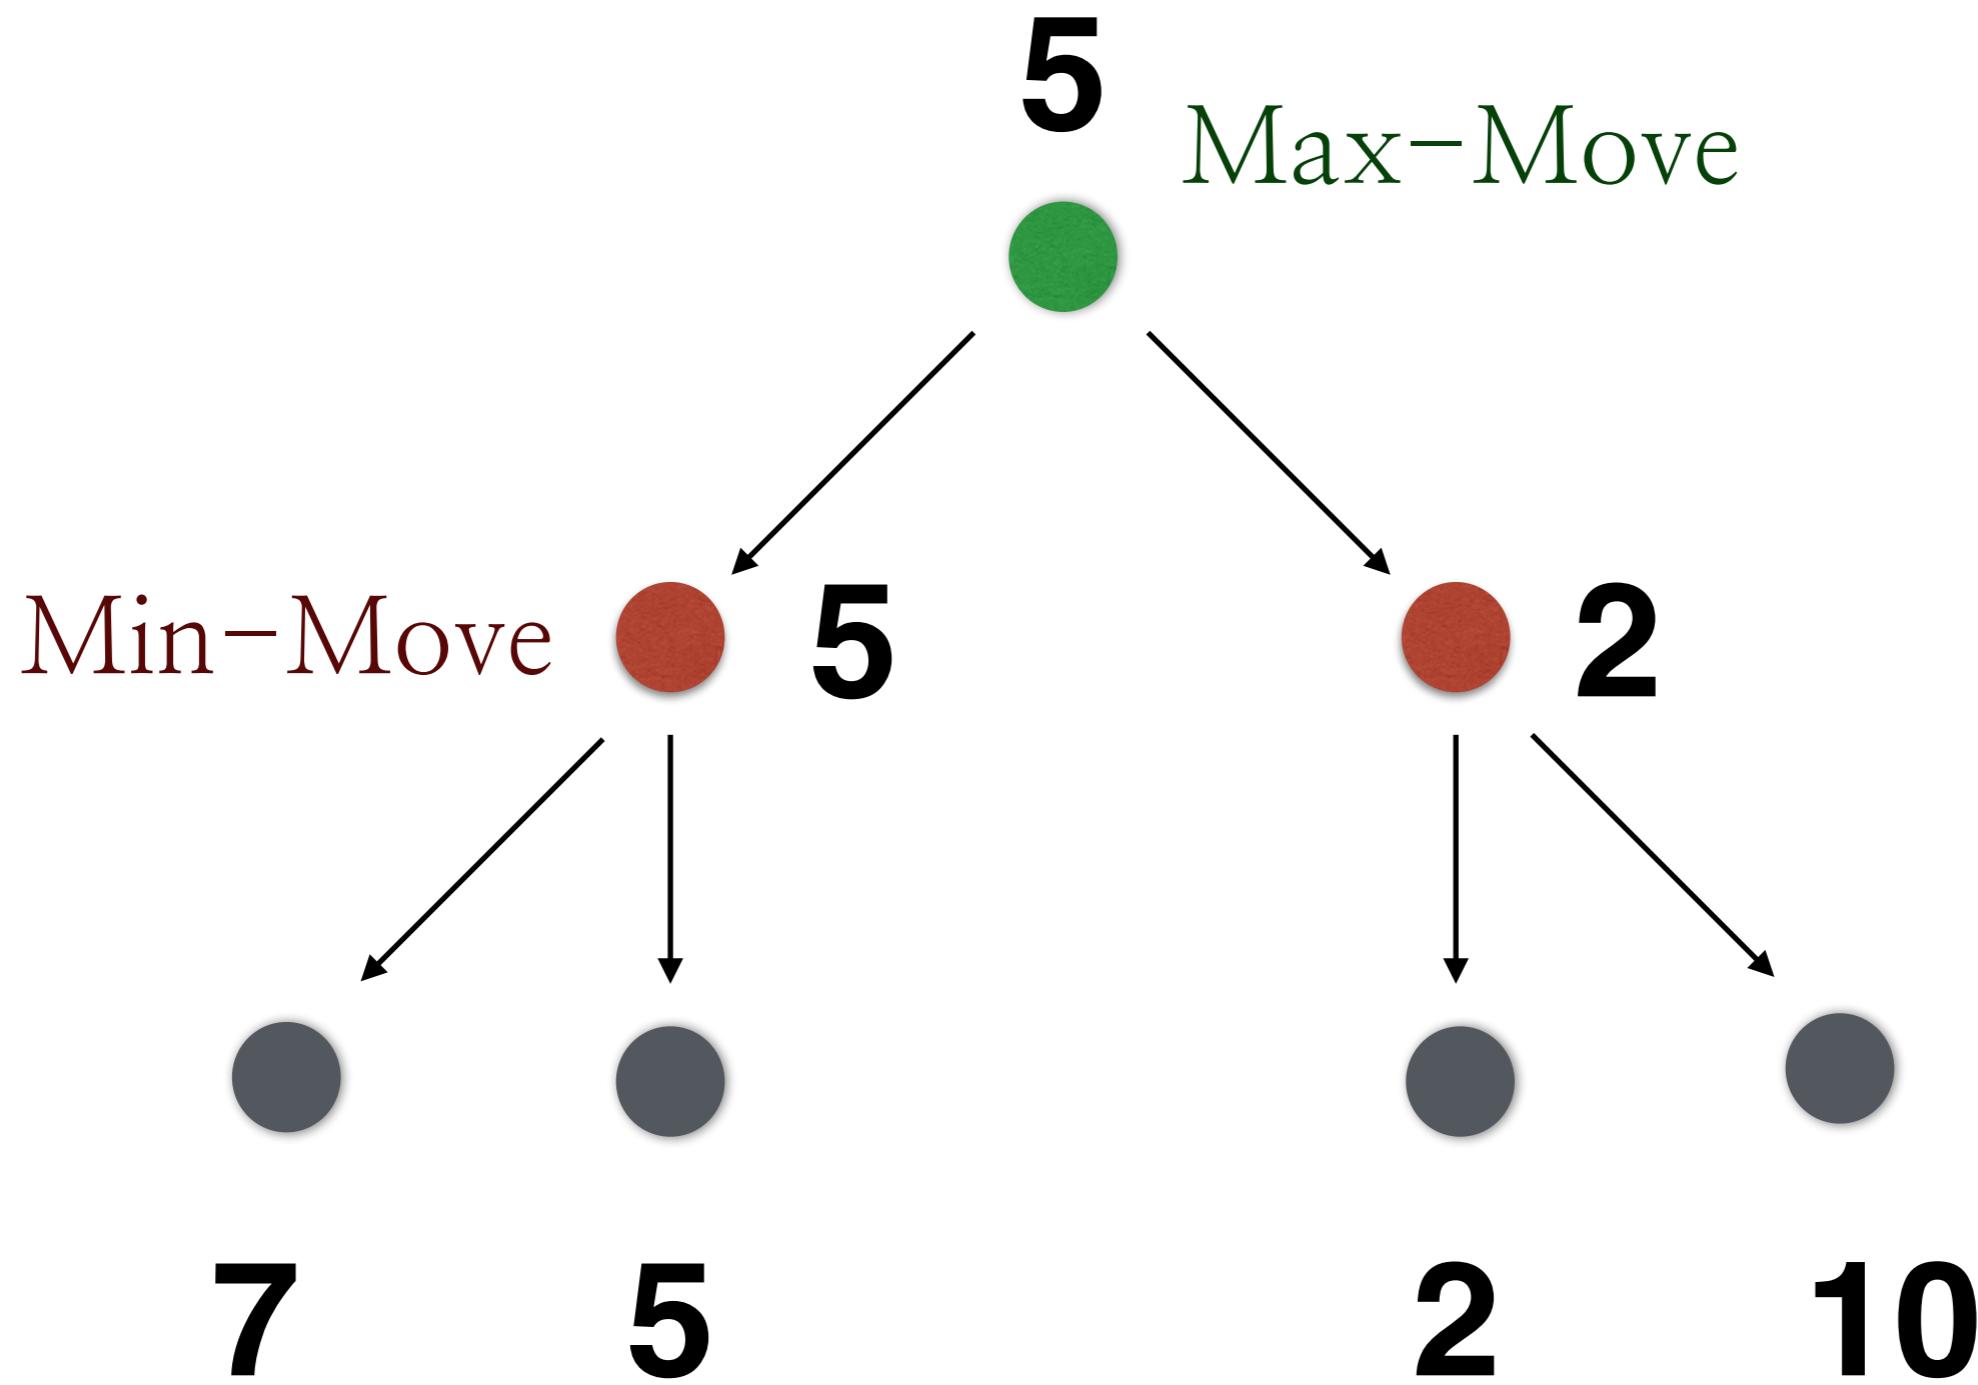
The optimal outcome that the max player can achieve by going into the subtree of that node, assuming that each player is perfectly rational.

C wants min, R wants max, C goes first

	C1	C2
R1	7	5
R2	2	10

Moves: (C1, R1)





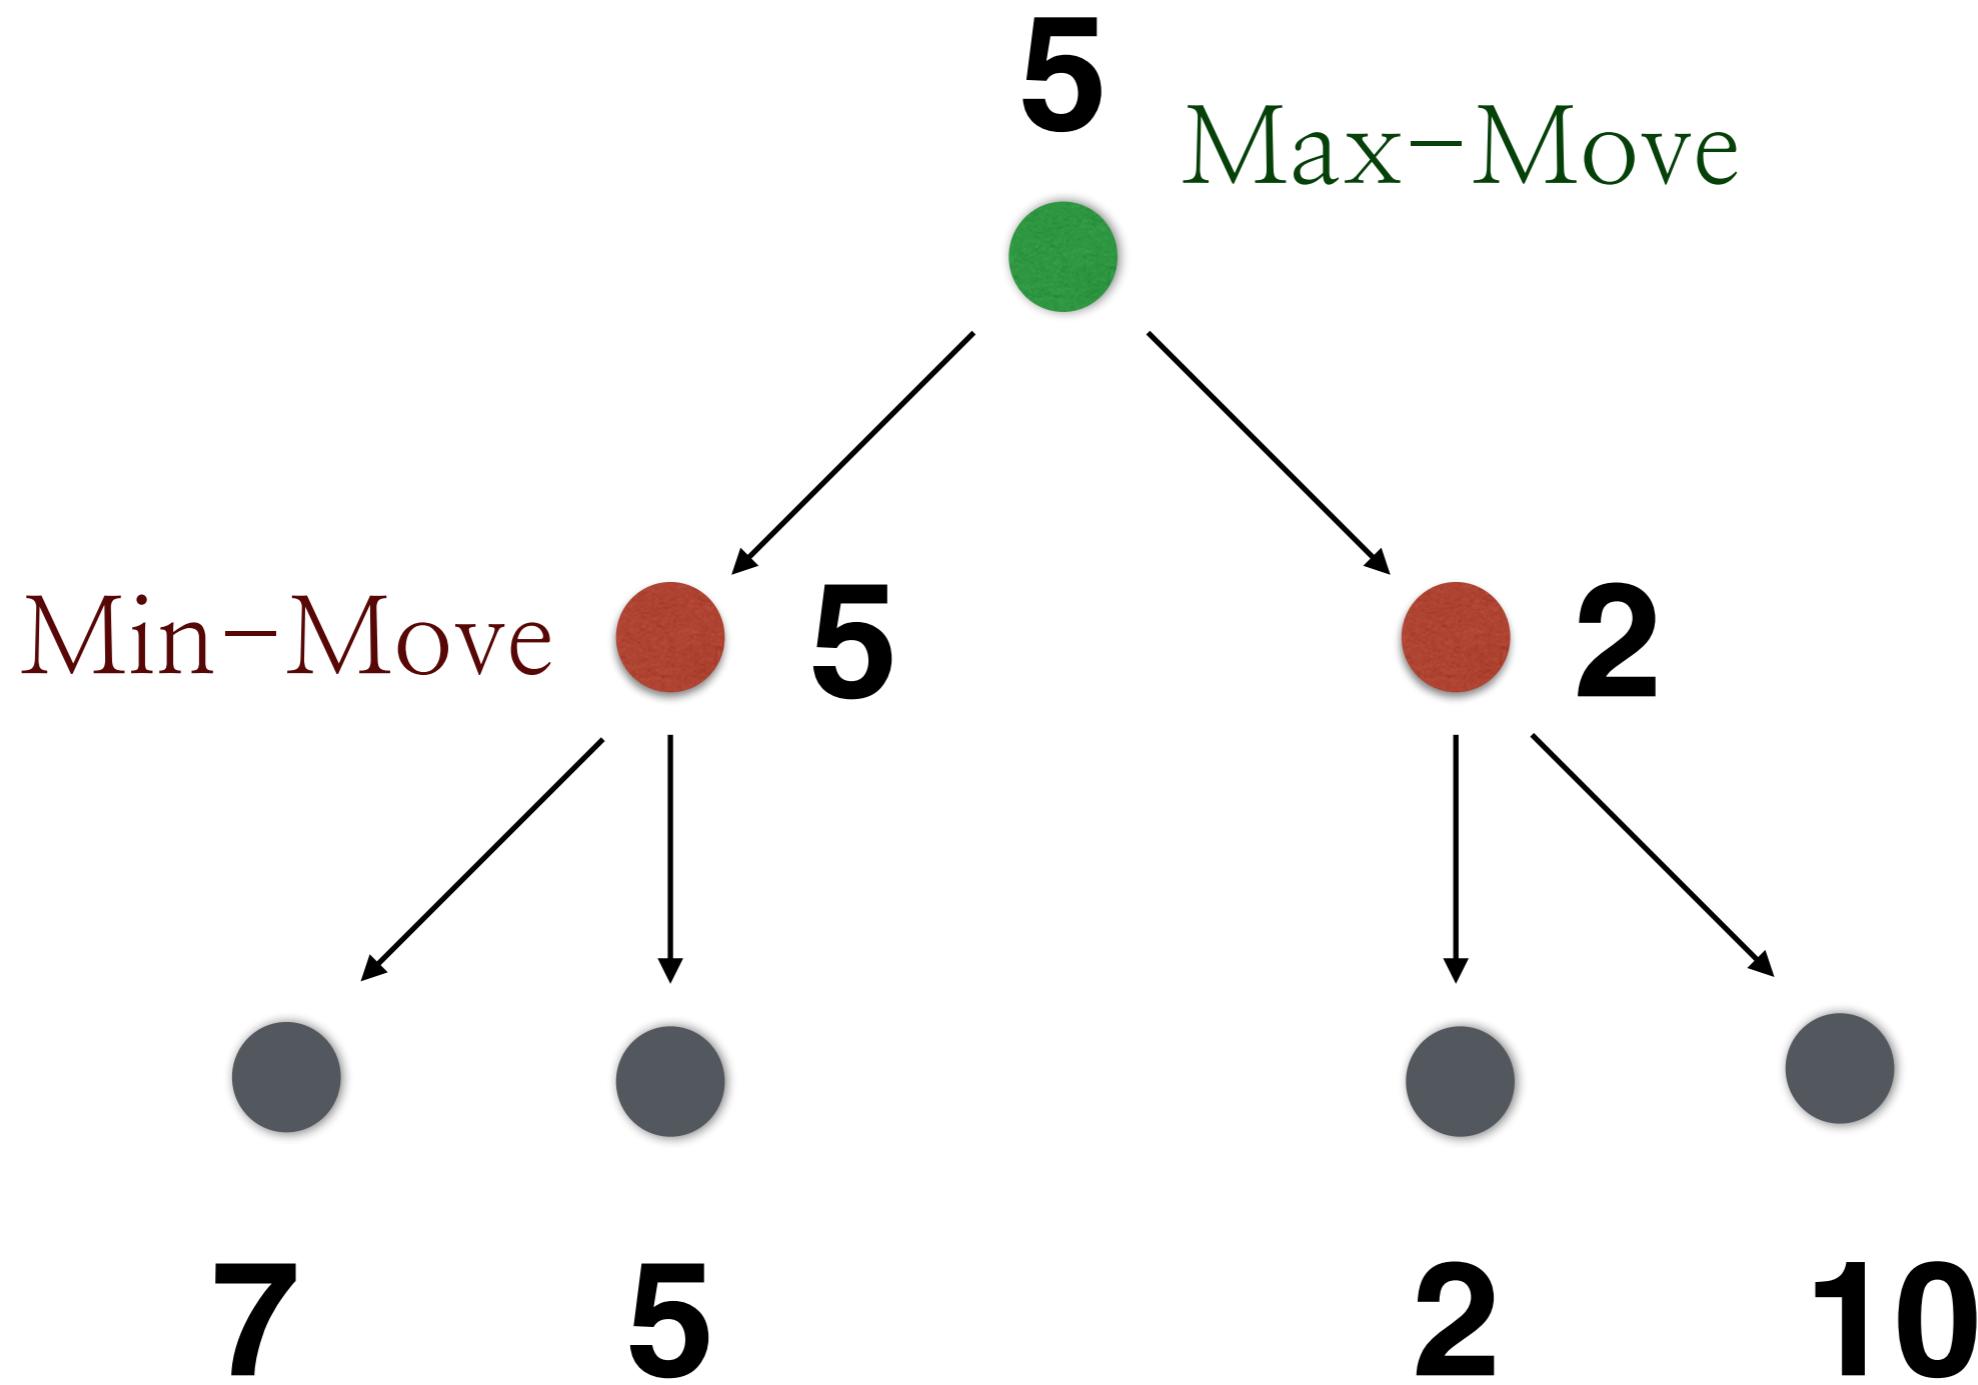
# Minimax Values

Minimax(node):

- If node is terminal, then Payoff(node)
- If node is controlled by the Max player, then the maximum of all Minimax(sub-node)
- If node is controlled by the Min player, then the minimum of all Minimax(sub-node)

# Minimax Algorithm

```
def minimax(node):
    if terminal(node):
        return payoff(node)
    elif max_player(node):
        value = -infinity
        for n in children(node):
            value = max(value, minimax(n))
        return value
    elif min_player(node):
        value = infinity
        for n in children(node):
            value = min(value, minimax(n))
        return value
    else:
        error
```



# Minimax Algorithm

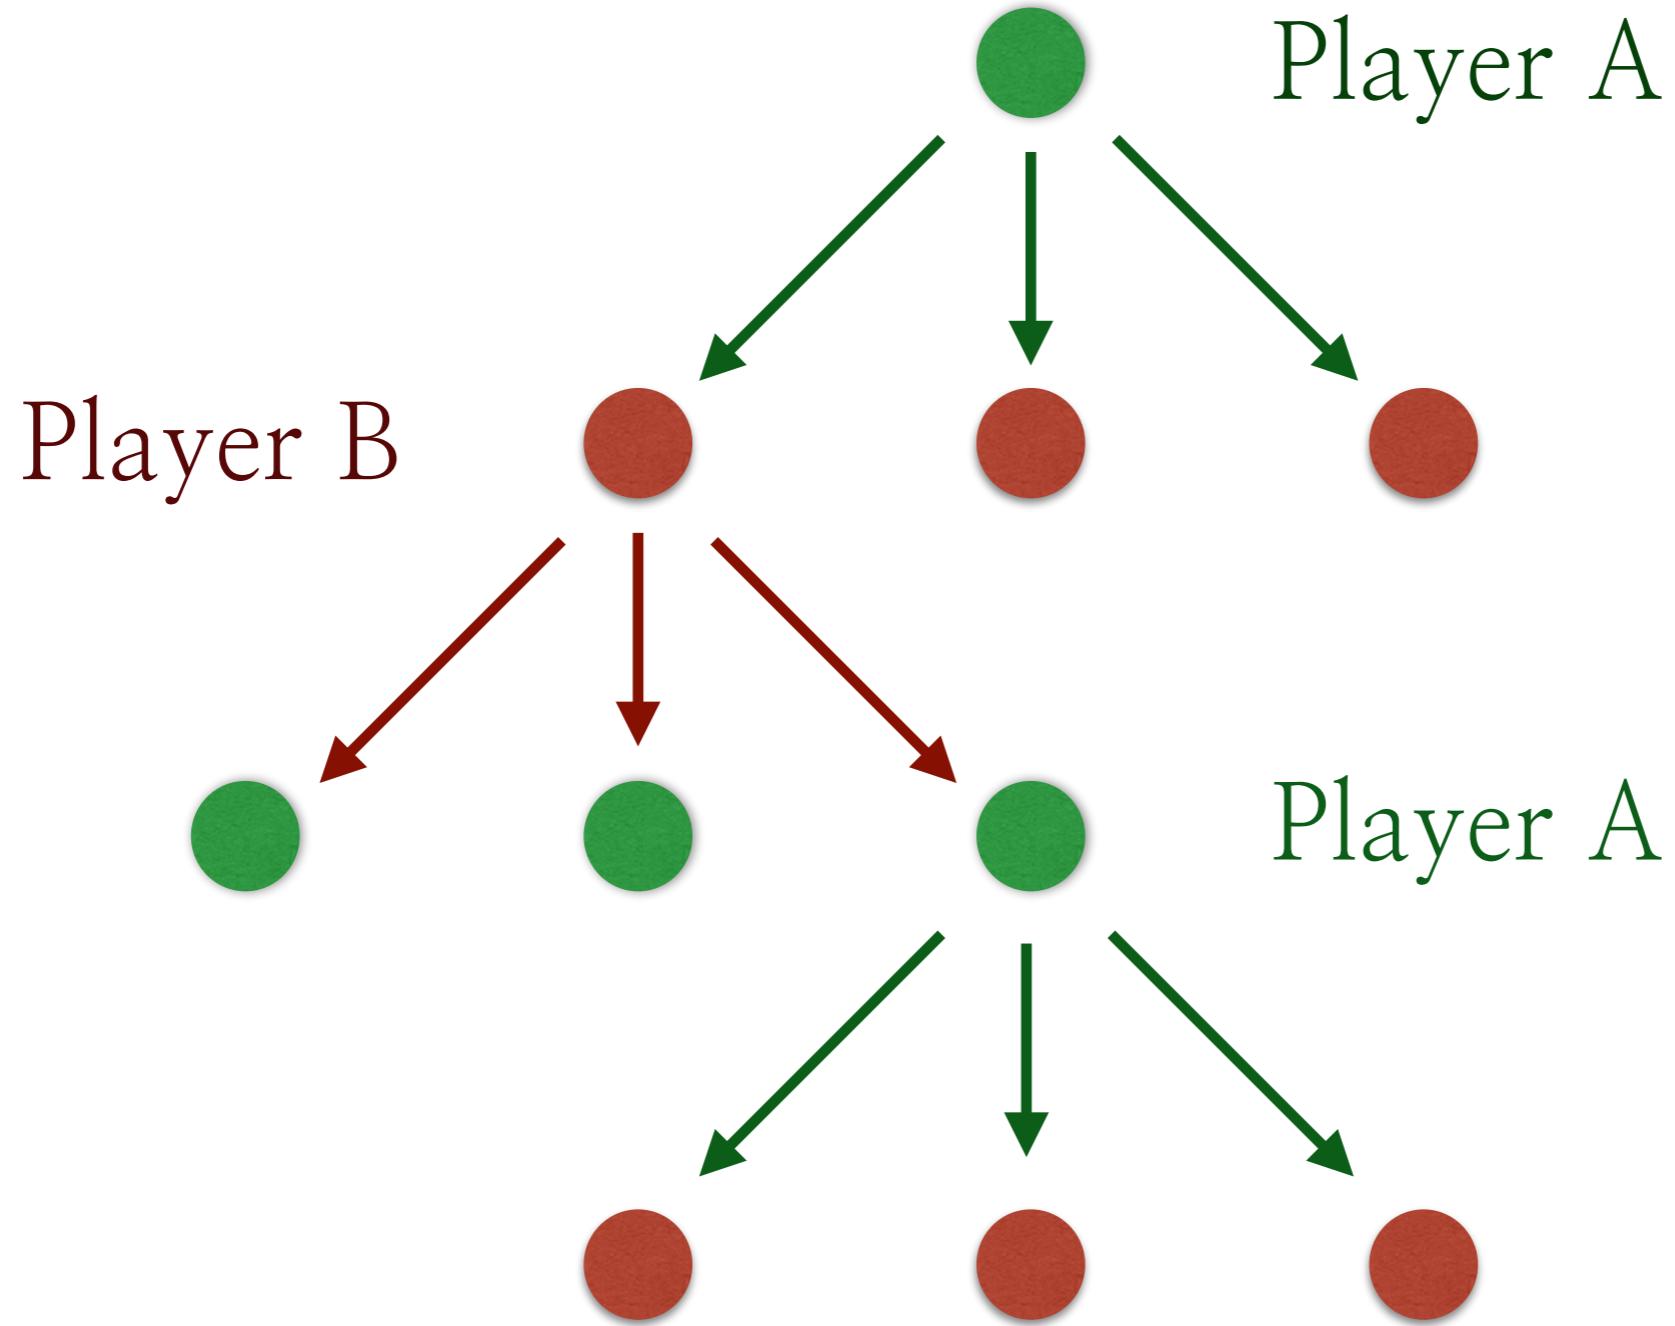
Remember:

We compute values for all these trees just to make **one decision**: what do we do (as the max player) at the root?

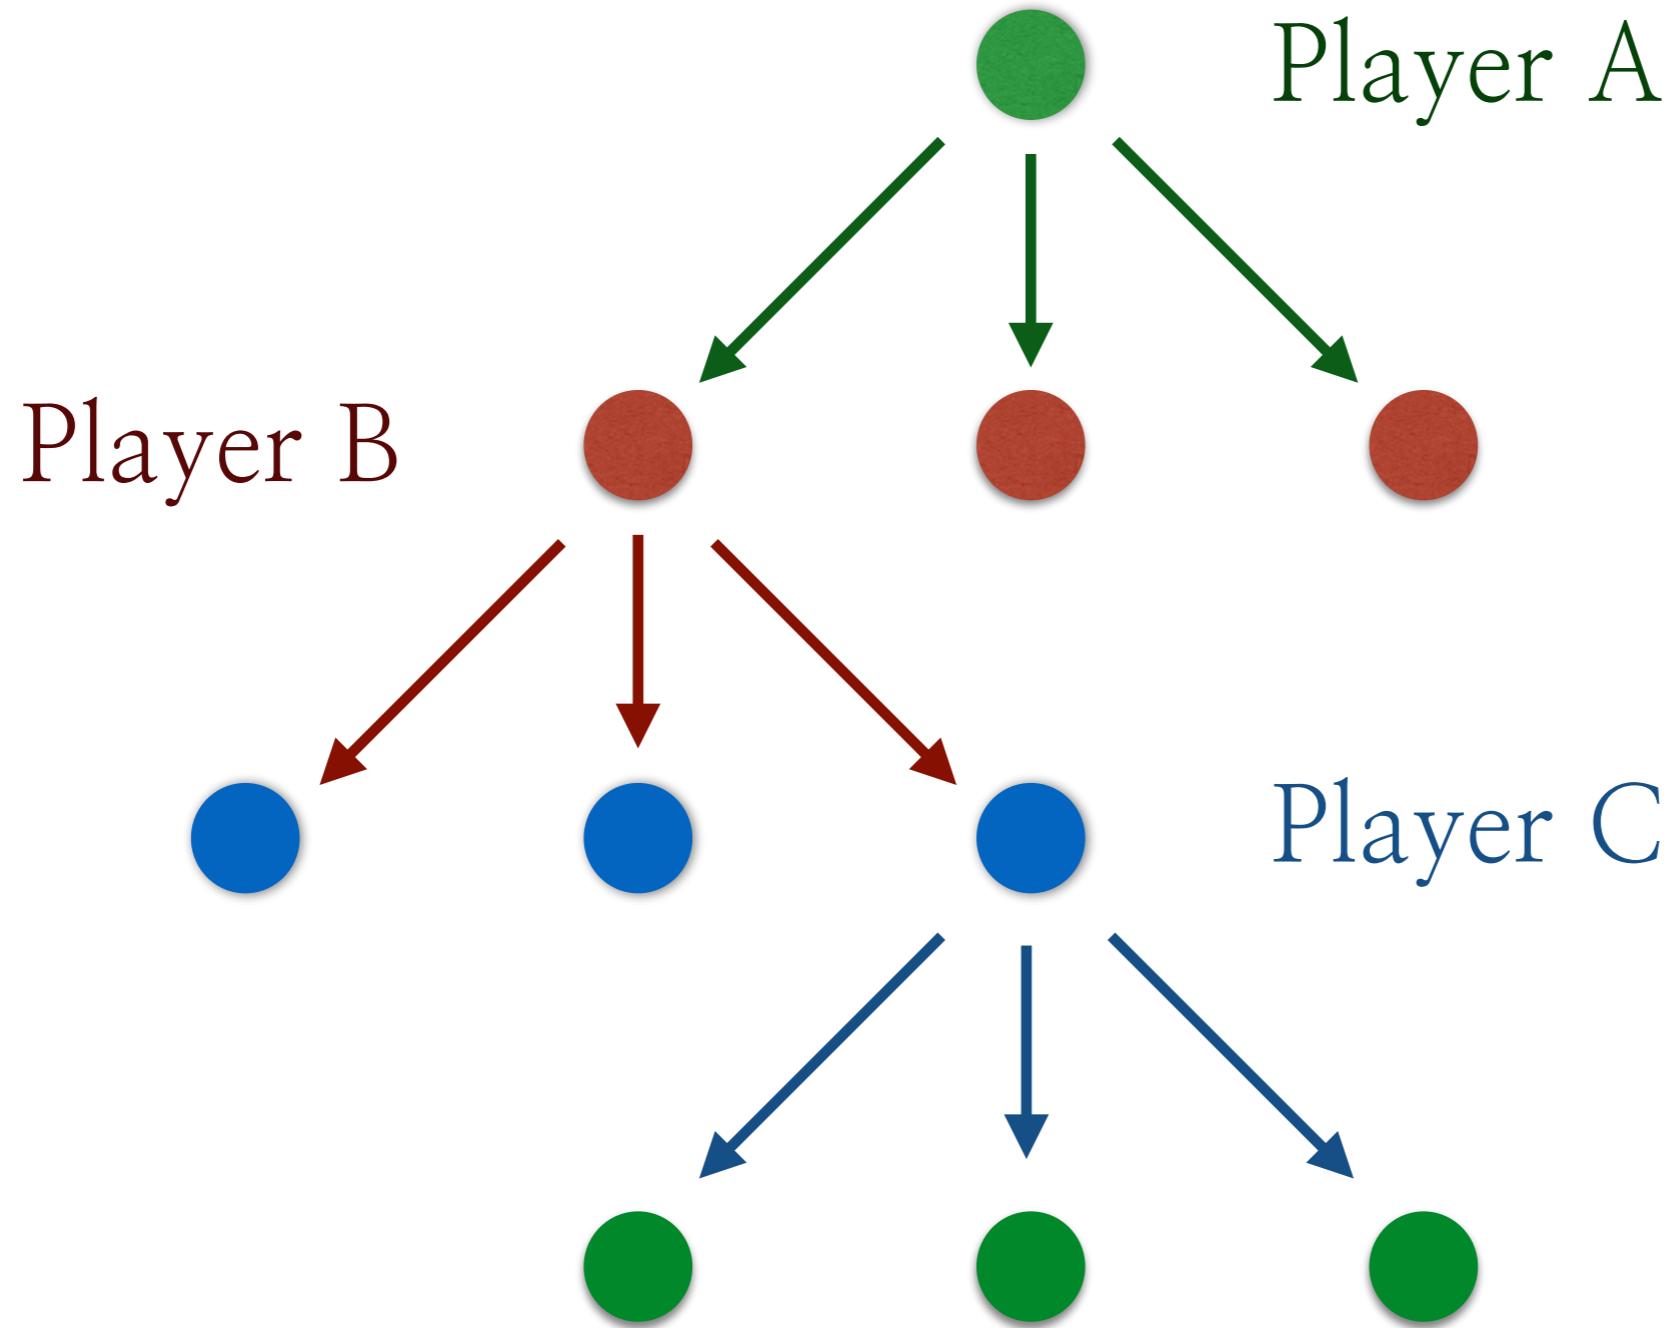
# Minimax Algorithm

- DFS? BFS?
- Complexity?
- Tree expansions are just simulations
- You don't need to understand the game!

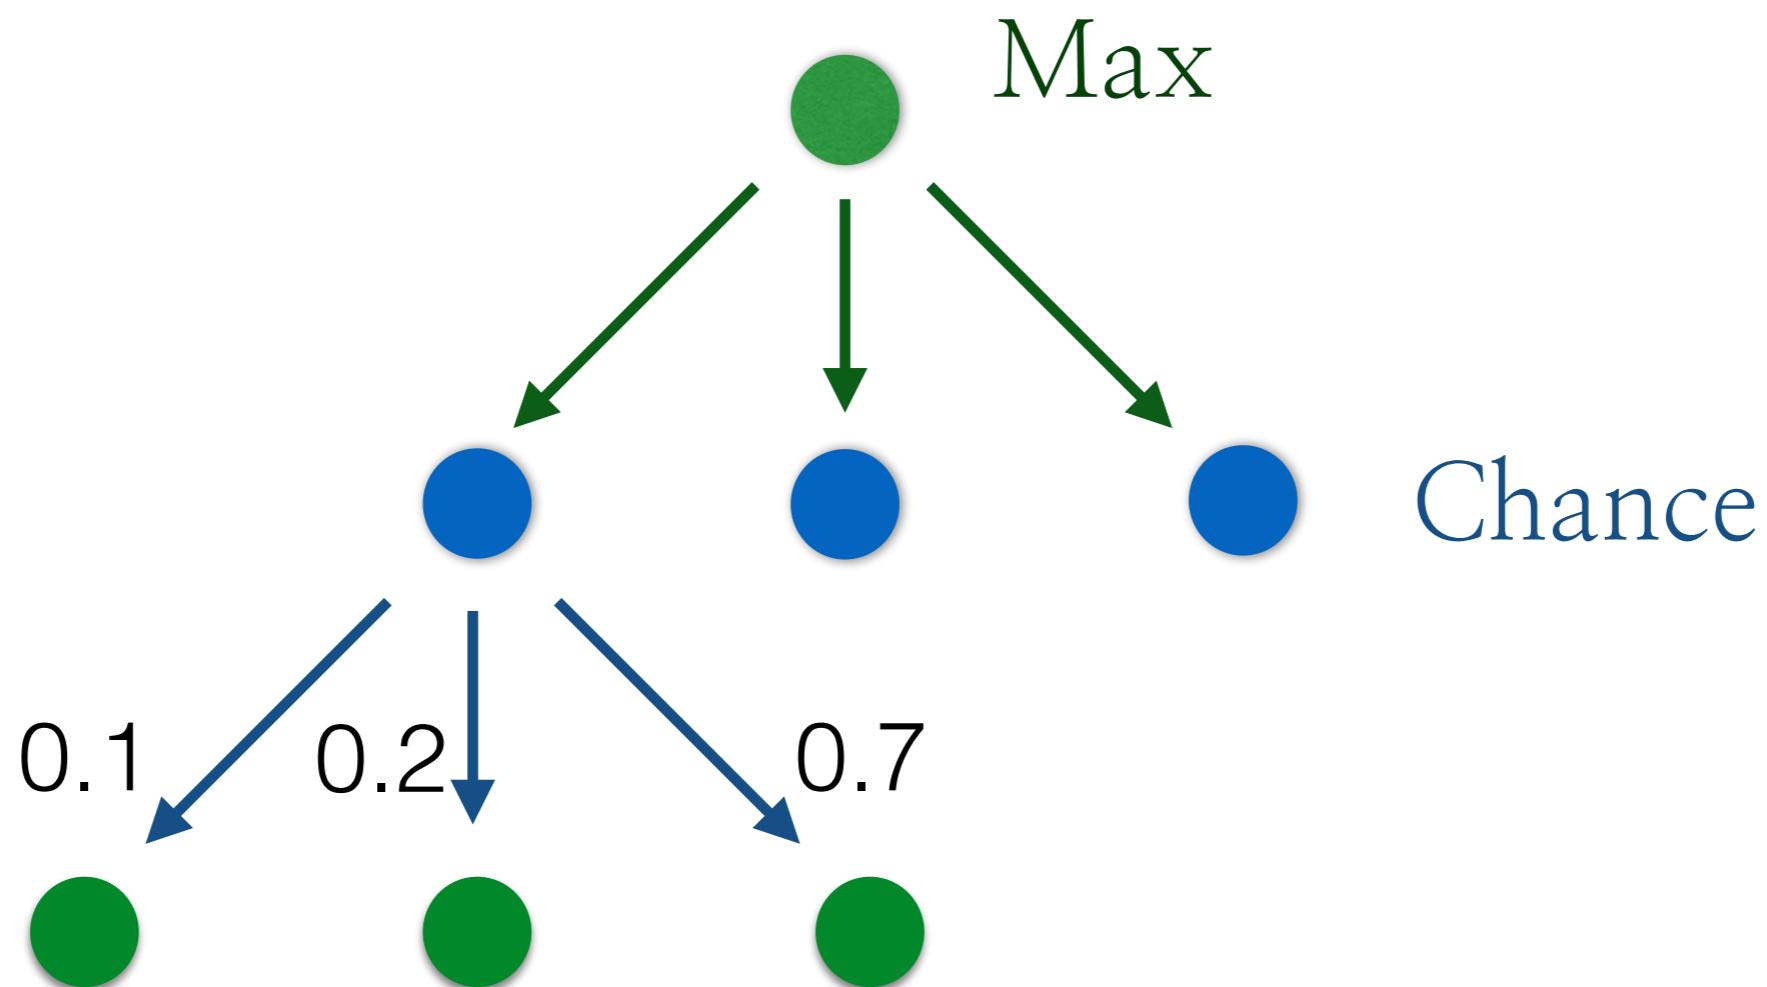
# Extension I: Non Zero-Sum



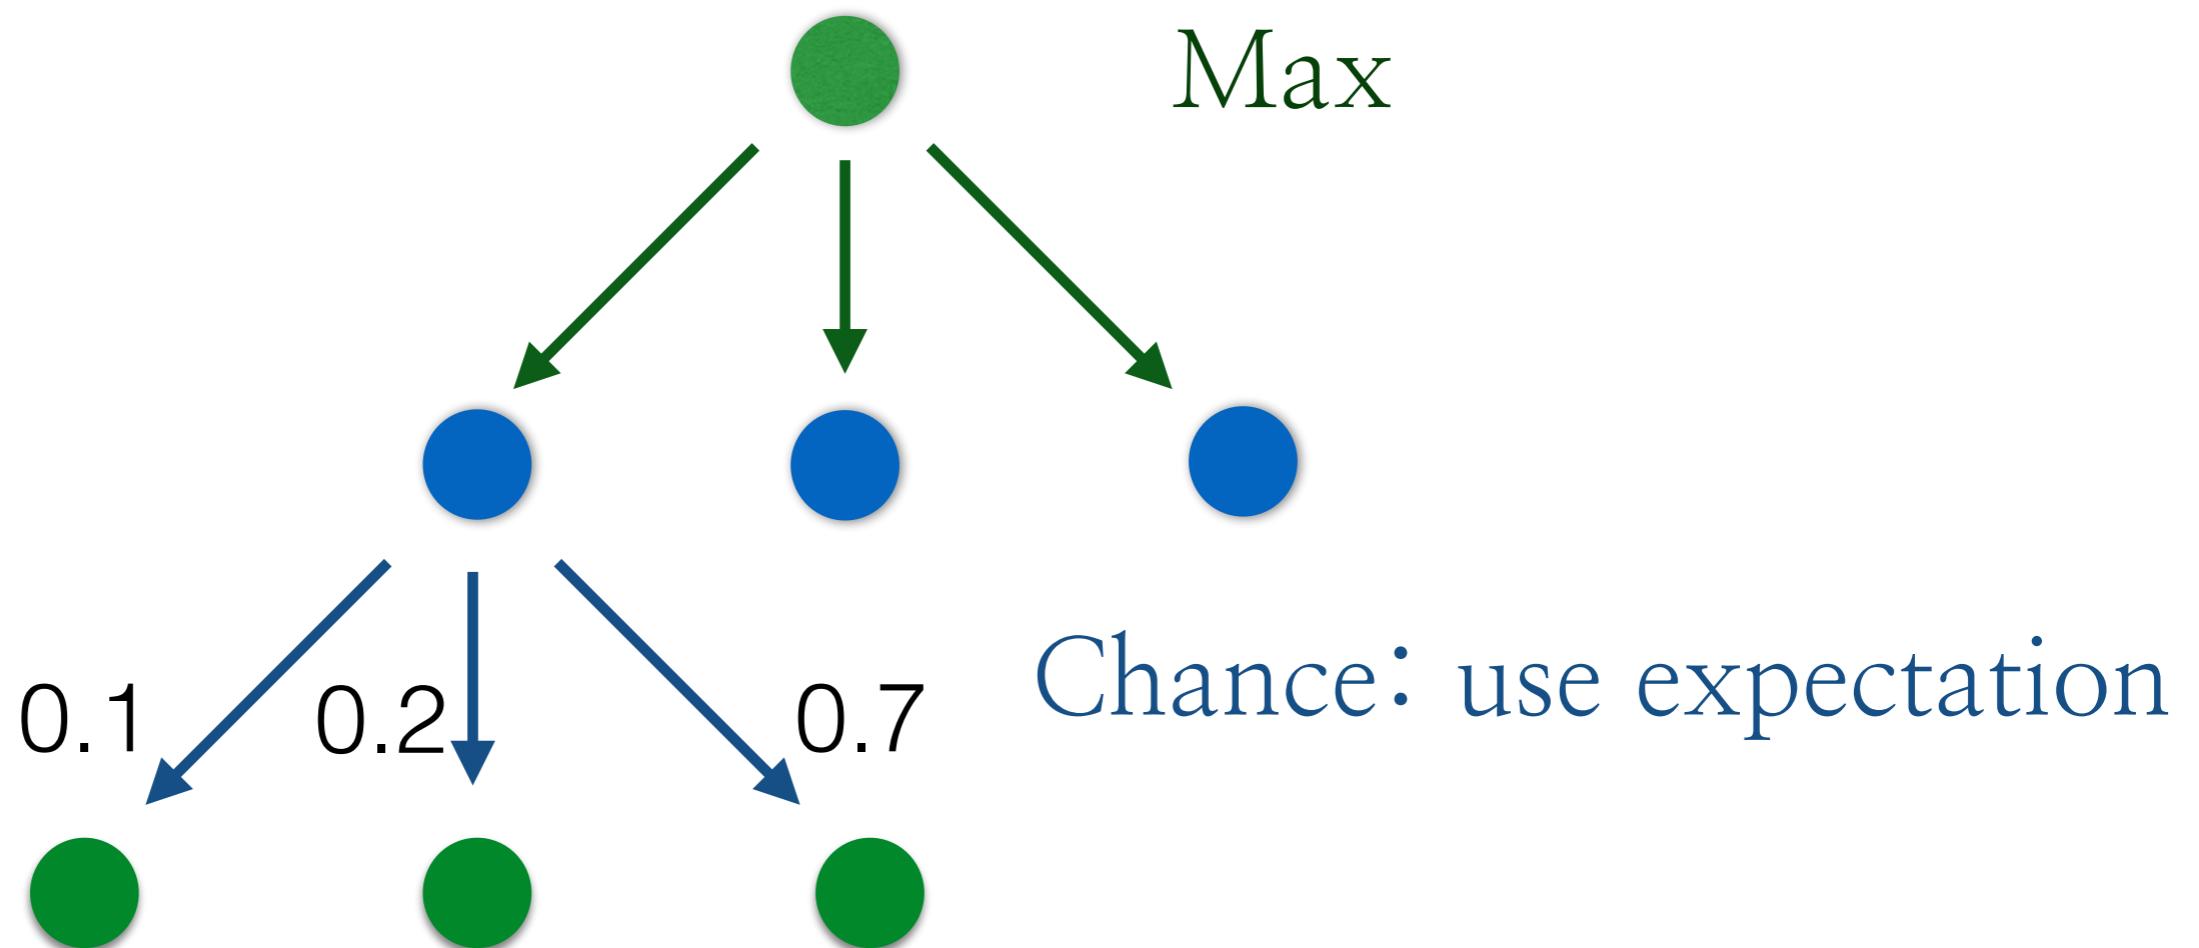
# Extension II: Multiple Players



# Extension III: Stochastic Games



# Expectimax



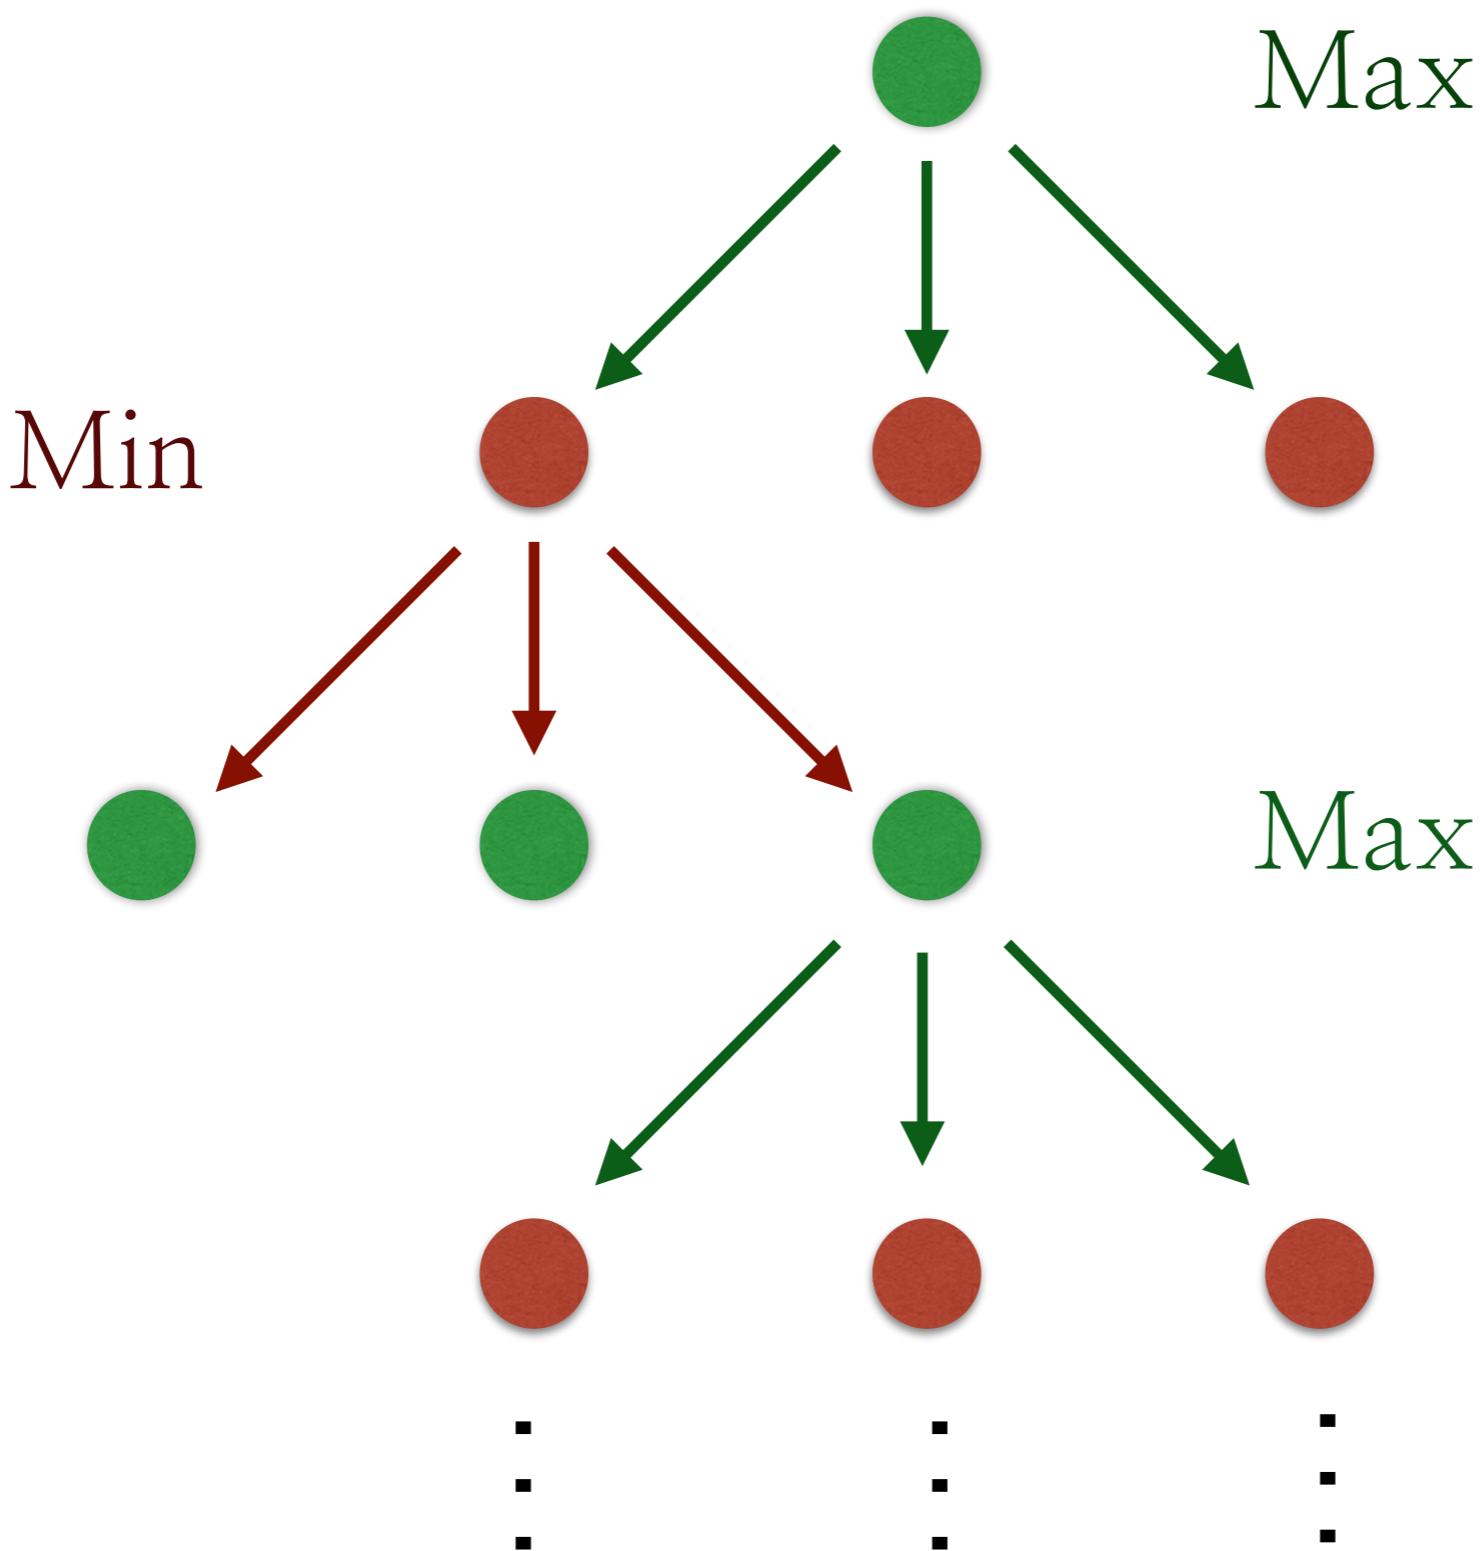
# Expectimax Algorithm

```
def expectimax(node):
    if terminal(node):
        return payoff(node)
    elif max_player(node):
        value = -infinity
        for n in children(node):
            value = max(value, expectimax(n))
        return value
    elif chance_player(node):
        value = 0
        for n in children(node):
            value = value + expectimax(n)*chance(n)
        return value
    else:
        error
```

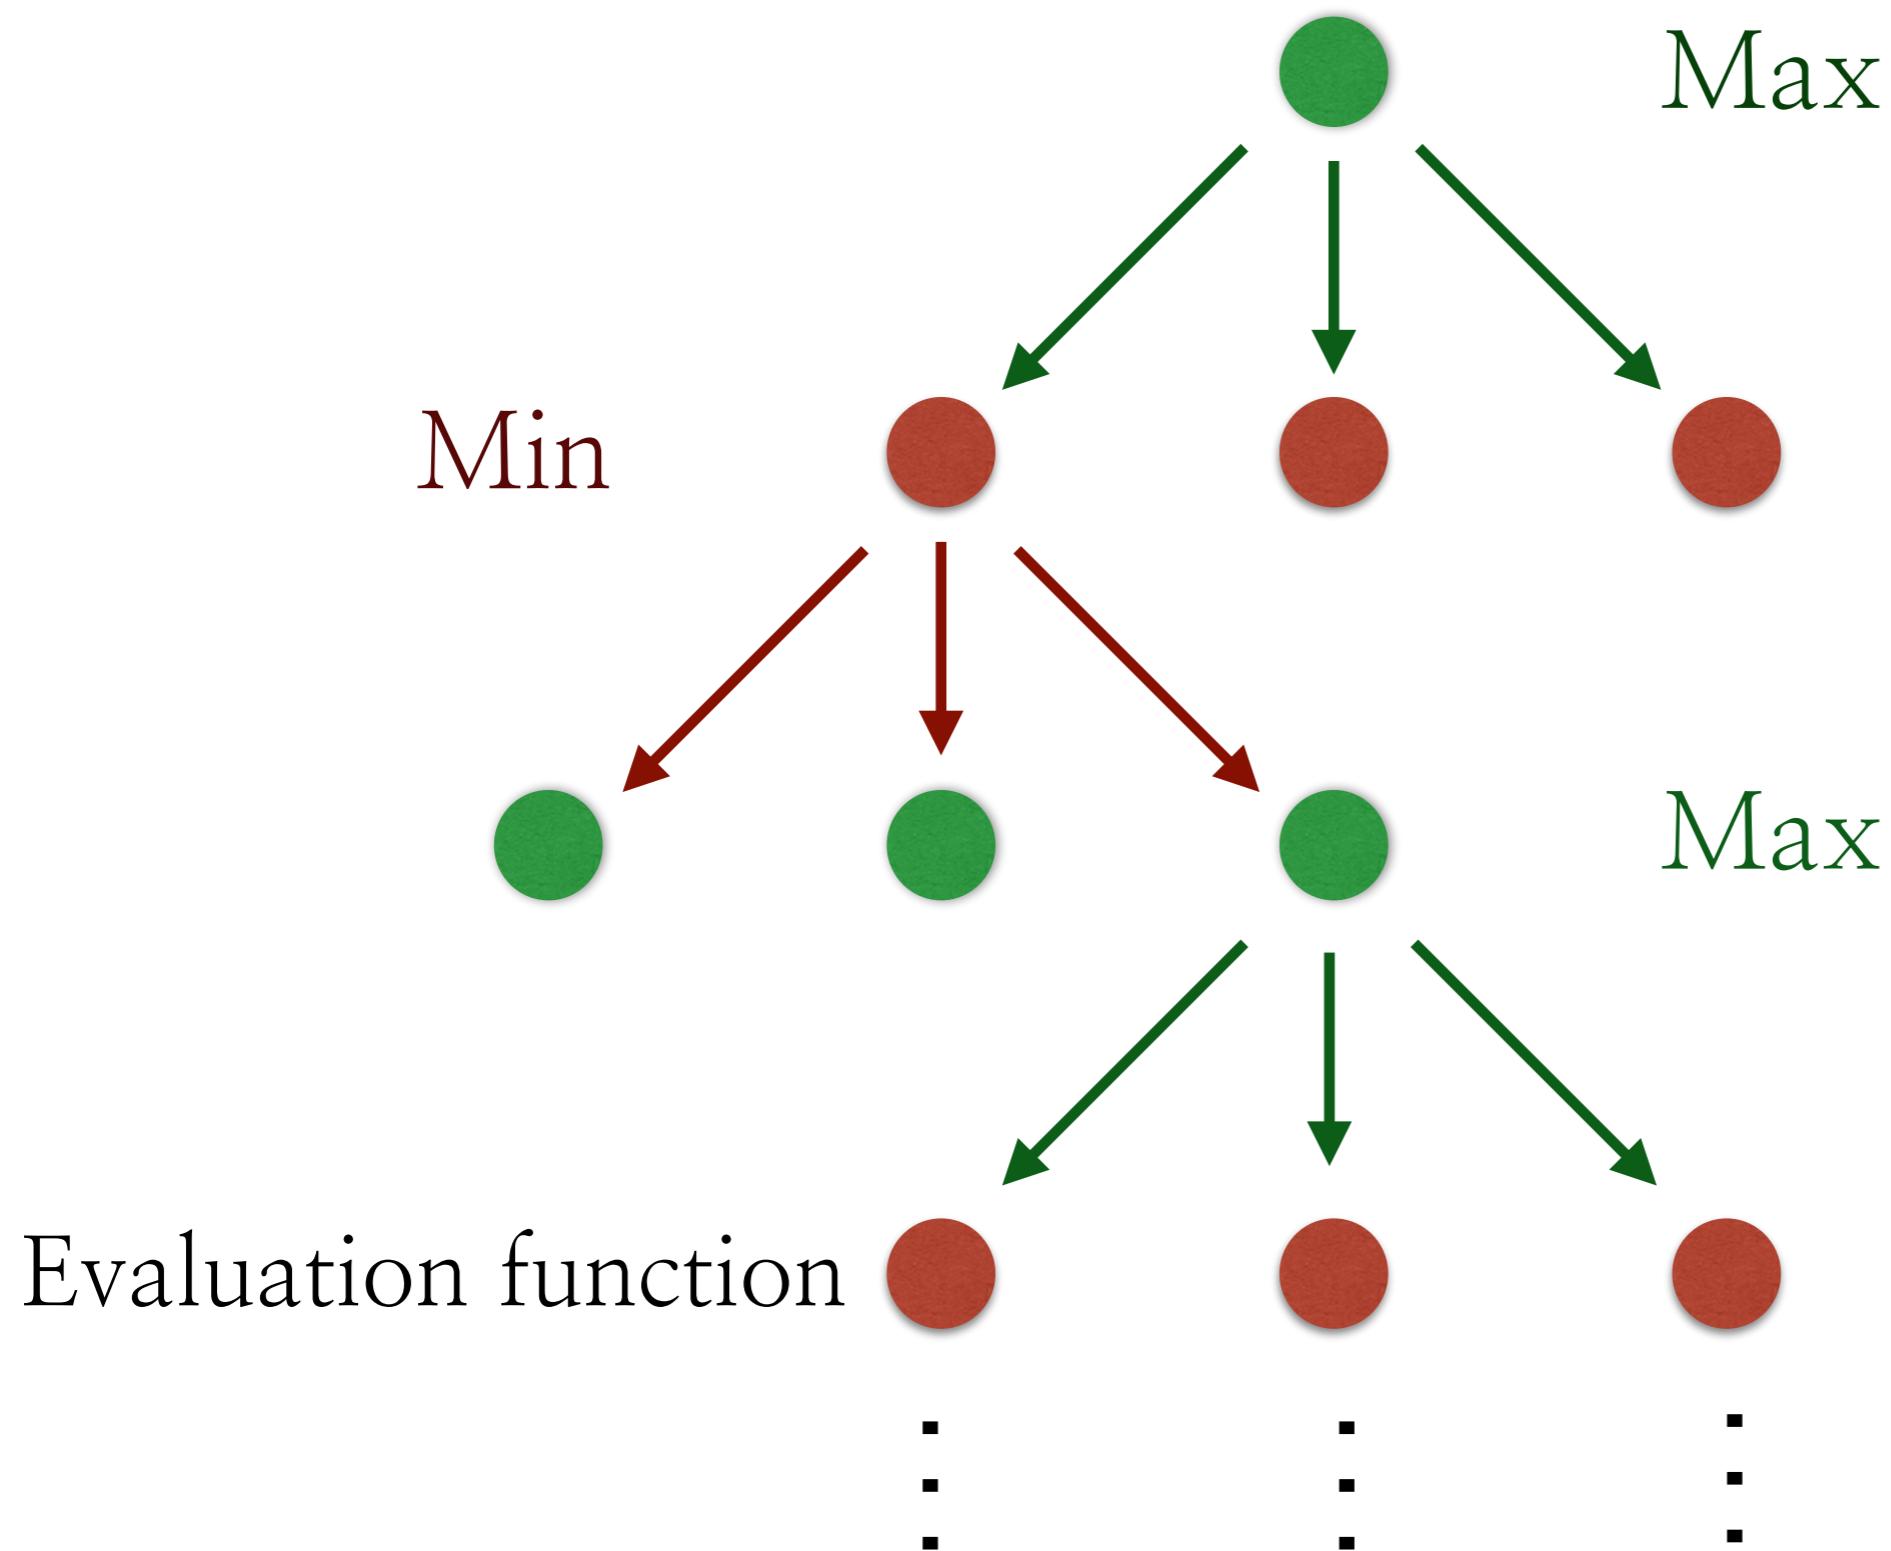
# Trees are too large!

- Vertical Pruning: Cut-off at some fixed depth and use heuristics to evaluate the intermediate nodes.
- Horizontal Pruning: Alpha-Beta Pruning

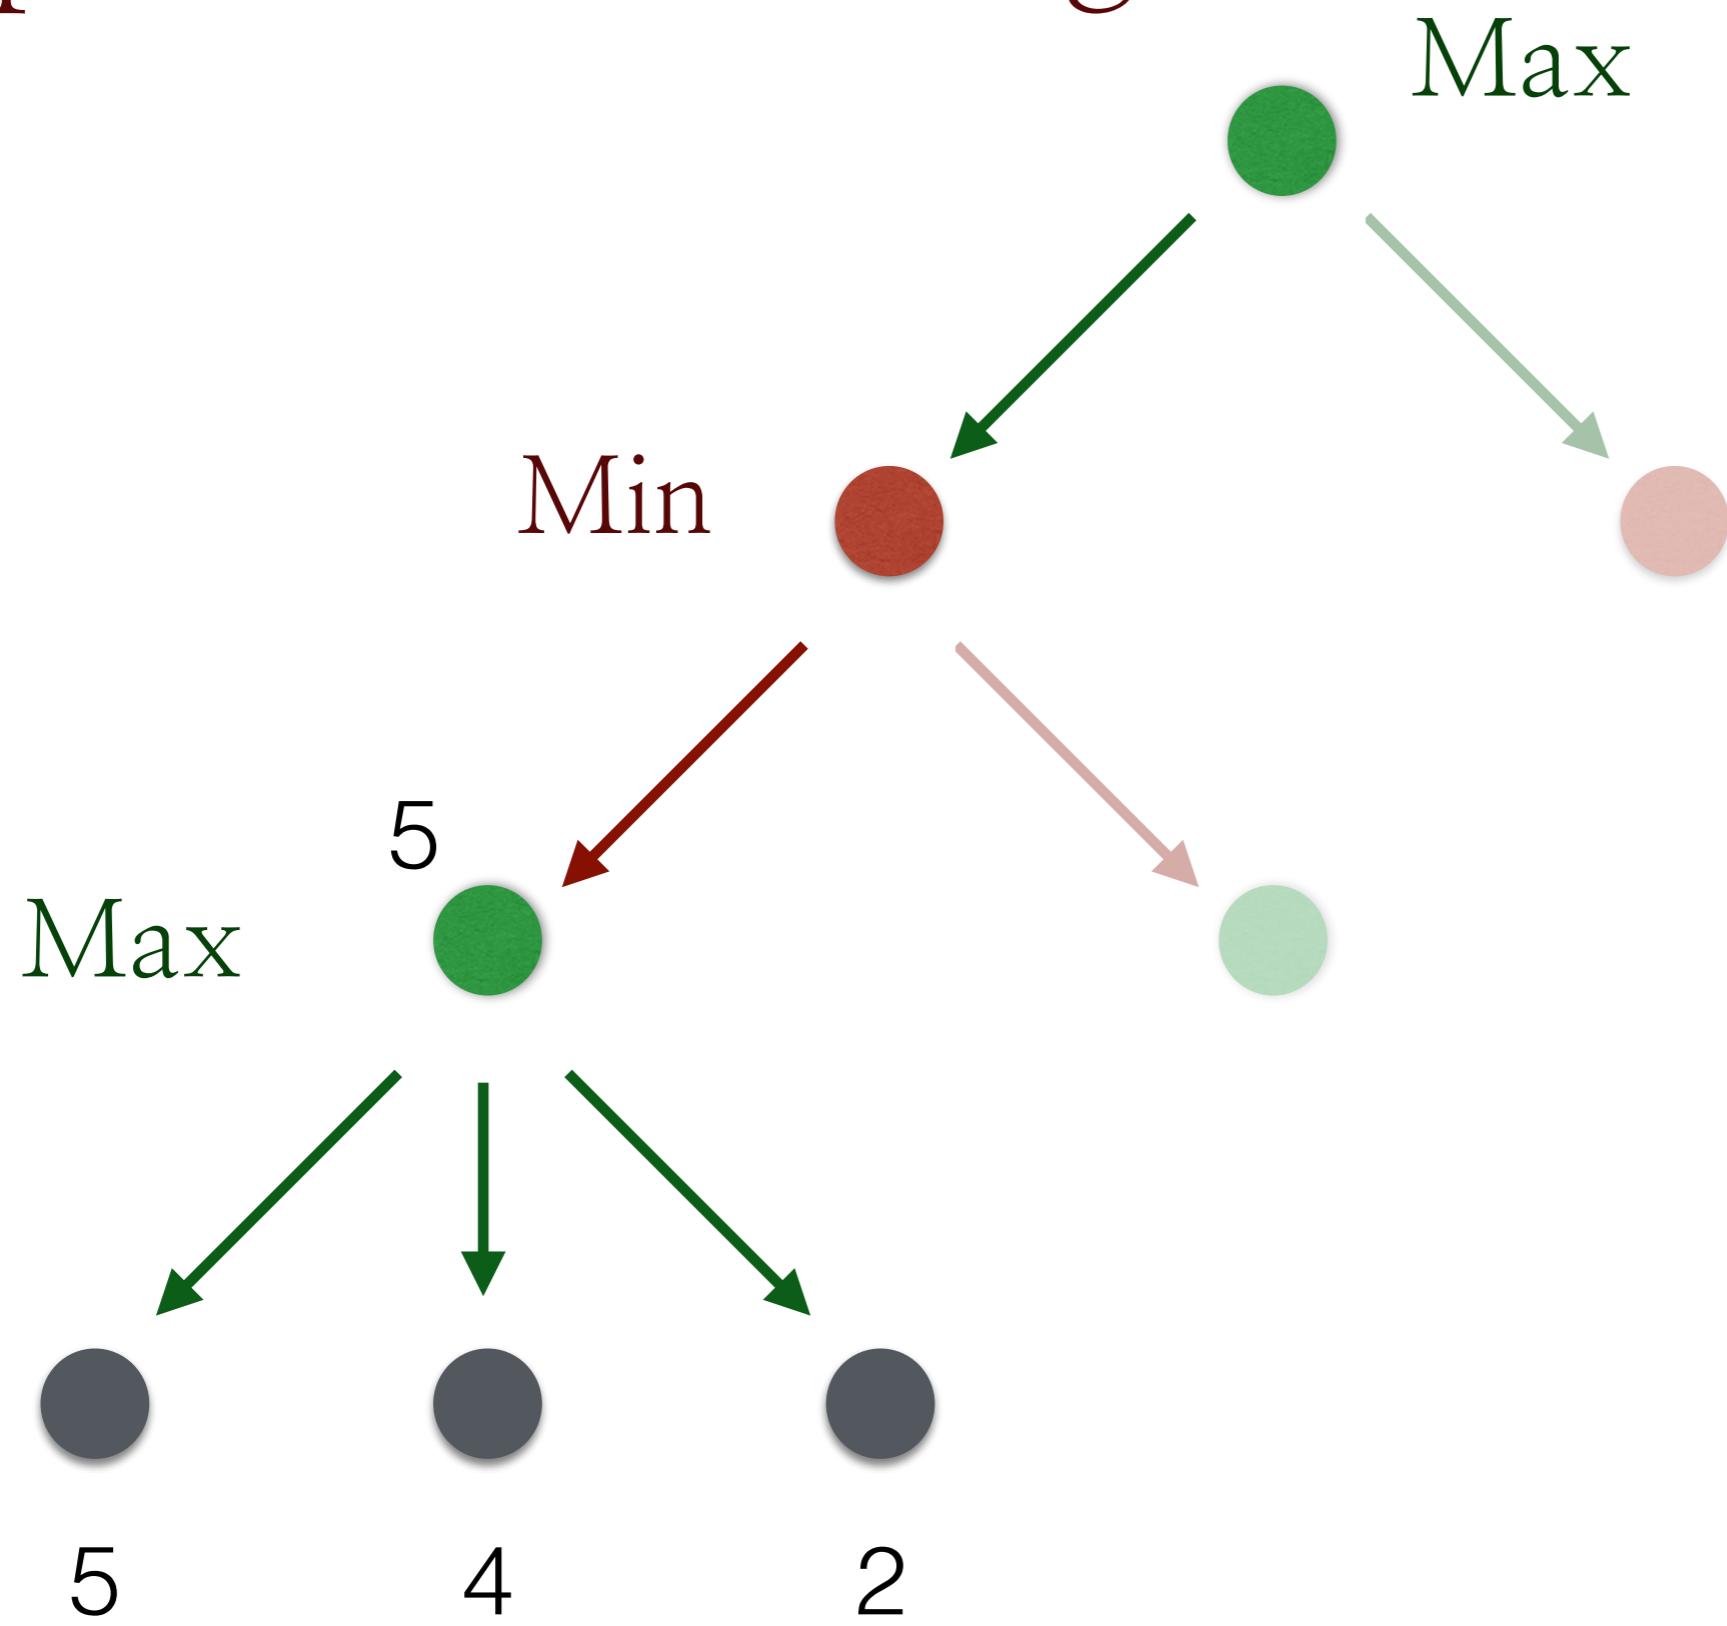
# Vertical Cut-off



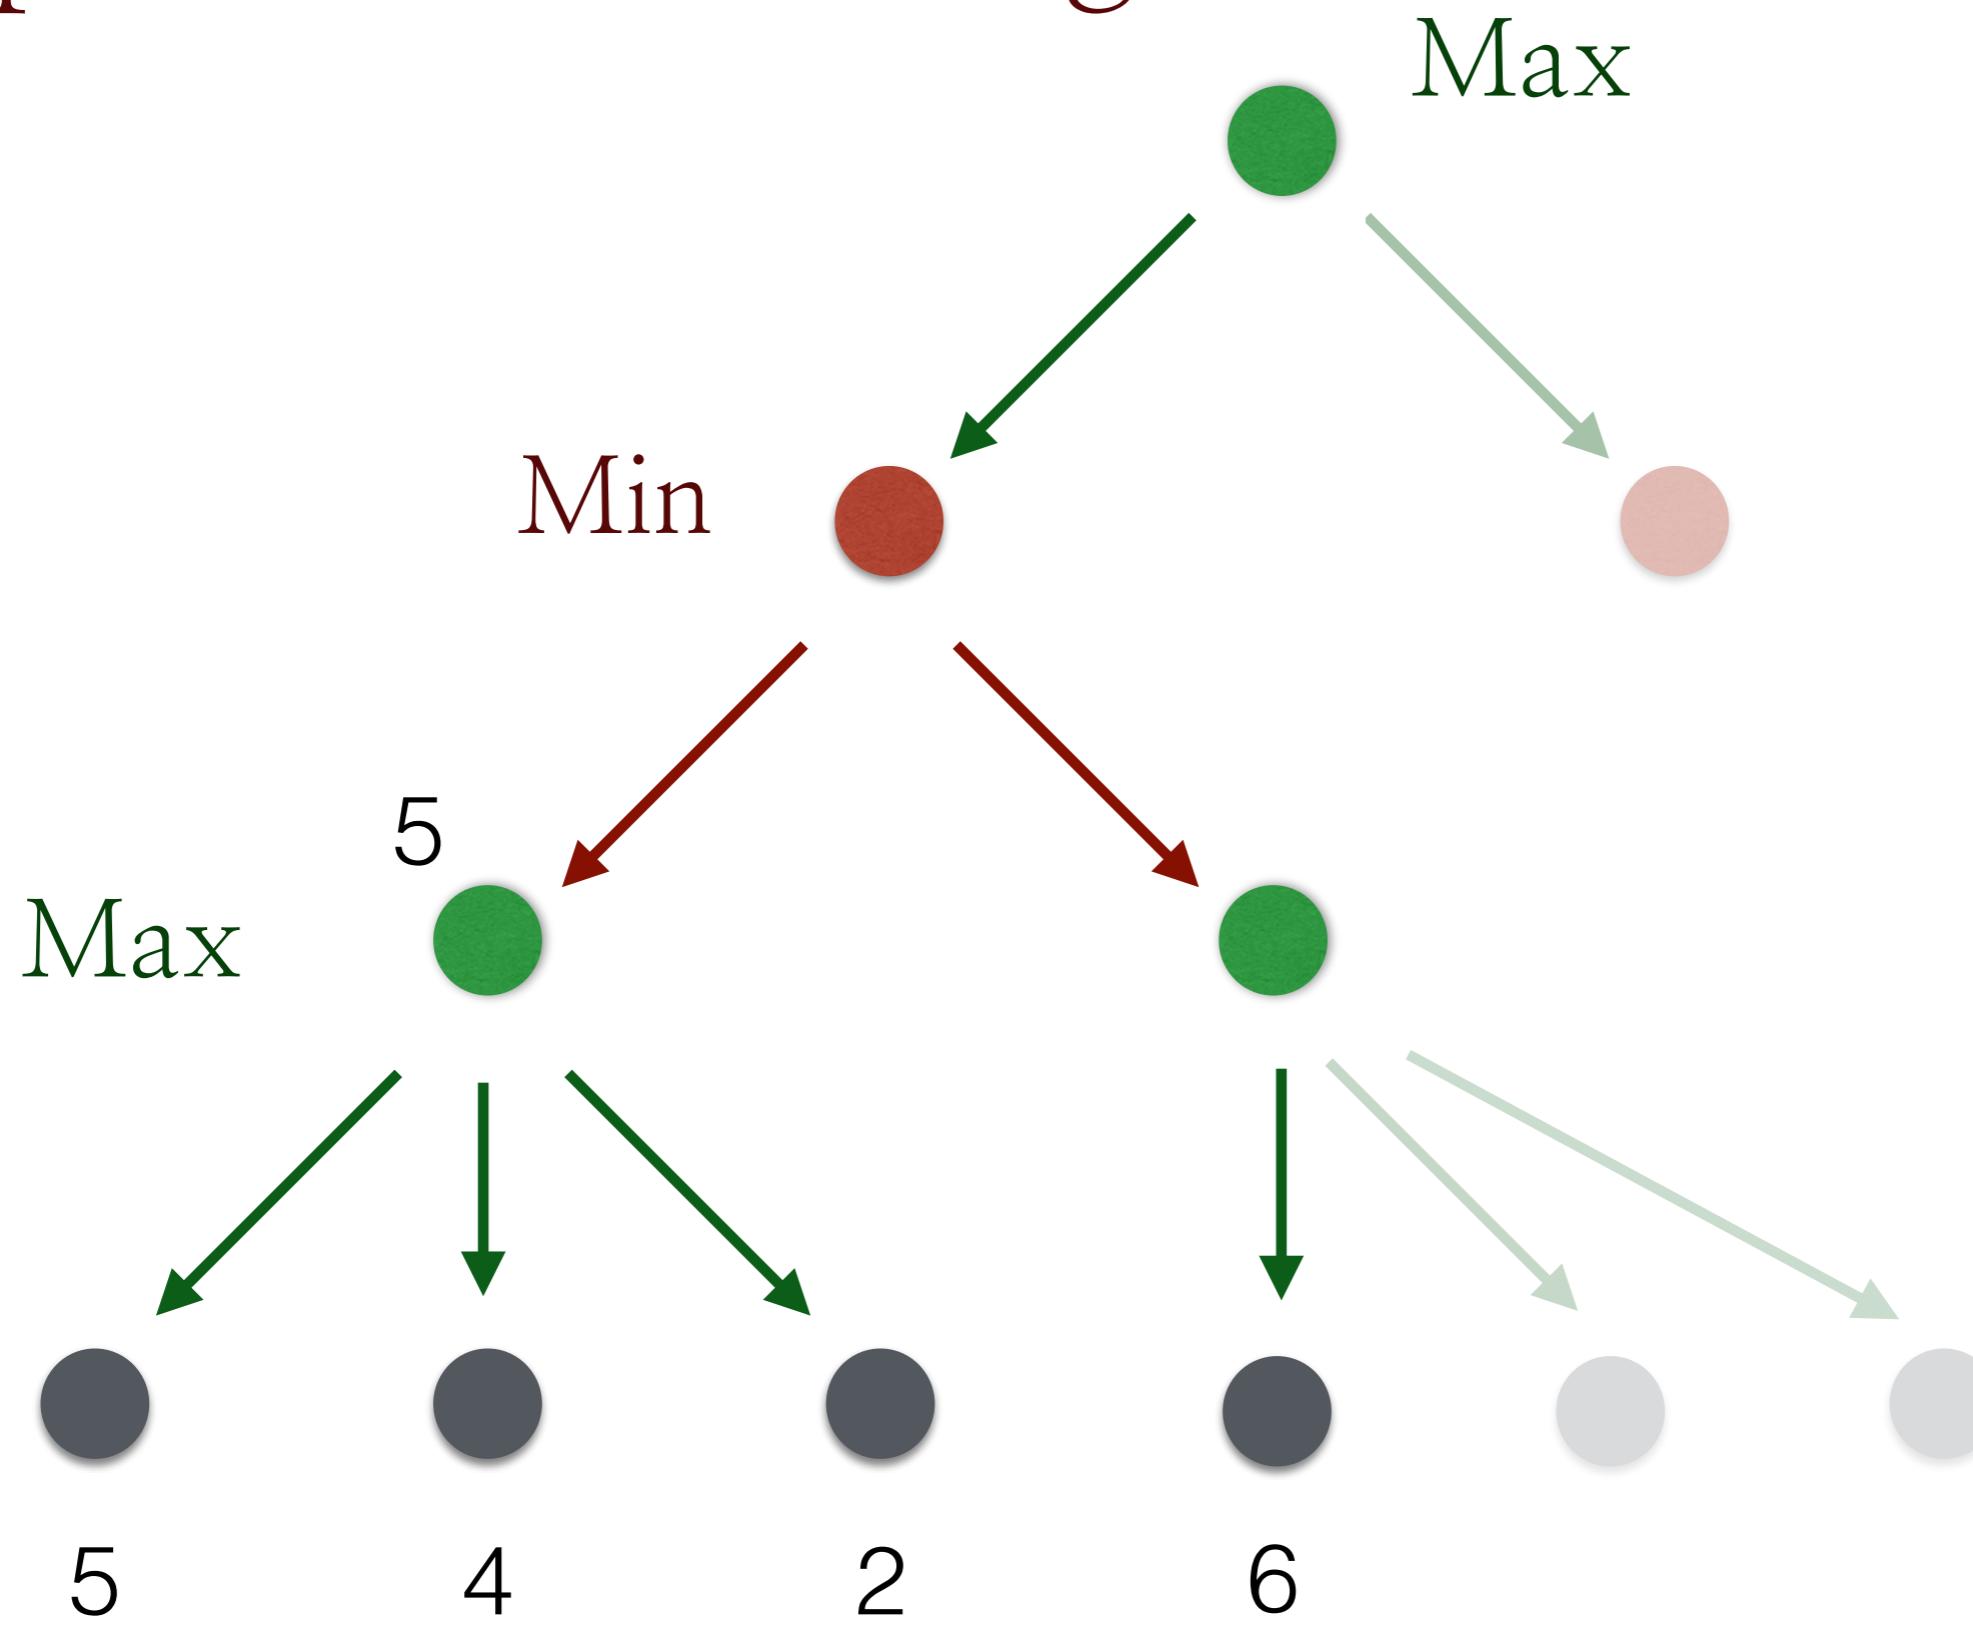
# Vertical Cut-off



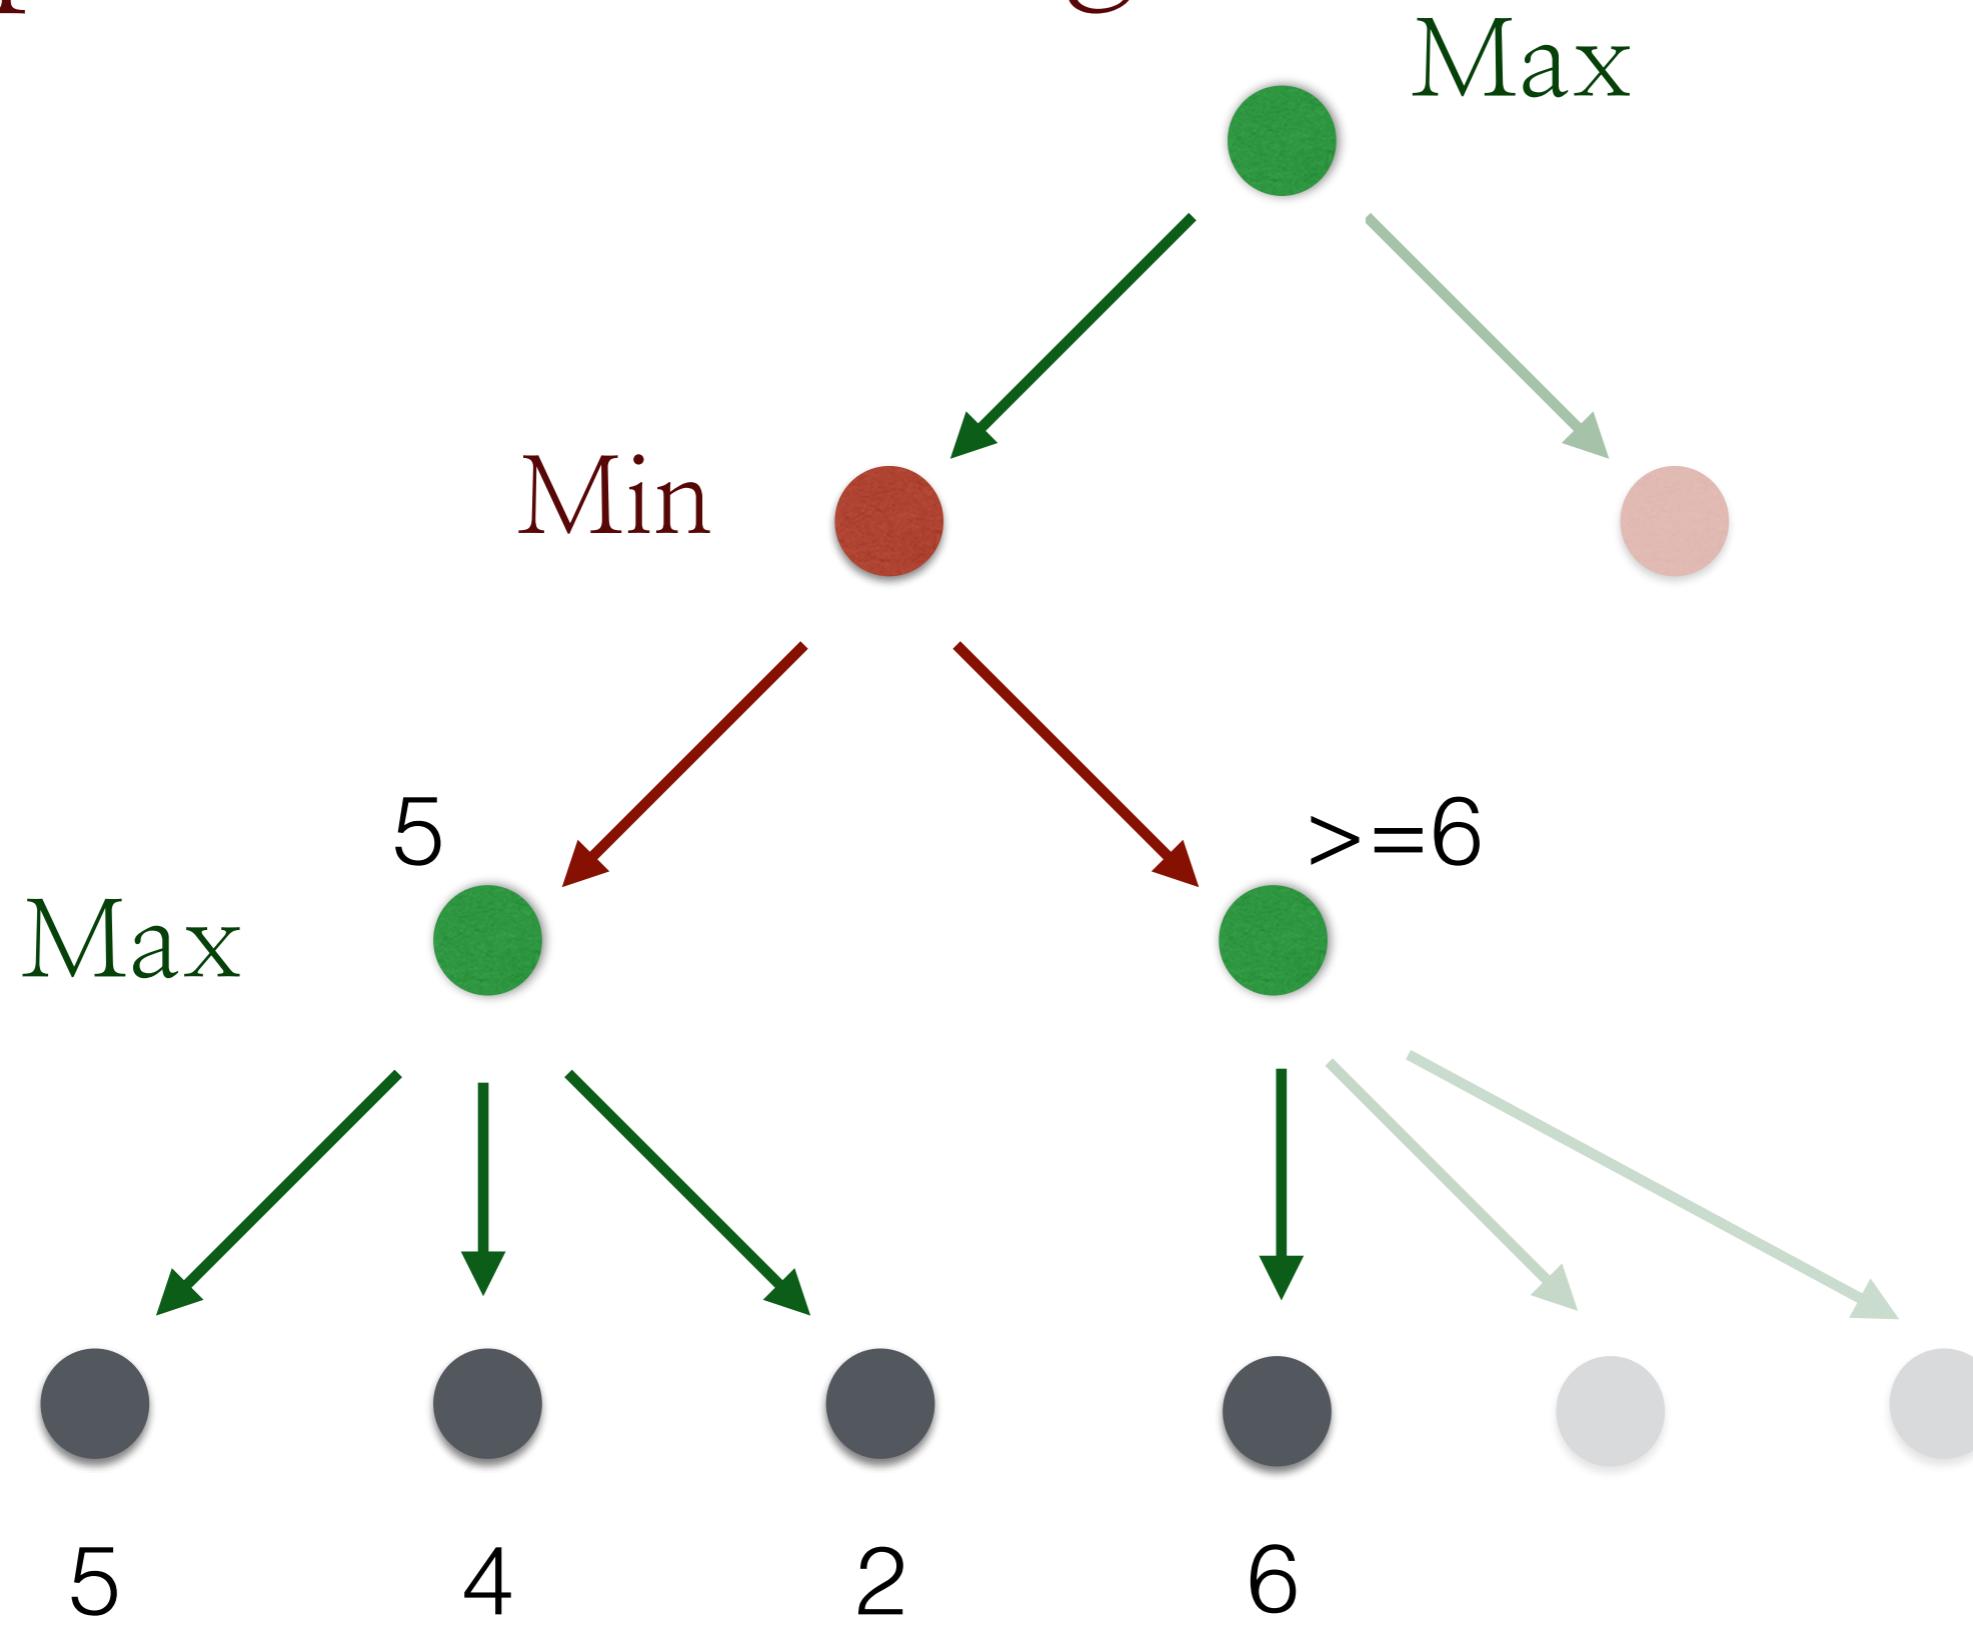
# Alpha–Beta Pruning



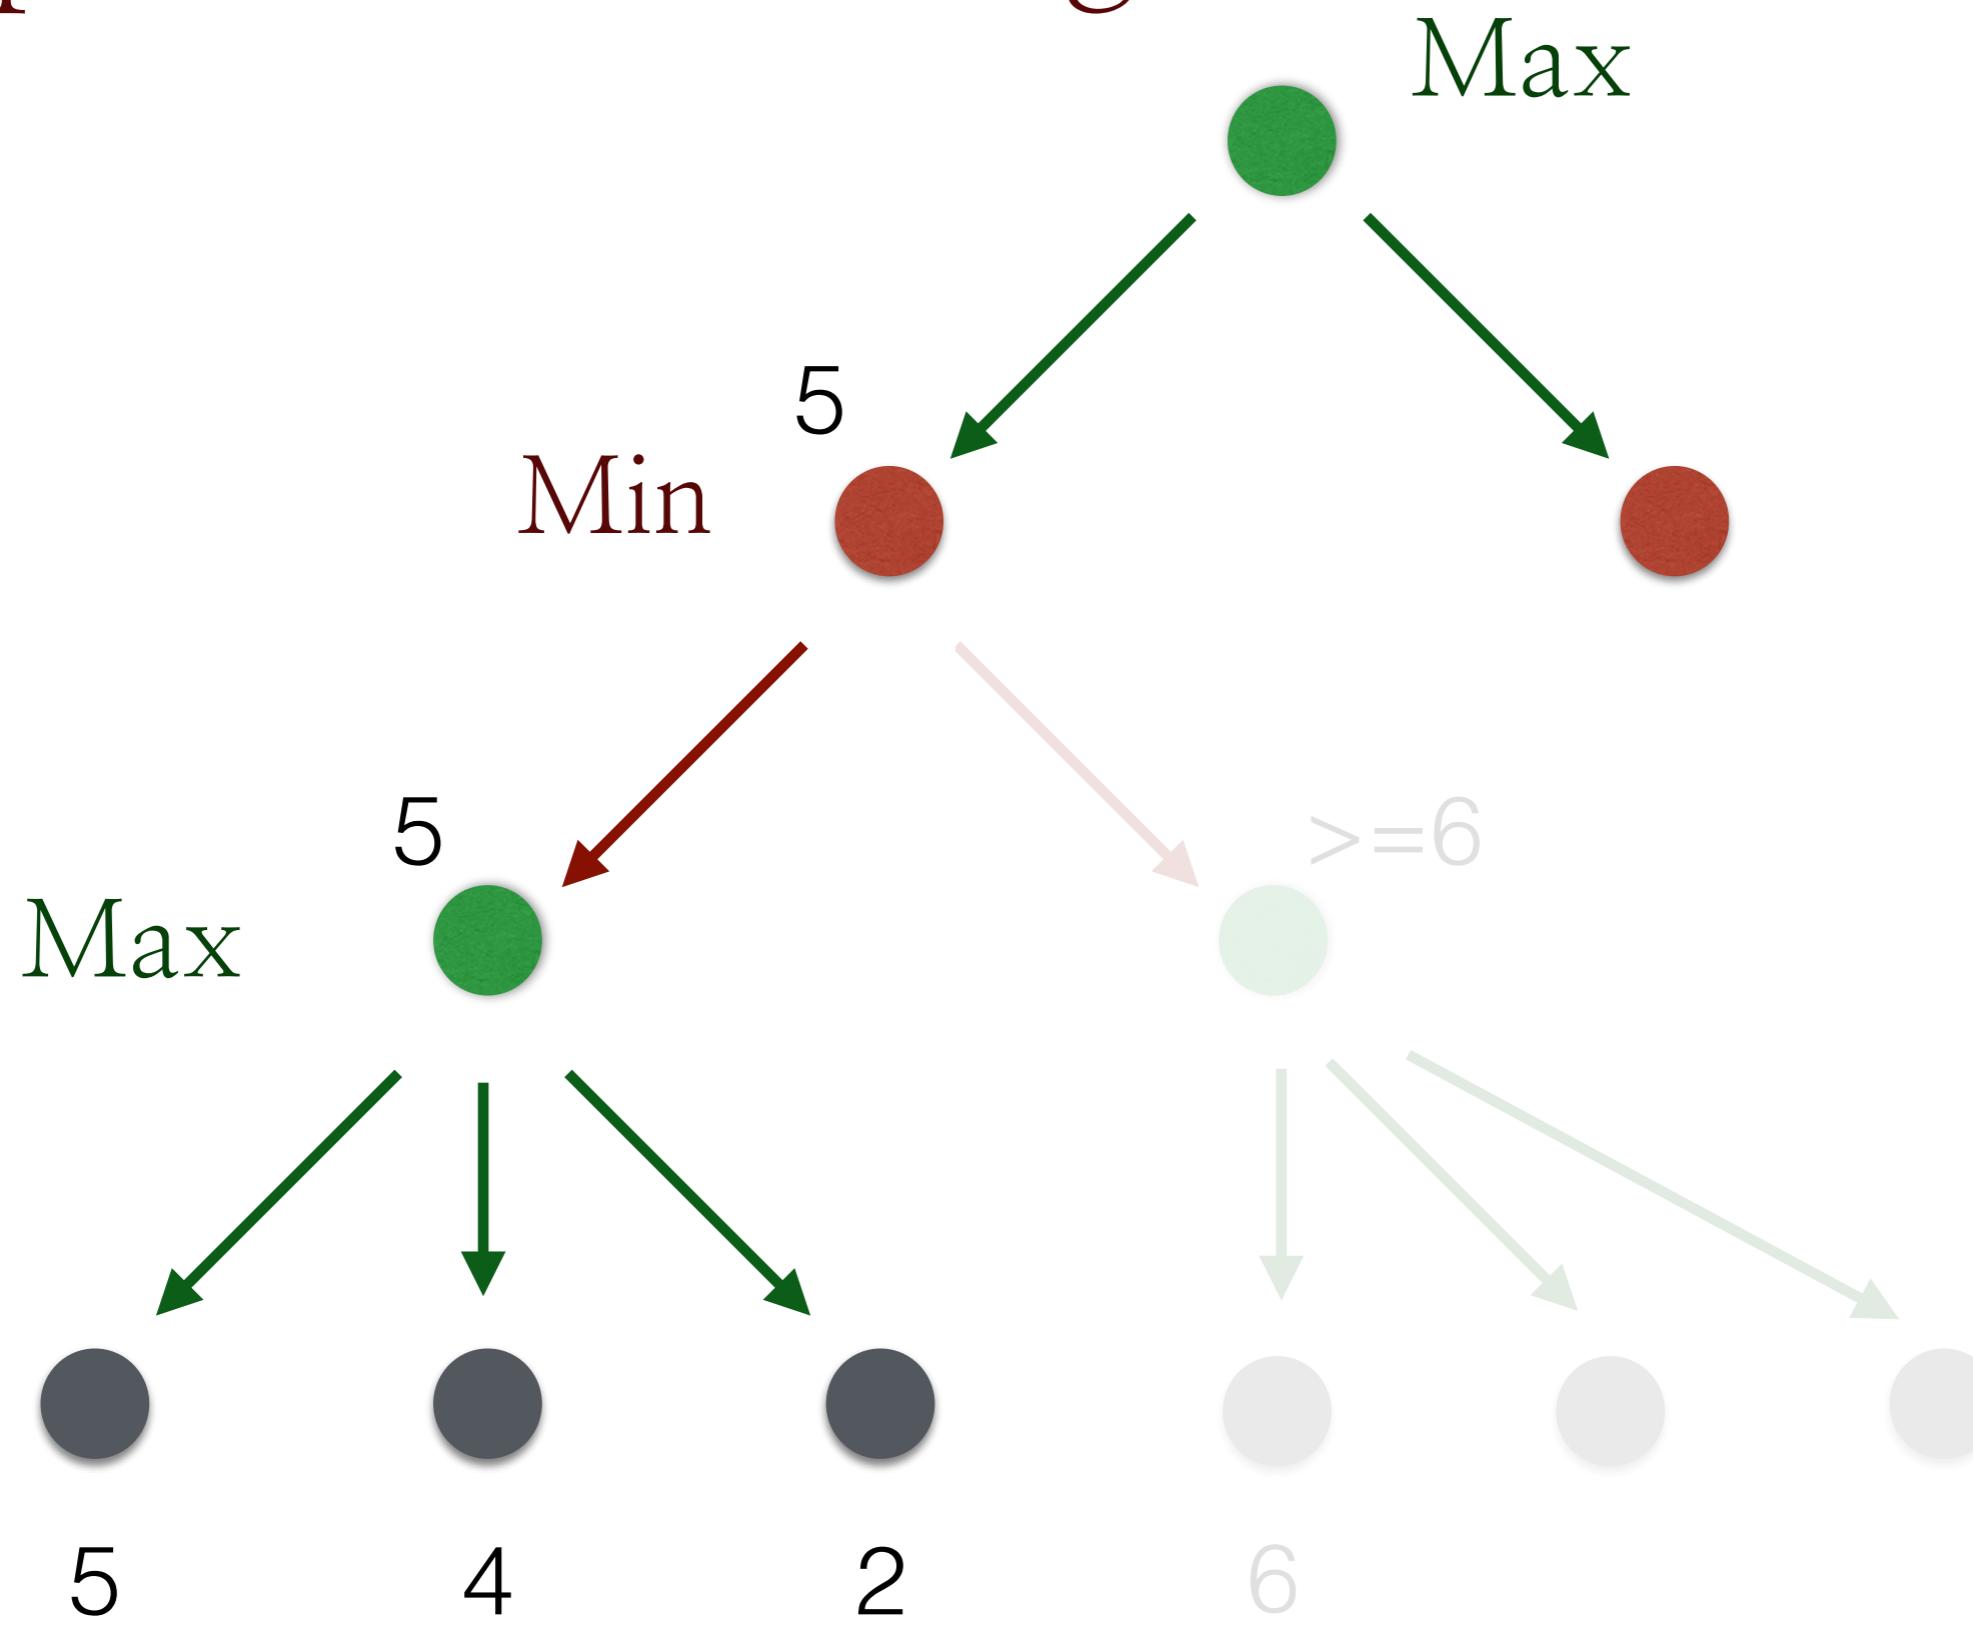
# Alpha–Beta Pruning



# Alpha–Beta Pruning



# Alpha–Beta Pruning



```
alpha = -infinity '''fallback for Max'''
beta = infinity '''fallback for Min'''
node = root

def alphabeta_minimax(node, alpha, beta):
    if terminal(node):
        return payoff(node)
    elif max_player(node):
        value = -infinity
        for n in children(node):
            value = max(value, alphabeta_minimax(n, alpha, beta))
        alpha = max(alpha, value) '''Try to push up'''
        if alpha >= beta: '''If alpha seems too big, stop'''
            break
        return value
    elif min_player(node):
        value = infinity
        for n in children(node):
            value = min(value, alphabeta_minimax(n, alpha, beta))
        beta = min(beta, value) '''Try to push down'''
        if beta <= alpha: '''If beta looks too small, stop'''
            break
        return value
```



